



# Bayesian Statistics and Bayesian Cognitive Modeling (BayesCog)

## Part 2

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)

Faculty of Psychology, University of Vienna

[https://github.com/lei-zhang/BayesCog\\_Part2](https://github.com/lei-zhang/BayesCog_Part2)

lei.zhang@univie.ac.at  
lei-zhang.net  
@lei\_zhang\_lz



# Schedule

## Part 1

04.05: 9:30-13:30

11.05: 9:30-13:30

18.05: 9:30-13:30

25.05: 9:30-13:30

## Part 2

10.06: 9:30-13:30 (different date!)

**16.06: 9:30-13:00** (different date and time)

**21.06: 10:00-13:30** (different date and time)

29.06: 9:30-13:30

## Statistical modeling

Introduction

R Basics

Probability Basics

Bayes' theorem

MCMC and Stan

Single-Parameter Model – Binomial Model

Multiple-Parameter Model – Linear Regression

Inference, Posterior Predictive Check

## Cognitive modeling

Reinforcement Learning Model

Hierarchical Models

Optimizing Stan Codes

Model Comparison

Stan Style Tip and Debugging

Model-Based fMRI

Capstone Project: Delay Discounting Task

# About me

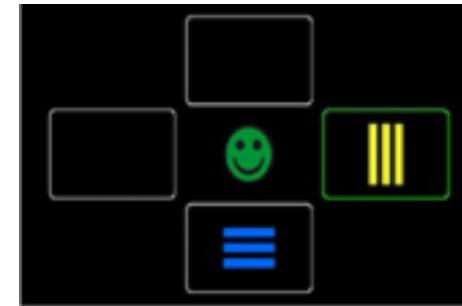
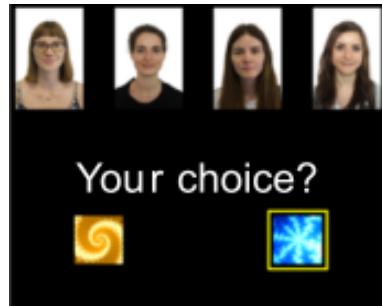
- Current: Postdoc @ [SCAN-Unit](#), w/ [Claus Lamm](#) 
- Ph.D. Cognitive computational neuroscience, w/ [Jan Gläscher](#) 
- M.Sc. Cognitive neuroscience 
- B.Sc. Psychology 
- My journey through computational modeling
  - Started with MLE (@fminsearch in Matlab)
  - Switched to Bayesian: first JAGS, then Stan



# My research:

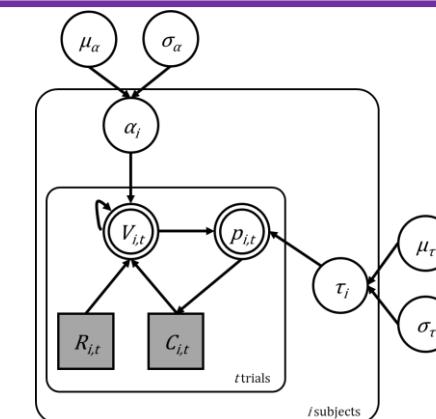
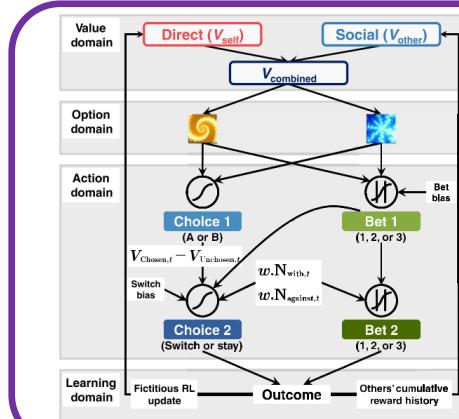
- I ask people to make decisions

Computation



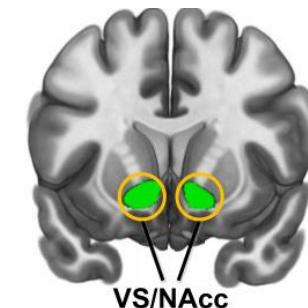
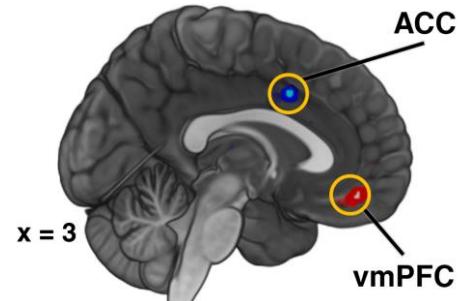
- I build computational models

Algorithm

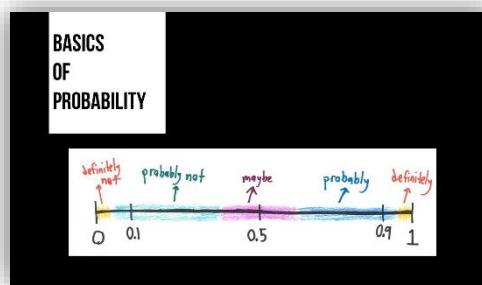


- I examine neural mechanisms

Implementation

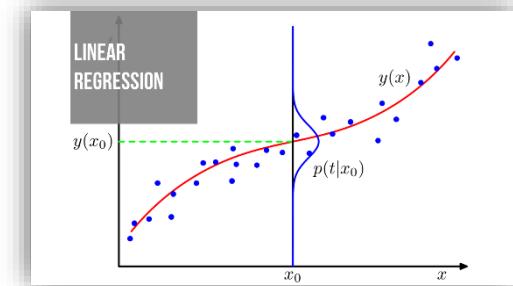
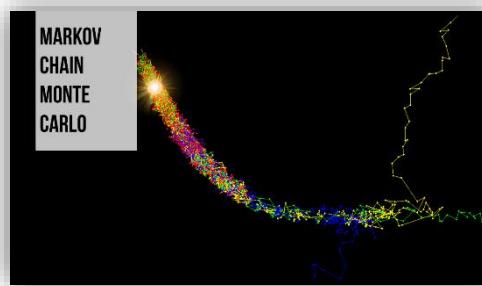
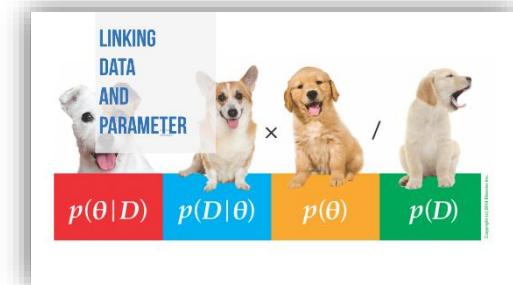


# Recap of Topics

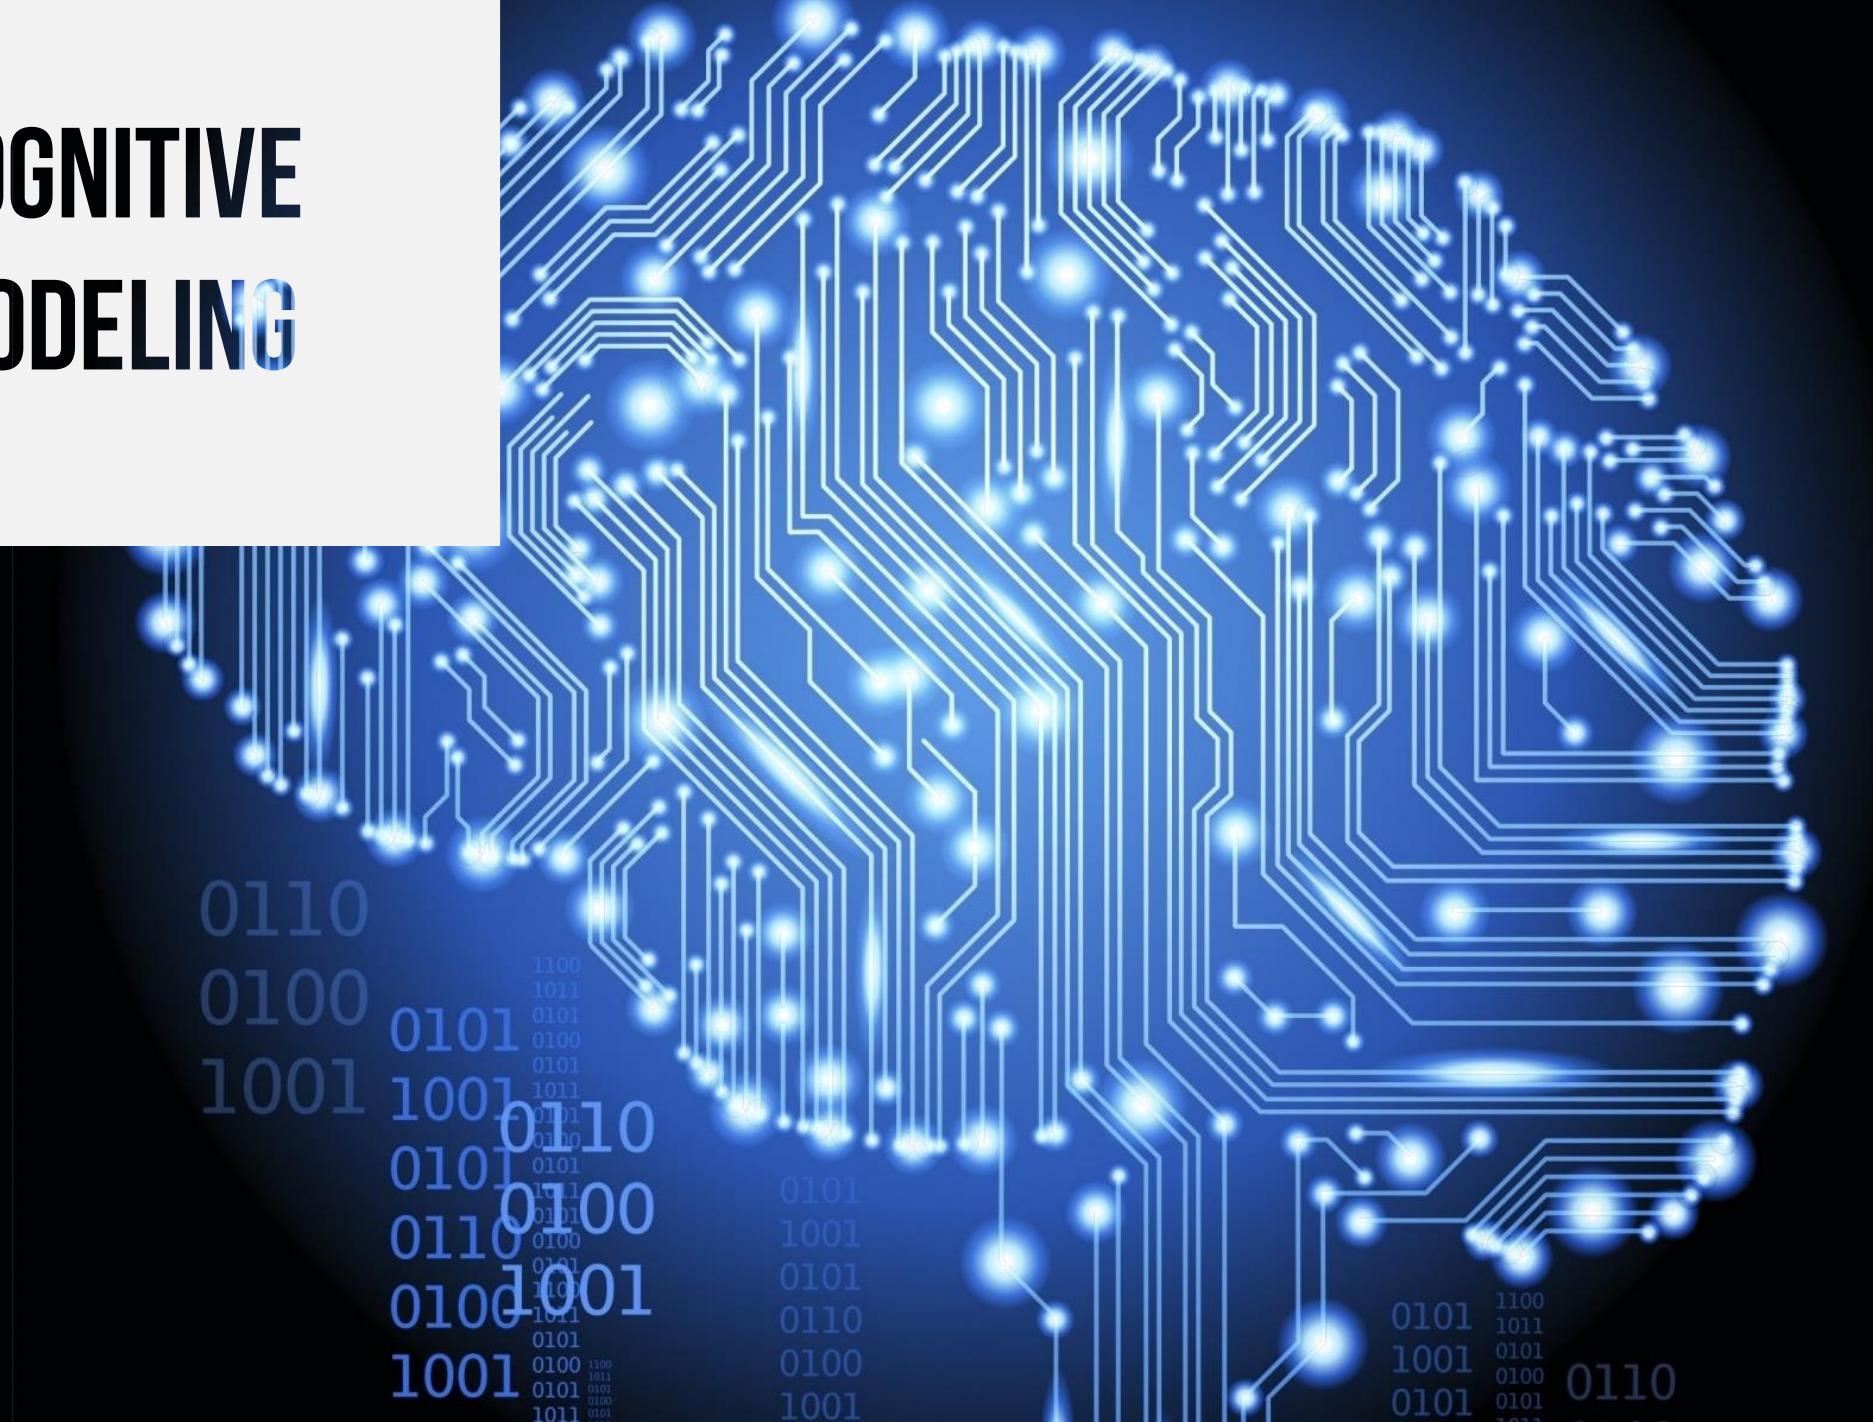


BAYES' THEOREM

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

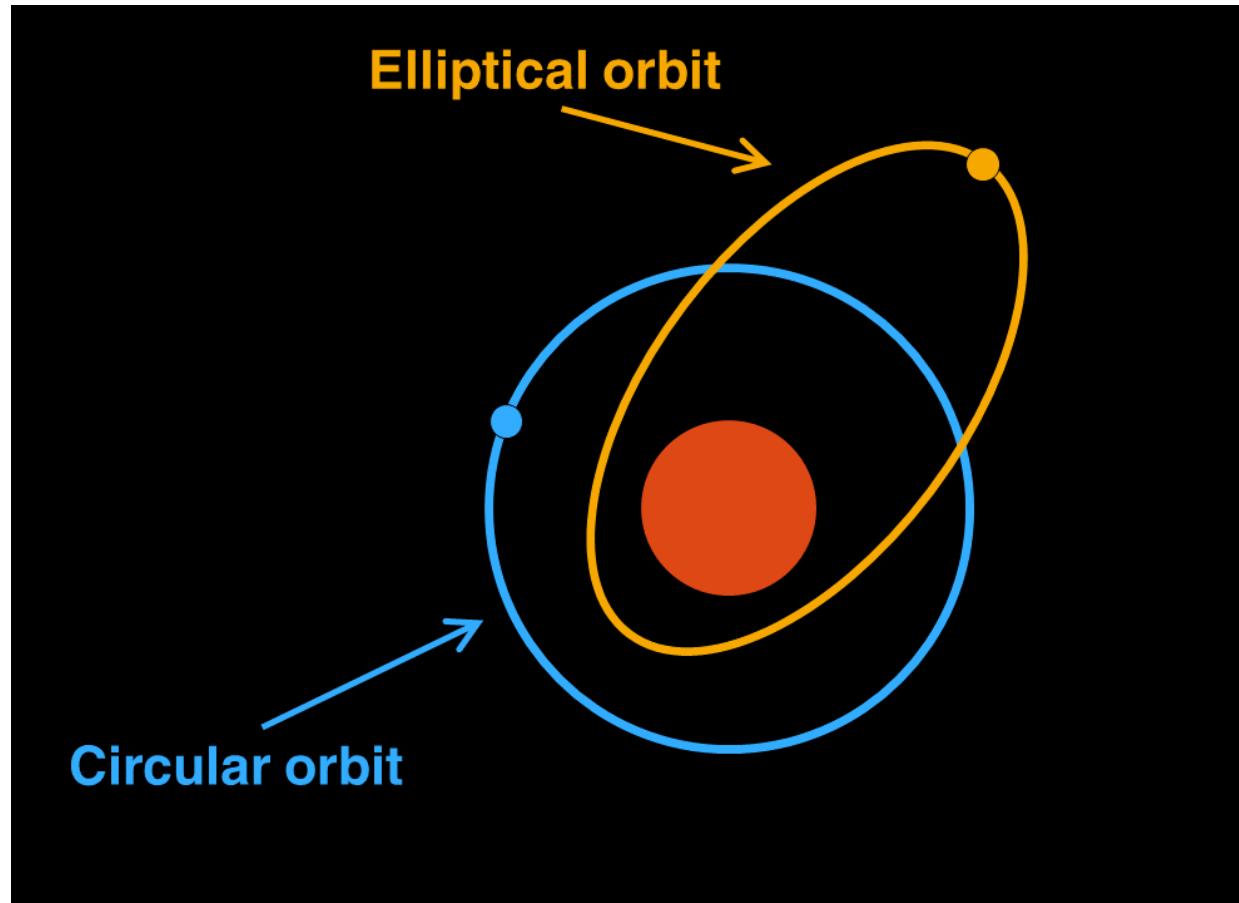


# COGNITIVE MODELING



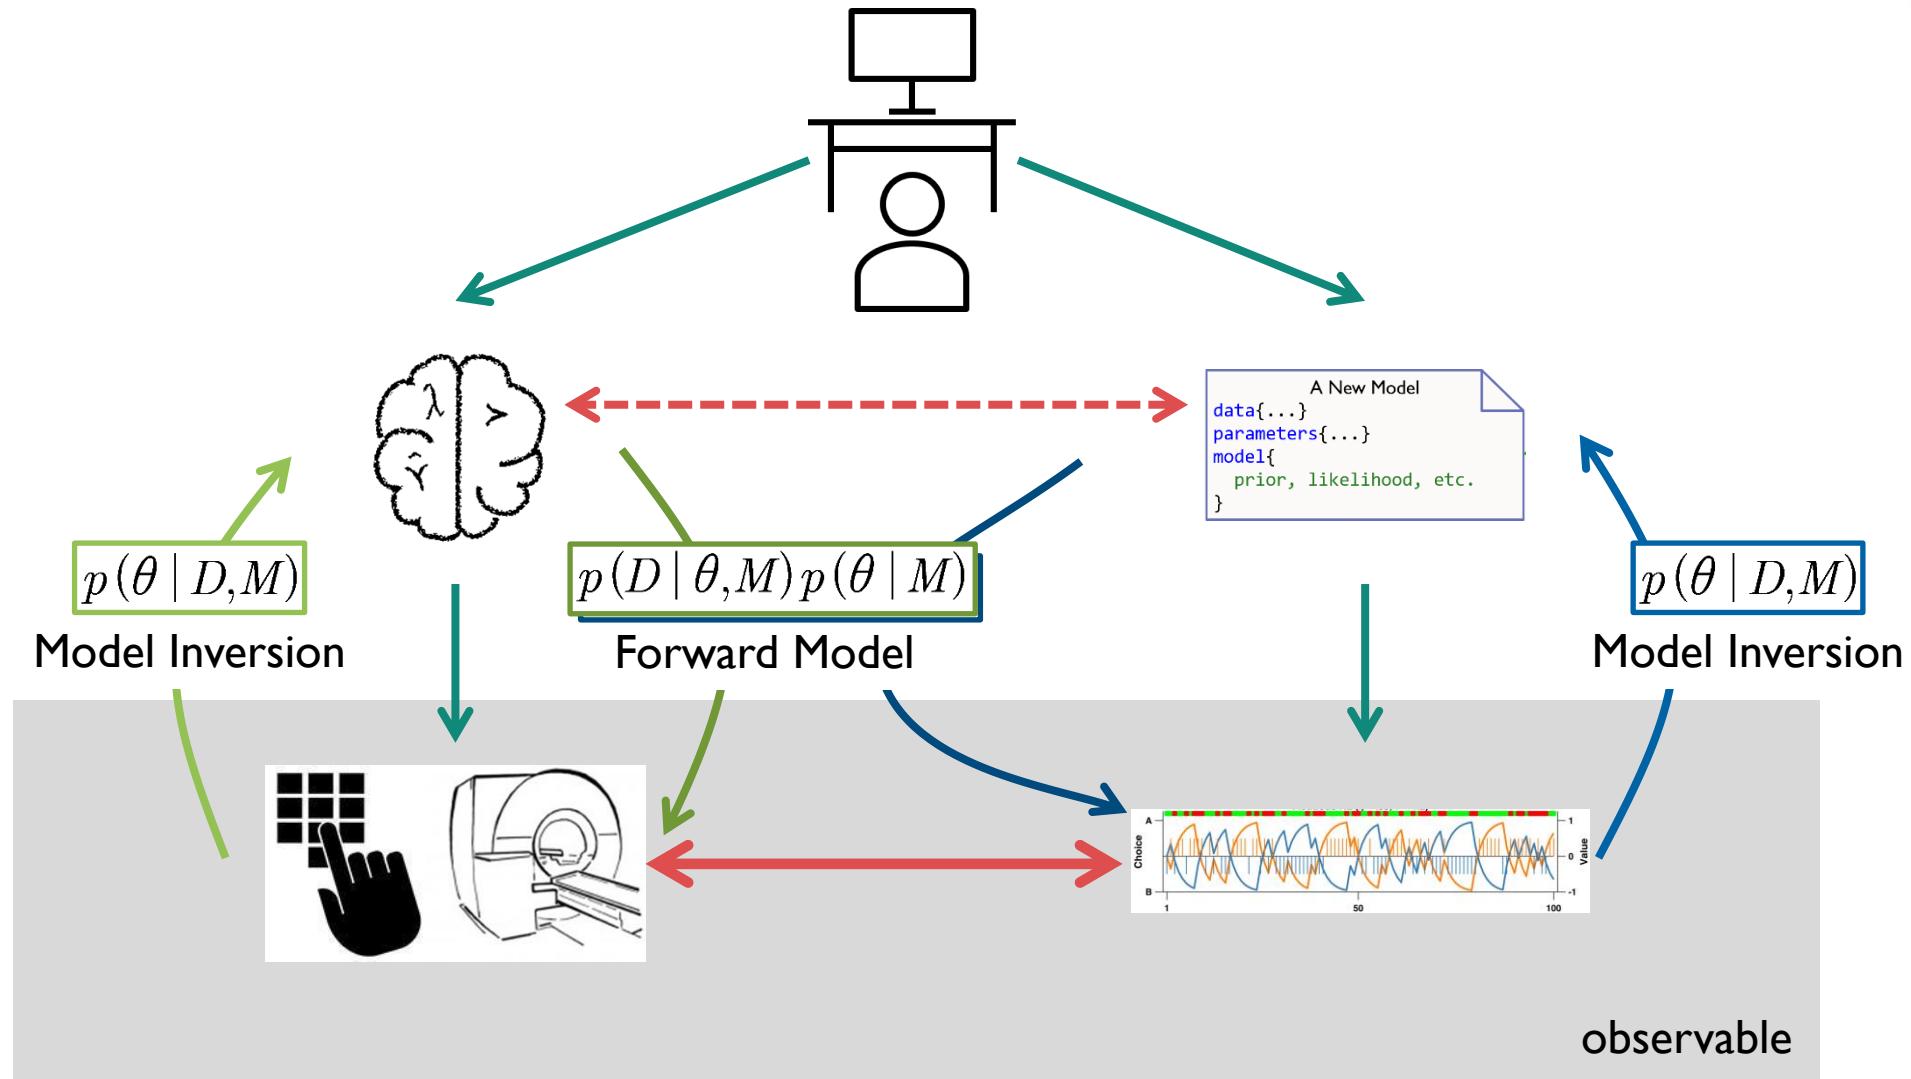
# The idea of computational modeling is never new

Scientists use mathematical models to approximate certain processes (physical or mental), in order to explain and to predict.



# What is Cognitive Modeling?

cognitive model  
statistics  
computing



Essentially, all the models are wrong, but some are useful.



– George E. P. Box

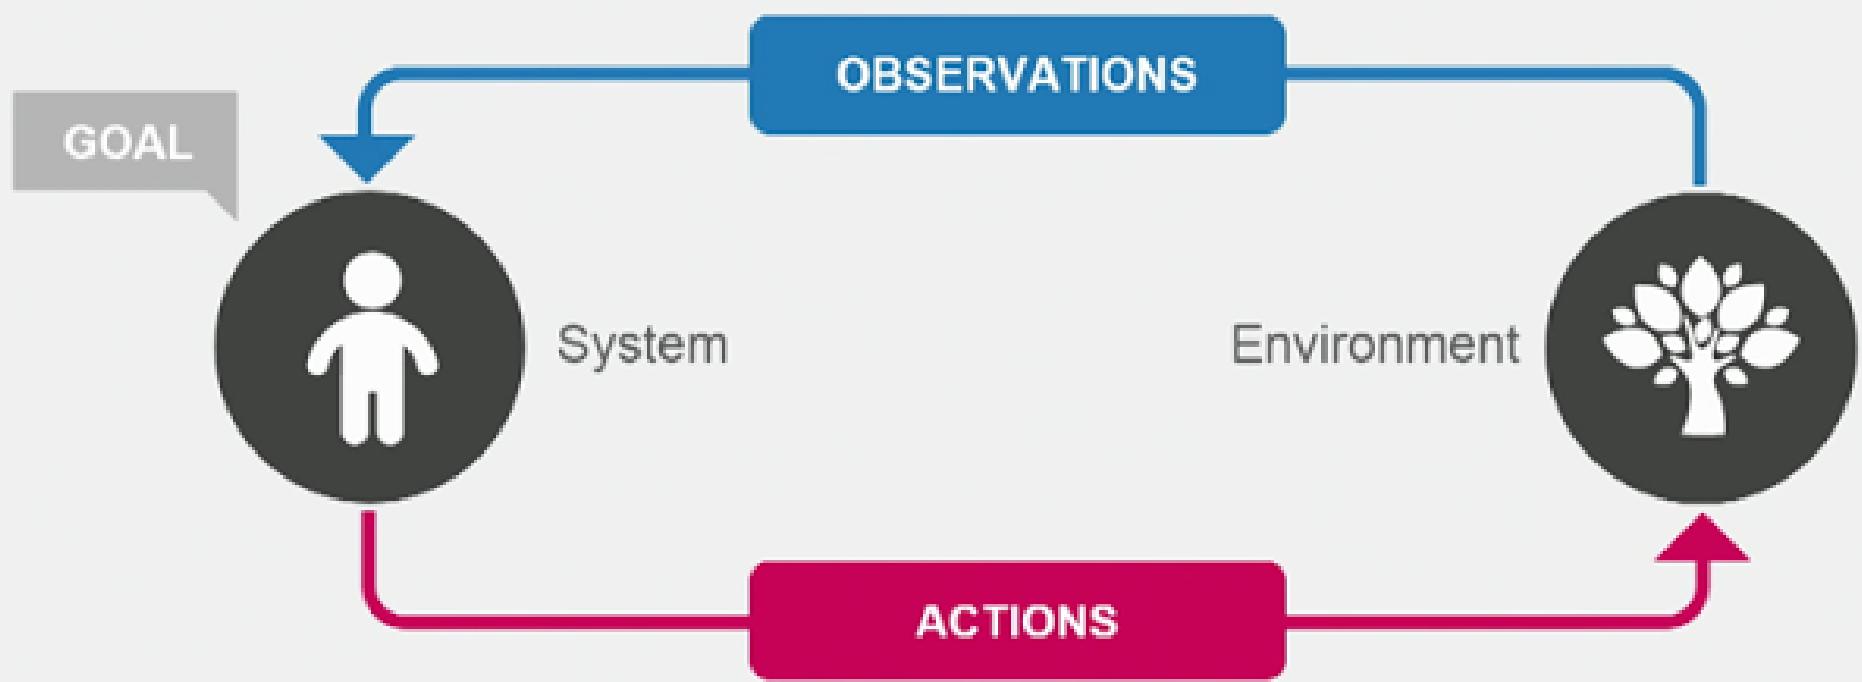
---

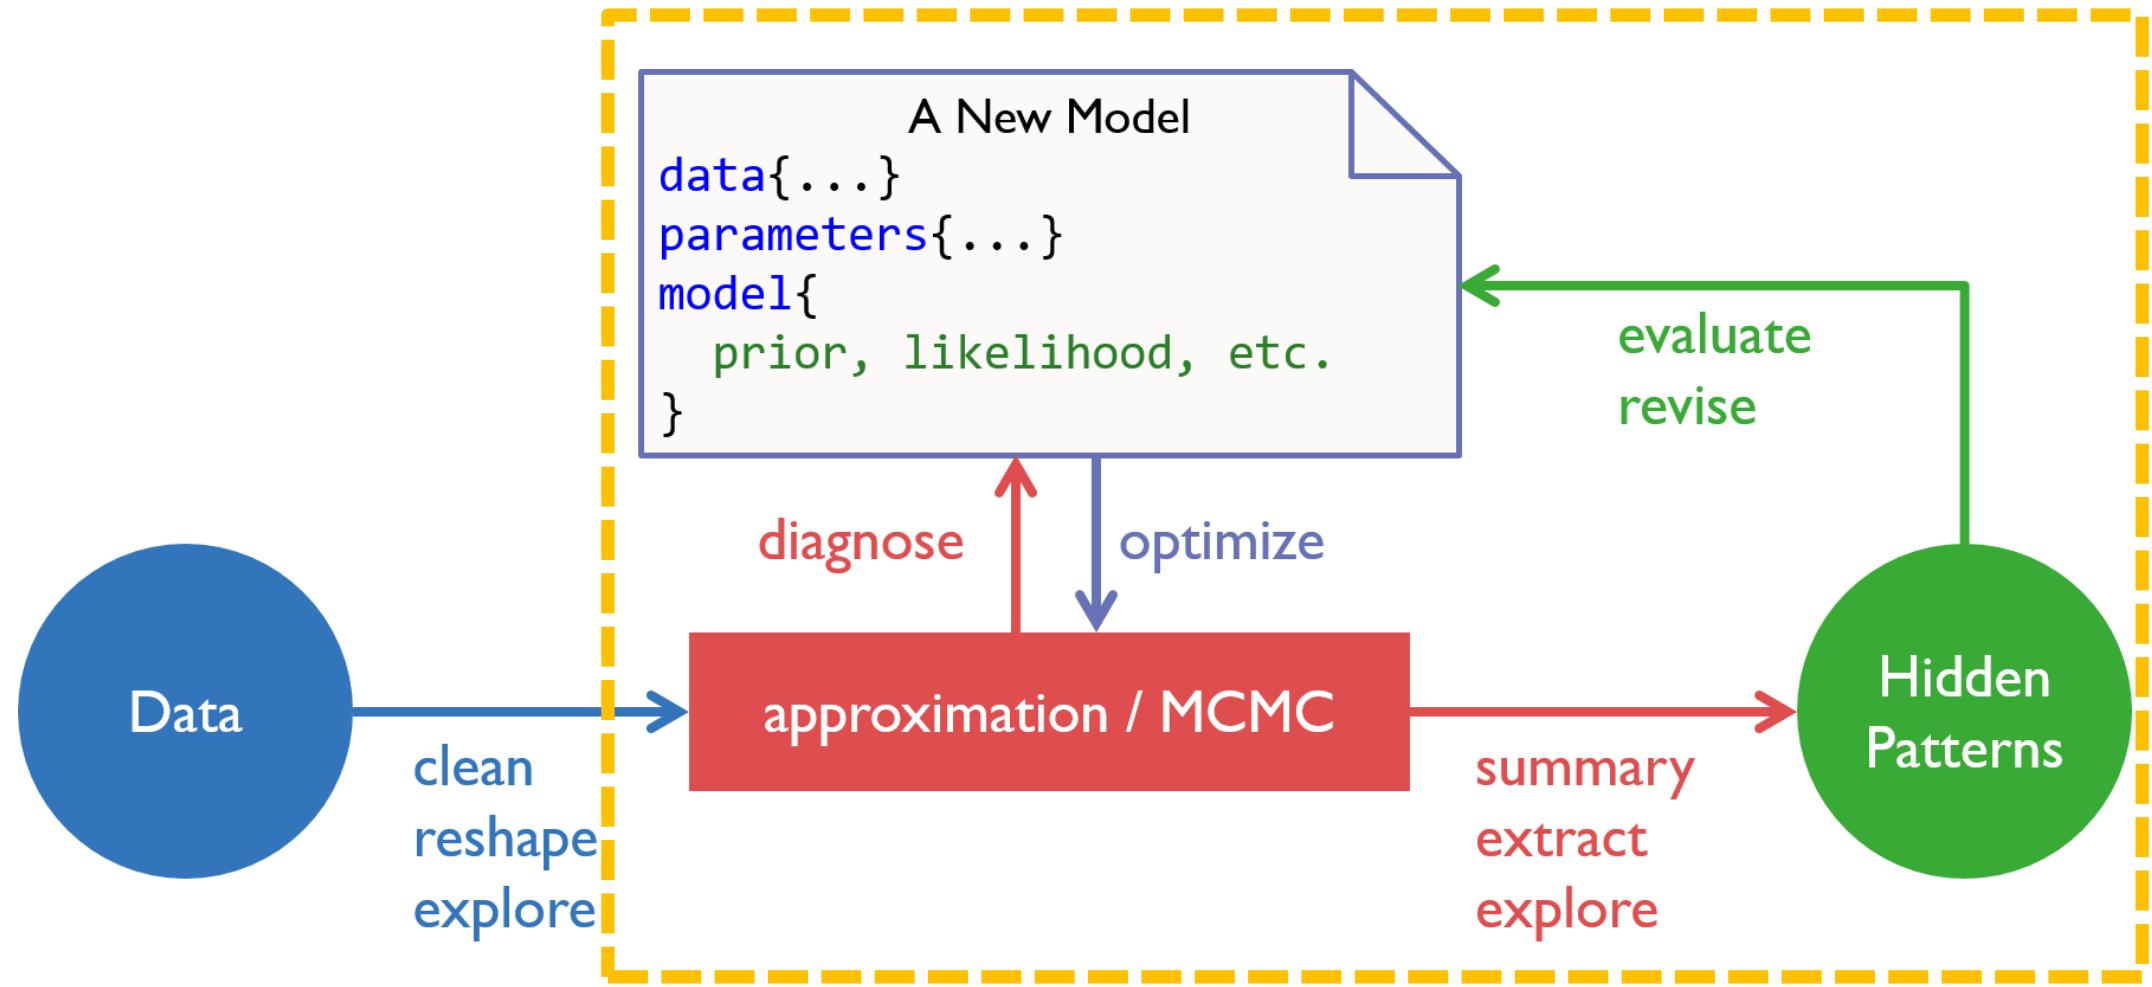
Essentially, all the models are ~~wrong~~ imperfect, but some are useful.

# Common cognitive models in decision-making

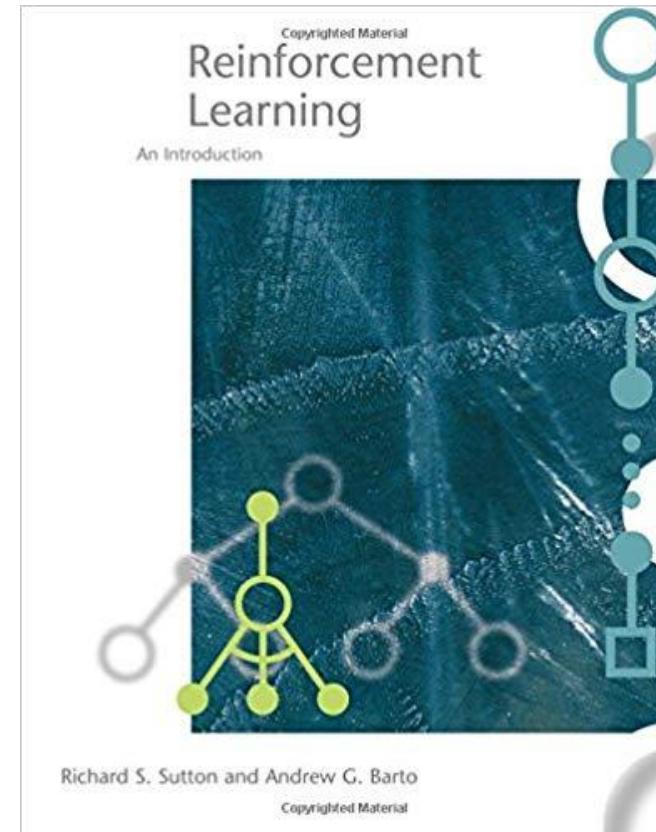
- Reinforcement learning model
- Bayesian learning model
- Risk-aversion model
- Hyperbolic delay discounting model
- Fehr-Schmidt inequity aversion model
- Sequential sampling model
- Experience-weighted attraction model
- ...

# REINFORCEMENT LEARNING FRAMEWORK

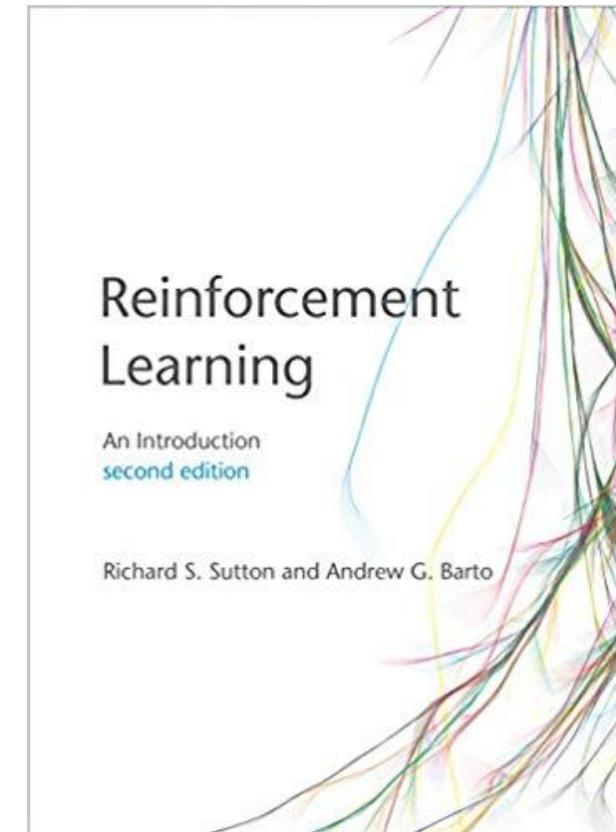




# The very short history



1998

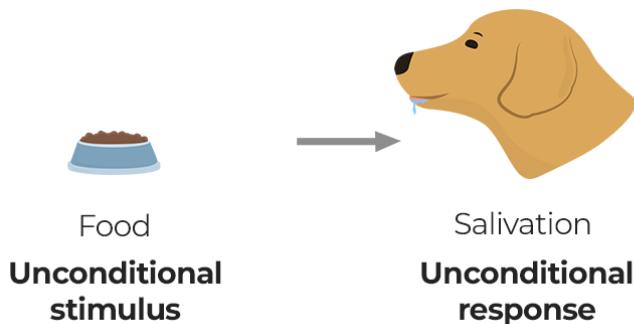


2018

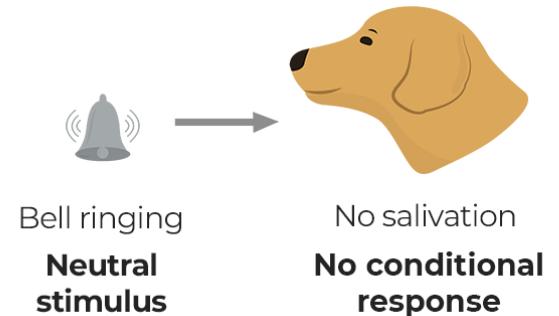
# why is it relevant?

cognitive model  
statistics  
computing

1. Before conditioning



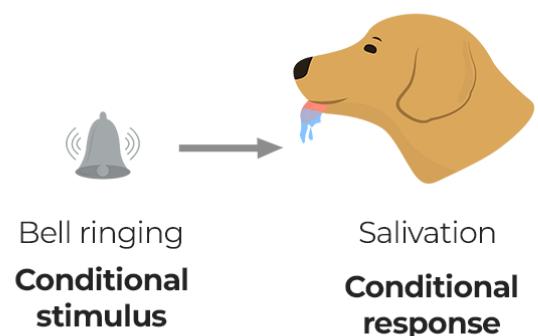
2. Before conditioning



3. During conditioning



4. After conditioning



# Classic examples

**REPORT**

## Dissociable Roles of Ventral and Dorsal Striatum in Instrumental Conditioning

John O'Doherty<sup>1,\*</sup>, Peter Dayan<sup>2</sup>, Johannes Schultz<sup>1</sup>, Ralf Deichmann<sup>1</sup>, Karl Friston<sup>1</sup>, Raymond J. Dolan<sup>1</sup>

+ See all authors and affiliations

Science 16 Apr 2004:  
Vol. 304, Issue 5669, pp. 452-454  
DOI: 10.1126/science.1094285

Published: 06 November 2005

## Uncertainty-based competition between prefrontal and dorsolateral striatal systems for behavioral control

Nathaniel D Daw✉, Yael Niv & Peter Dayan

Nature Neuroscience 8, 1704–1711 (2005) | Cite this article

**Neuron**  
Volume 66, Issue 4, 27 May 2010, Pages 585-595

Article

## States versus Rewards: Dissociable Neural Prediction Error Signals Underlying Model-Based and Model-Free Reinforcement Learning

Jan Gläscher <sup>1, 3</sup>✉, Nathaniel Daw <sup>4</sup>, Peter Dayan <sup>5</sup>, John P. O'Doherty <sup>1, 2, 6</sup>

Open Access | Published: 07 June 2011

## Empirical support for an involvement of the mesostriatal dopamine system in human fear extinction

K A Raczka, M-L Mechias, N Gartmann, A Reif, J Deckert, M Pessiglione & R Kalisch✉

Translational Psychiatry 1, e12 (2011) | Cite this article

# Very recent examples

cognitive model

statistics

computing

REPORT

## Dopamine promotes cognitive effort by biasing the benefits versus costs of cognitive work

A. Westbrook<sup>1,2,3,\*</sup>, R. van den Bosch<sup>2,3</sup>, J. I. Määttä<sup>2,3</sup>, L. Hofmans<sup>2,3</sup>, D. Papadopetraki<sup>2,3</sup>, R. Cools<sup>2,3,†</sup>, M. J. Frank<sup>1,4,†</sup>

+ See all authors and affiliations

Science 20 Mar 2020:  
Vol. 367, Issue 6484, pp. 1362-1366  
DOI: 10.1126/science.aaz5891

Neuron

Available online 17 March 2020

In Press, Corrected Proof 

Article

## A Neuro-computational Account of Arbitration between Choice Imitation and Goal Emulation during Human Observational Learning

Caroline J. Charpentier<sup>1, 2</sup>  , Kiyohito ligaya<sup>1</sup>, John P. O'Doherty<sup>1</sup>



3 out of 4 focused on Reinforcement Learning models!

nature reviews  
neuroscience

Review Article | Published: 12 March 2020

## The neural and computational systems of social learning

Andreas Olsson , Ewelina Knapska & Björn Lindström

Nature Reviews Neuroscience 21, 197–212(2020) | Cite this article

Translational  
Psychiatry

Article | Open Access | Published: 17 March 2020

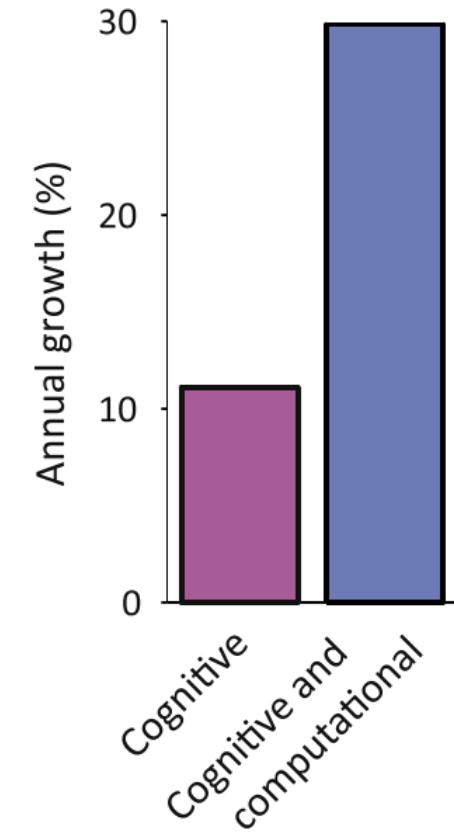
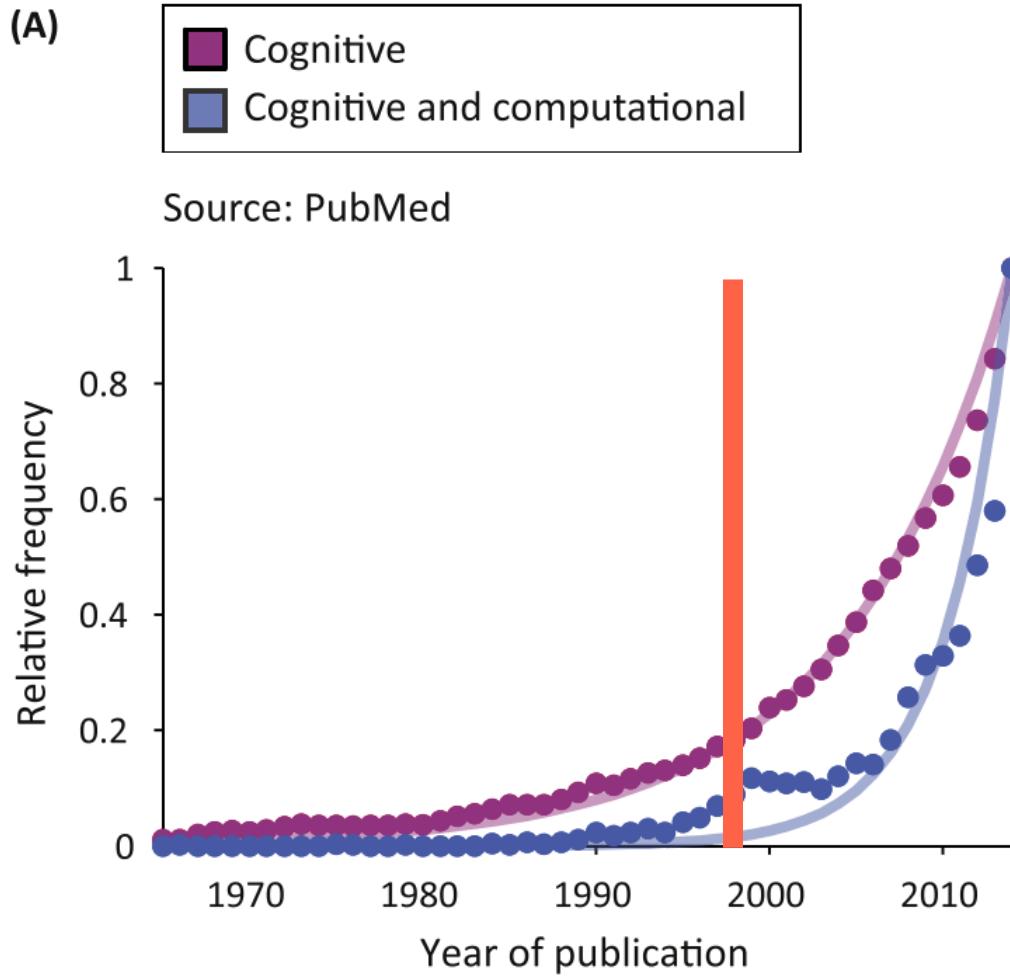
## Neurocomputational mechanisms underpinning aberrant social learning in young adults with low self-esteem

Geert-Jan Will , Michael Moutoussis, Palee M. Womack, Edward T. Bullmore, Ian M. Goodyer, Peter Fonagy, Peter B. Jones, NSPN Consortium, Robb B. Rutledge & Raymond J. Dolan

# Boom in Cognitive Modeling

cognitive model  
statistics  
computing

(A)



# 2-armed bandit task



a simple task often used in the laboratory:

- **repeated choice** between N options (**N-armed bandit**)
- ...whose properties (reward amounts, probabilities) are learned through **trial-and-error**
- ...with a **goal** in mind: maximize the overall reward

# 2-armed bandit task

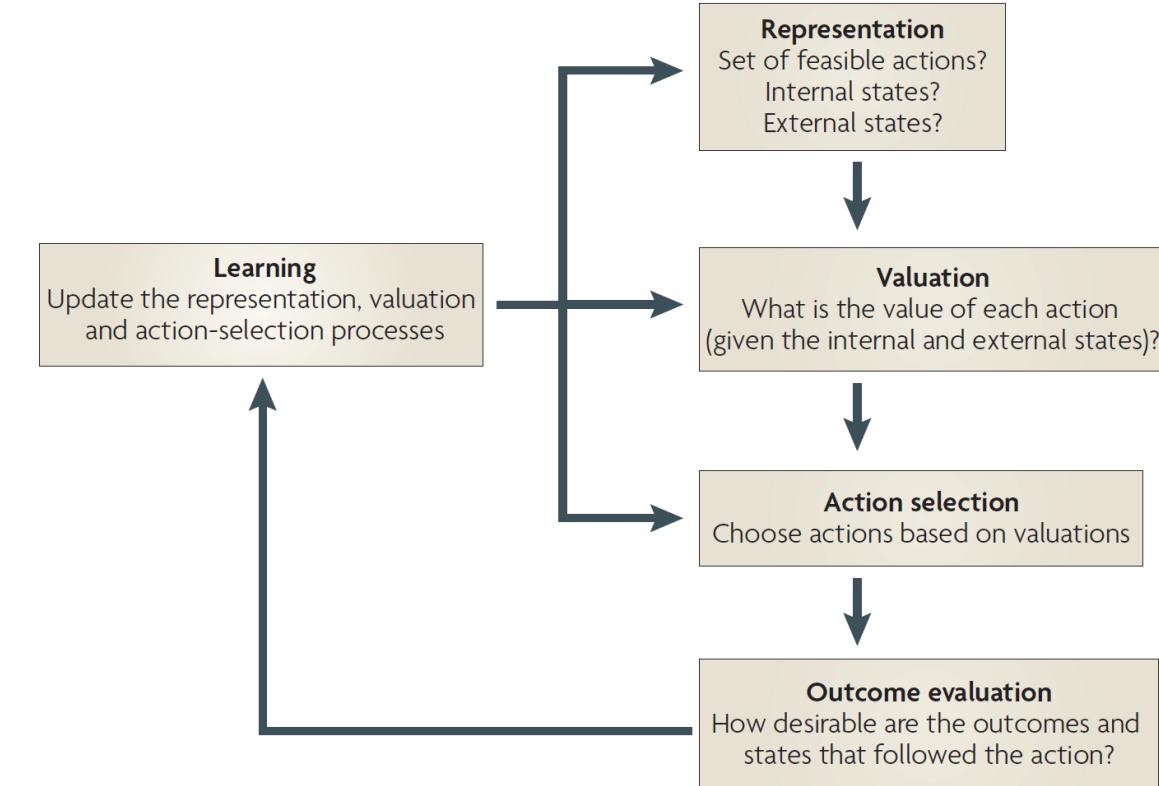
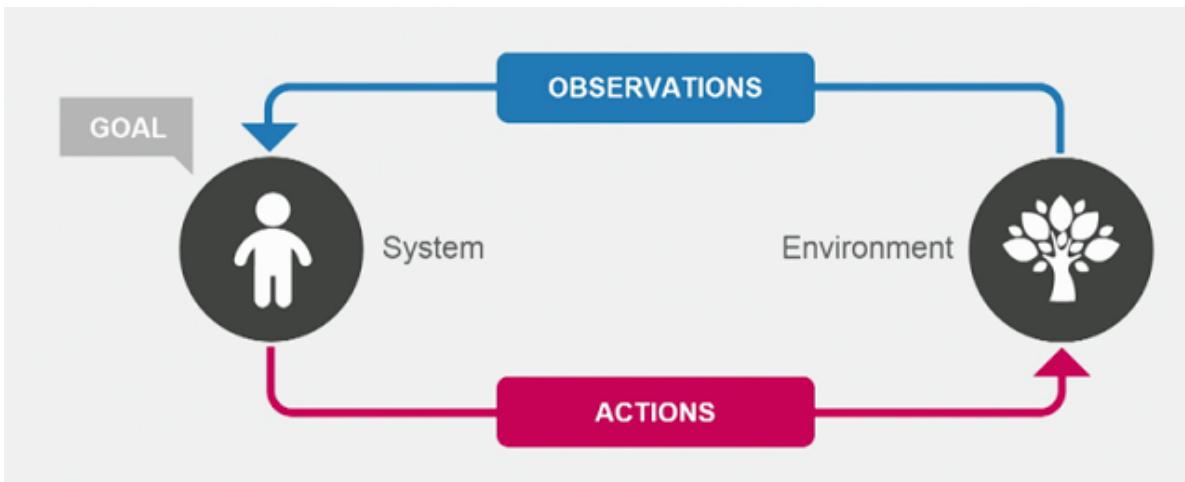


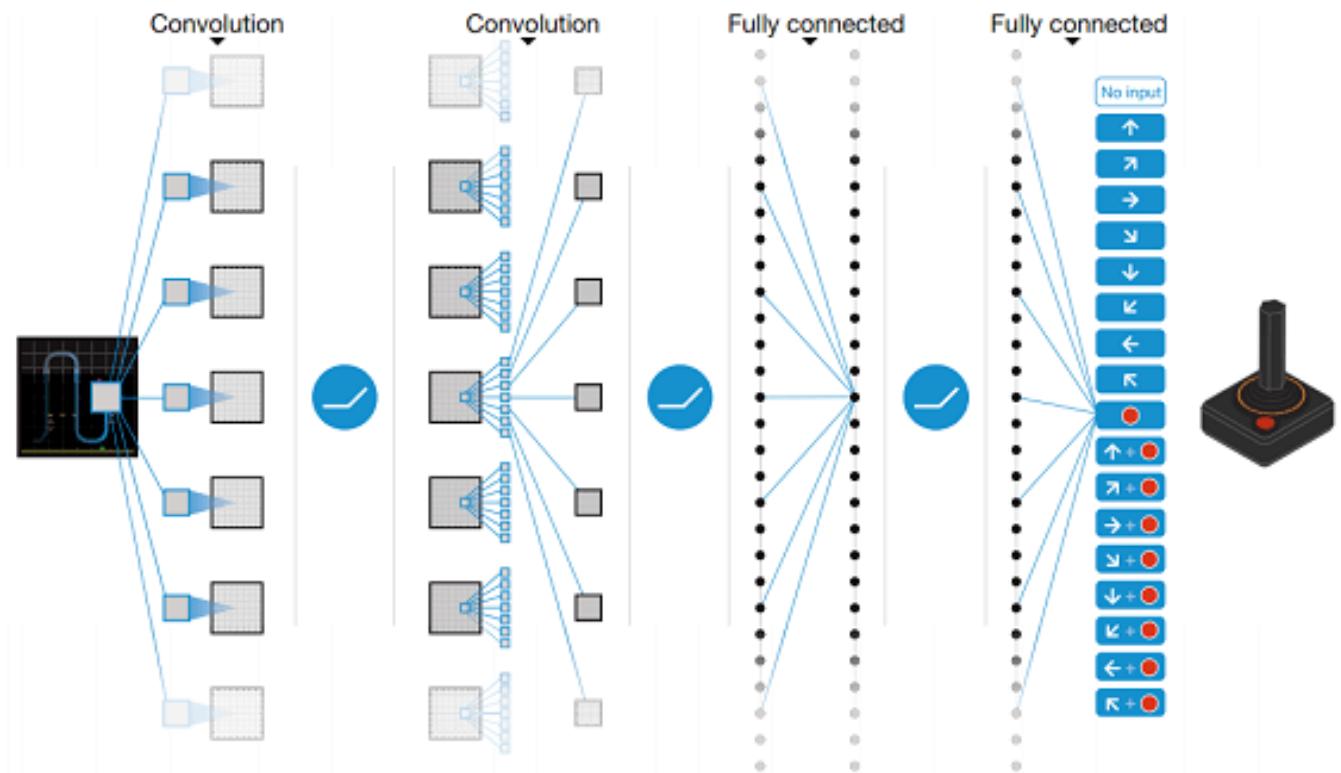
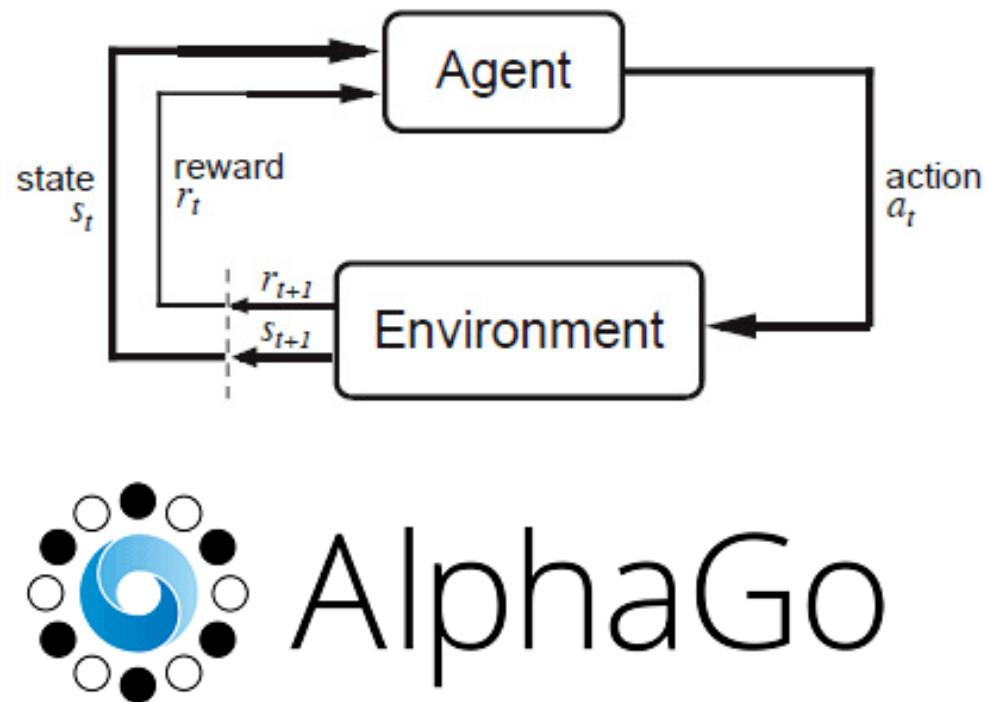
What can be your **strategies**:

1. **predict** the value of each deck
2. **choose** the best
3. **learn** from outcome to update predictions  
**(repeat)**

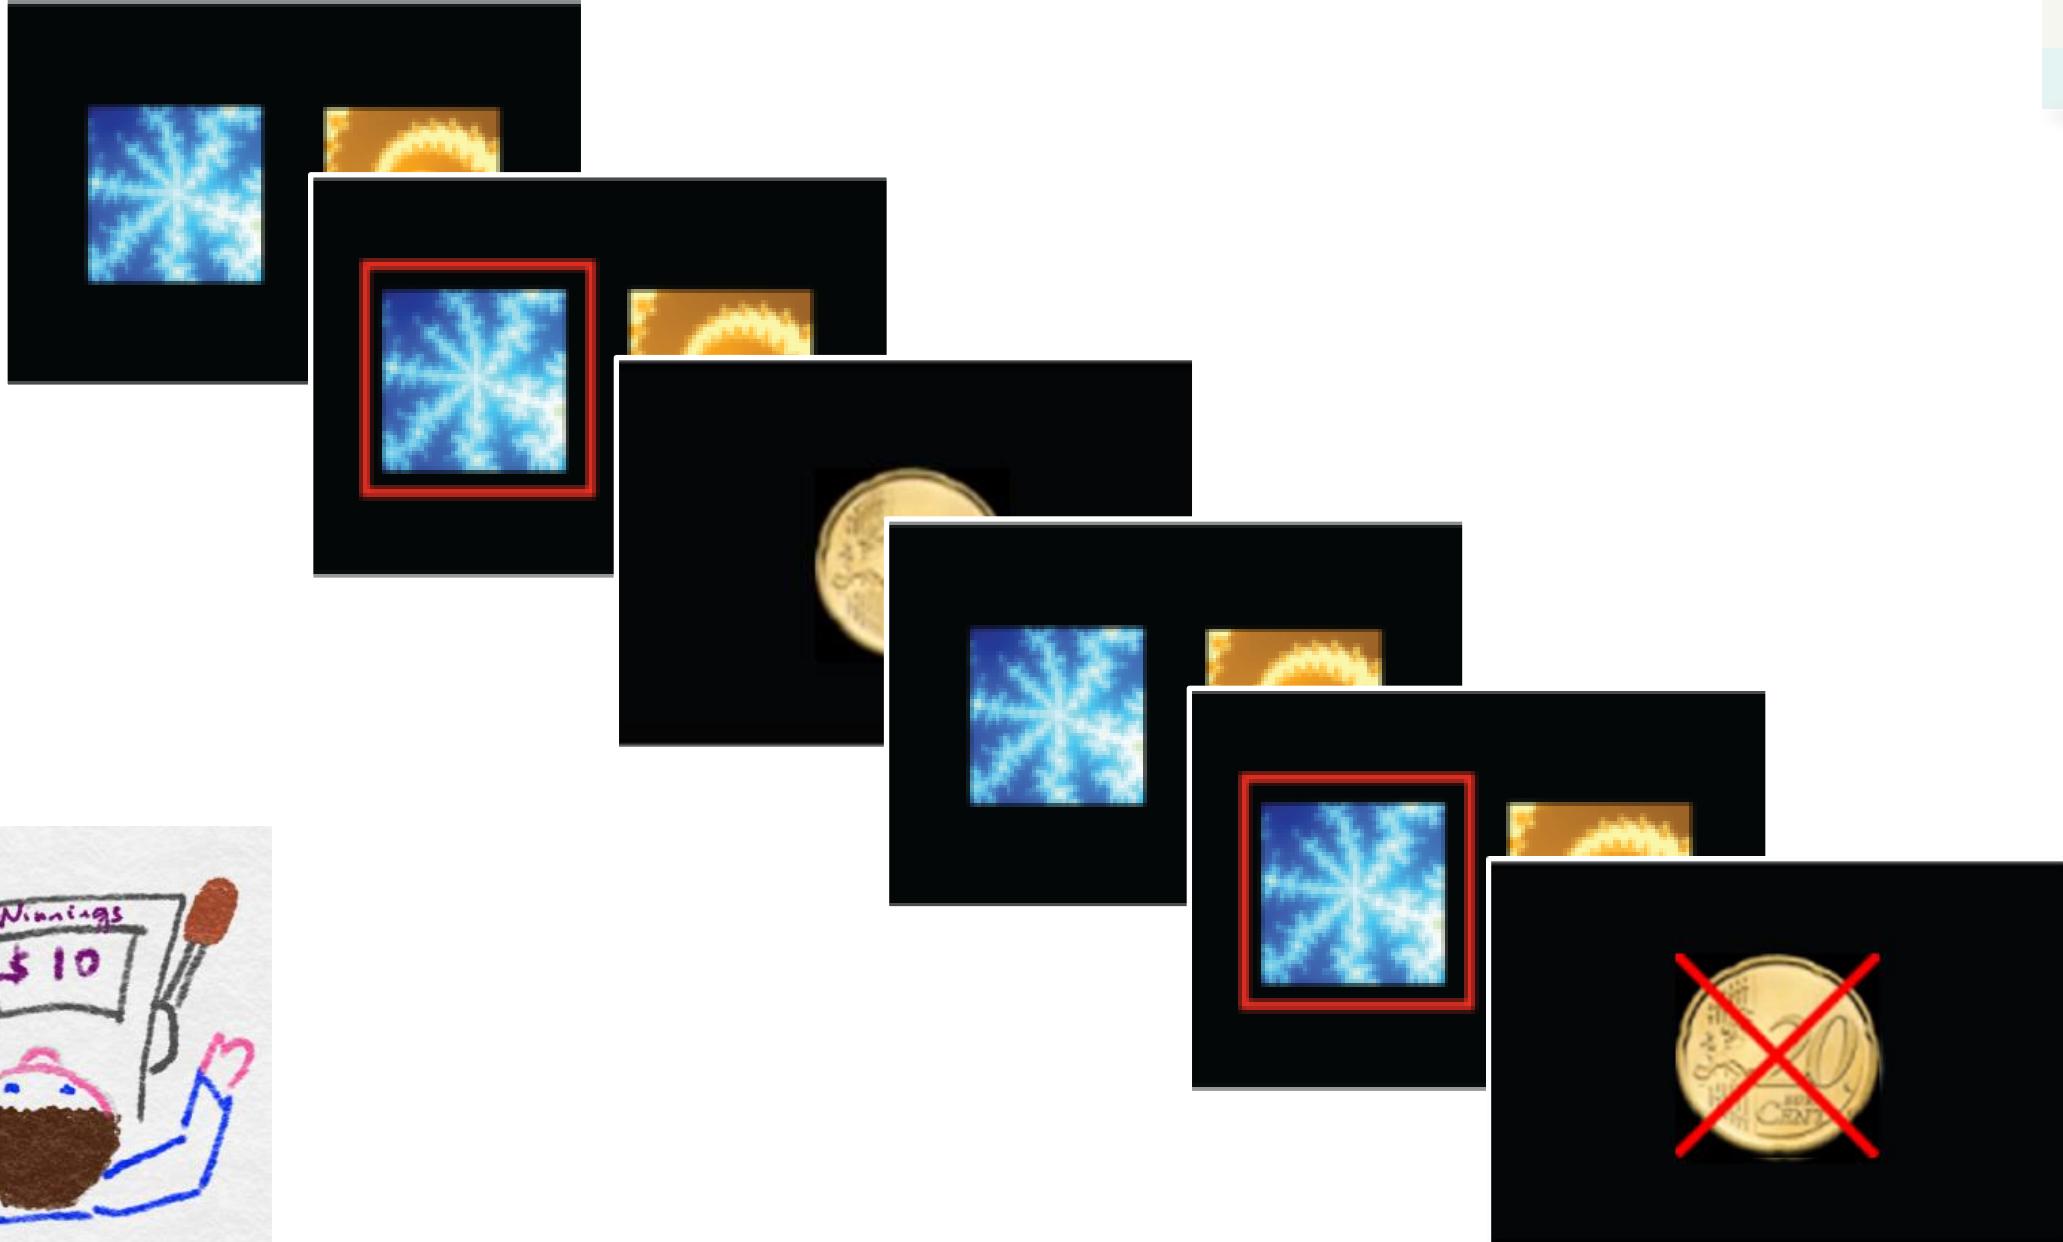
# How prediction is shaped by learning?

cognitive model  
statistics  
computing

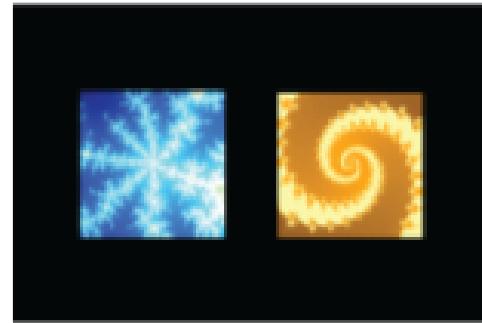




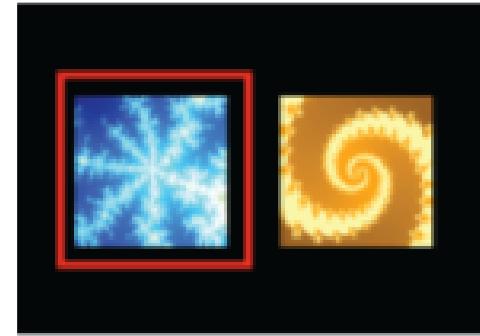
# AlphaGo



# One simple experiment: two choice task



choice presentation



action selection



outcome

what do we know?

what can we measure?

what do we not know?

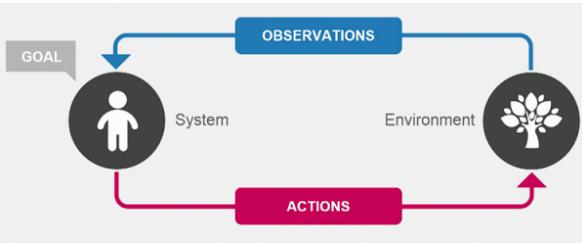
Data: choice & outcome

Summary stats: choice accuracy

Learning algorithm: RL update

$p(\text{choosing the better option})$

# Rescorla-Wagner Value Update



## Cognitive Model

- cognitive process
- using internal variables and free parameters

## Observation Model (Data Model)

- relate model to observed data
- has to account for noise

# Rescorla-Wagner (1972)

- The idea: **error-driven** learning
- Change in value is proportional to the difference between actual and predicted outcome



Robert A. Rescorla

Allan R. Wagner



Value update:  $V_t = V_{t-1} + \alpha * PE_{t-1}$

Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$

$\alpha$  - learning rate  
PE - reward prediction error  
V - value  
R - reward

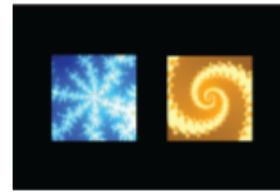
*Expectations on the next trial = the expectation on the current trial + learning rate \*  
prediction error (reward – current expectation)*

# Understand the learning rate

cognitive model  
statistics  
computing

Value update:  $V_t = V_{t-1} + \alpha * PE_{t-1}$

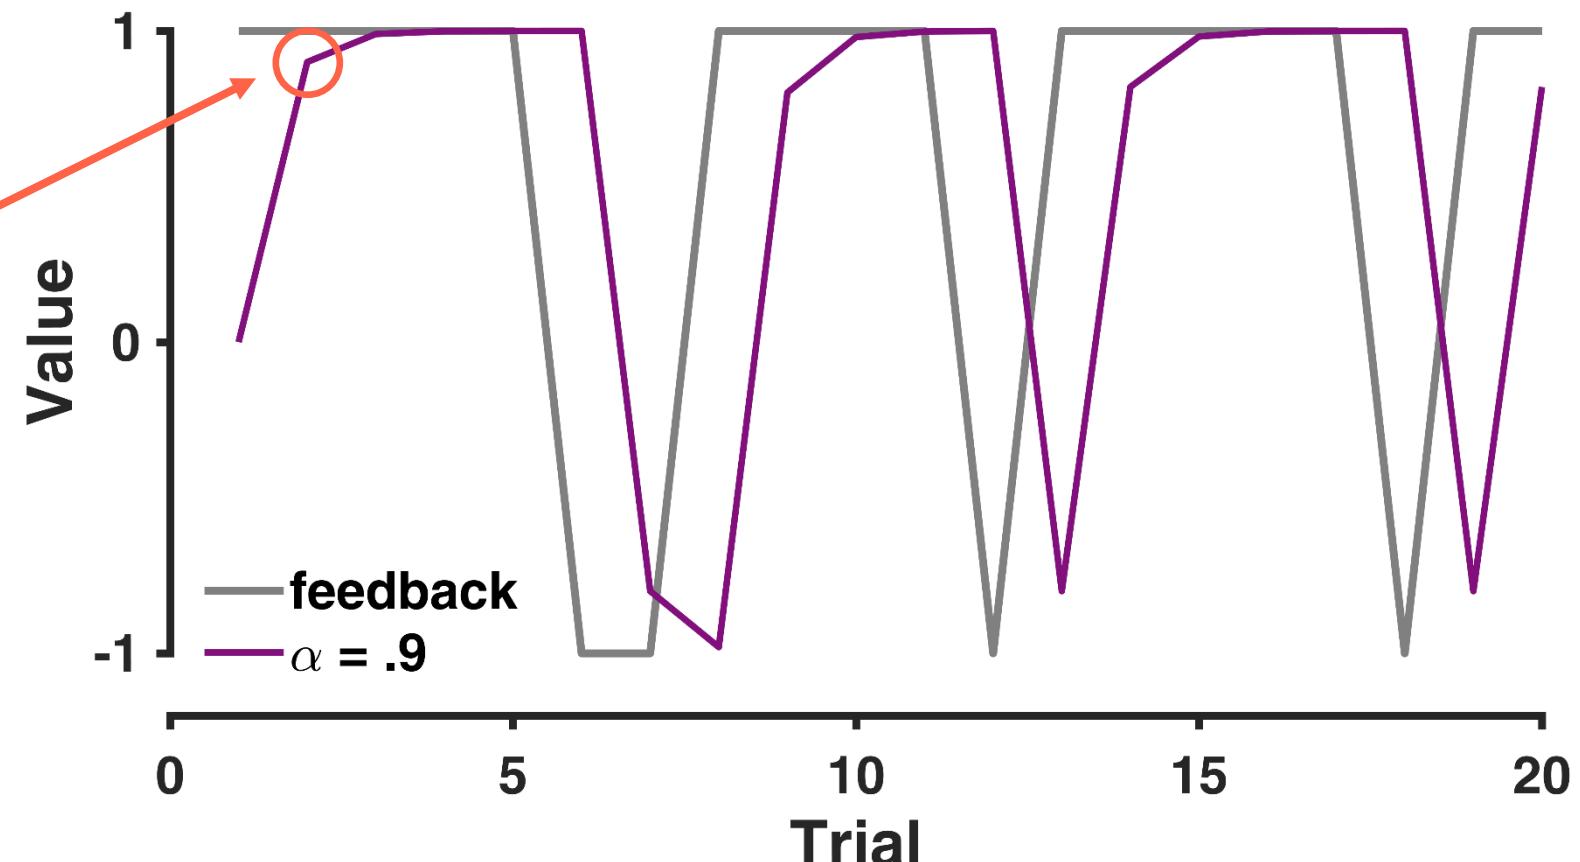
Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$



if  $\alpha = 0.9$

$$V_1 = 0$$

$$\begin{aligned} V_2 &= V_1 + 0.9 * (1 - V_1) \\ &= 0 + 0.9 * (1 - 0) \\ &= 0.9 \end{aligned}$$



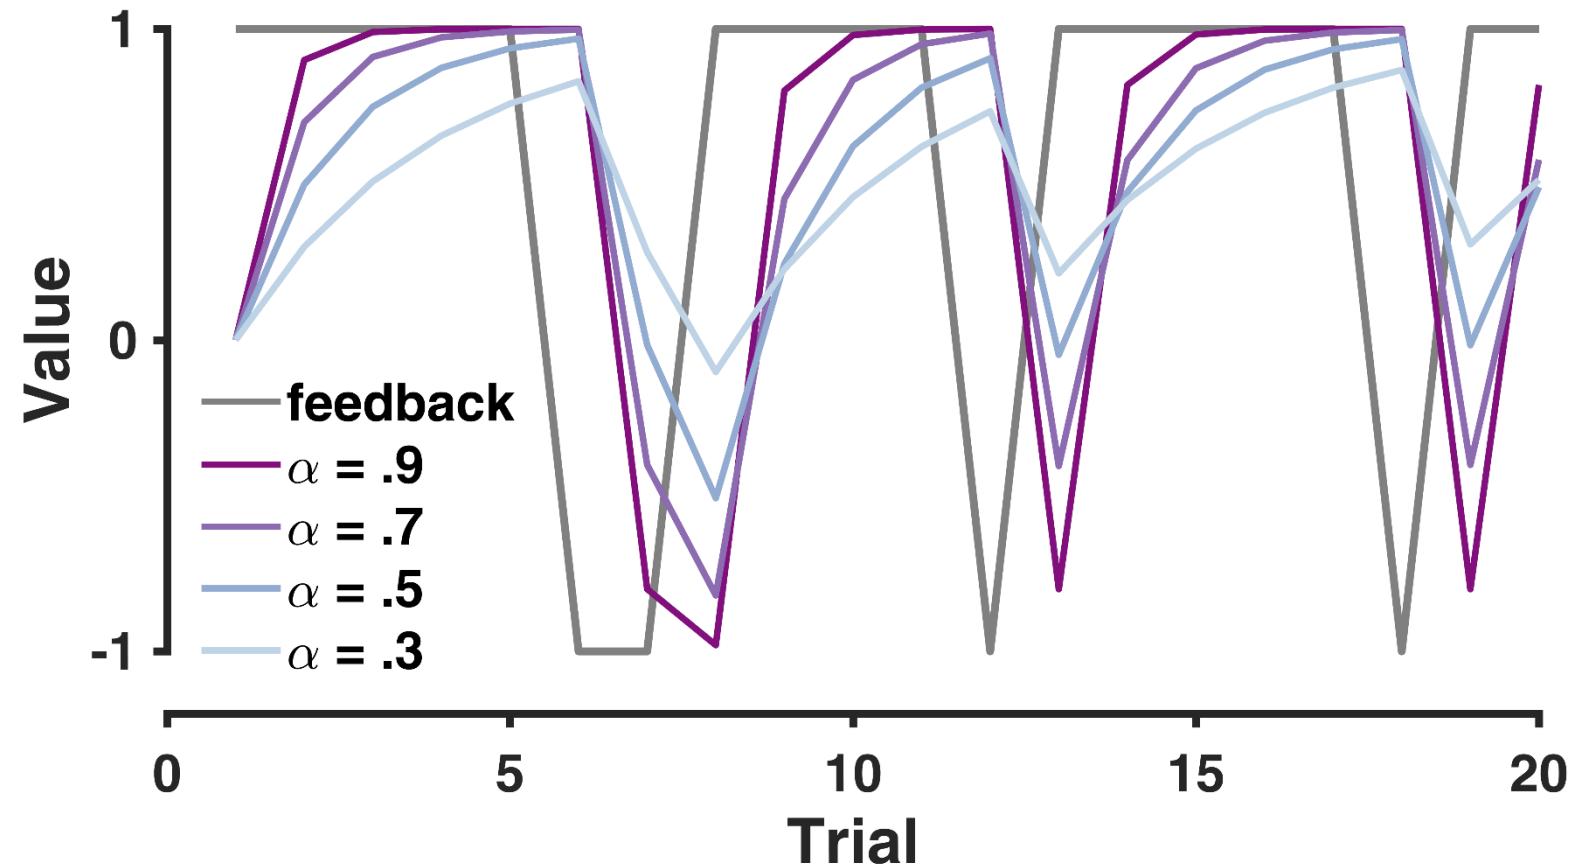
reward contingency – 80:20

# Understand the learning rate

cognitive model
statistics
computing

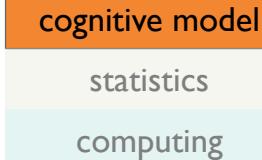
Value update:  $V_t = V_{t-1} + \alpha * PE_{t-1}$

Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$



reward contingency – 80:20

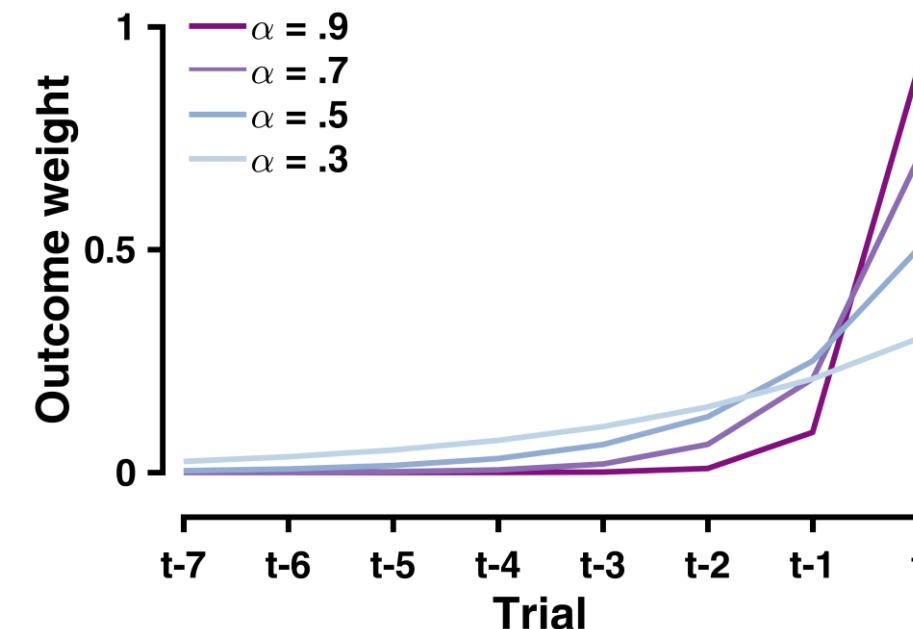
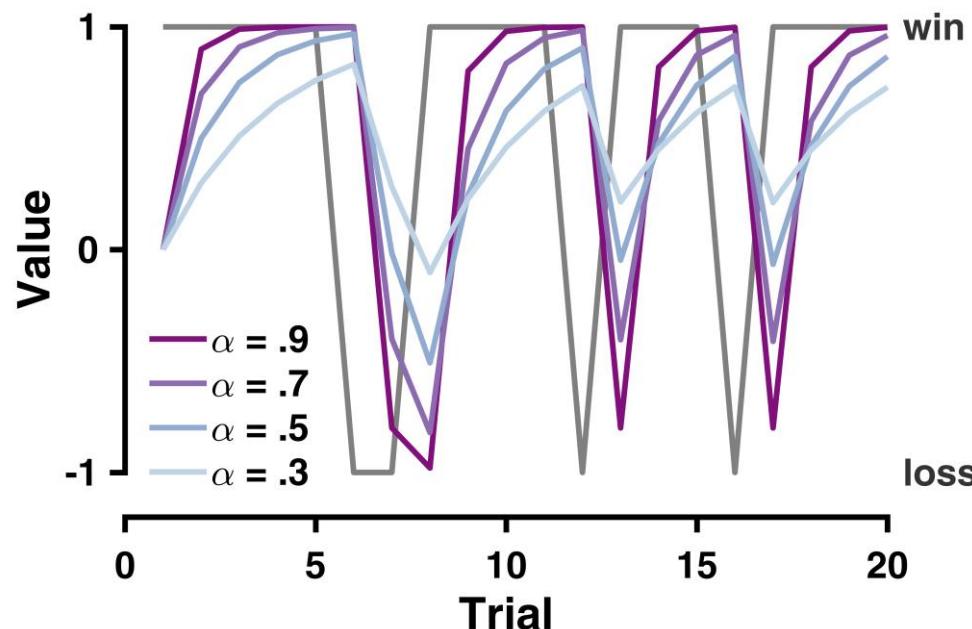
# Understand the learning rate



Value update:  $V_t = V_{t-1} + \alpha * PE_{t-1}$

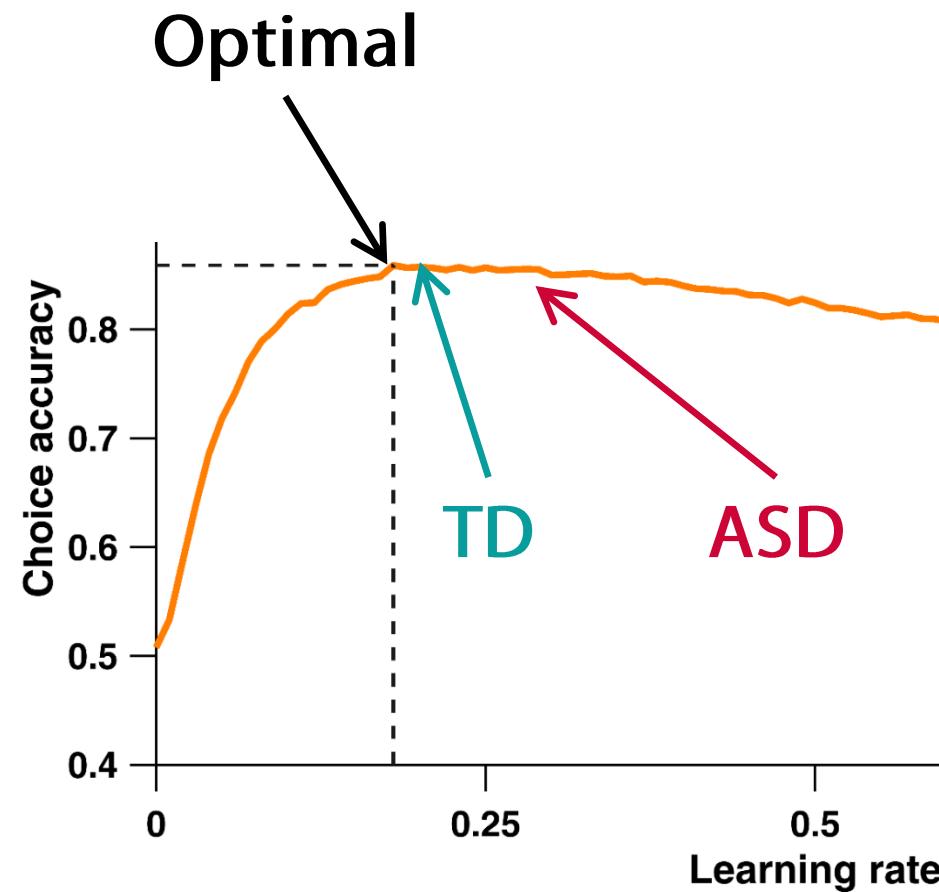
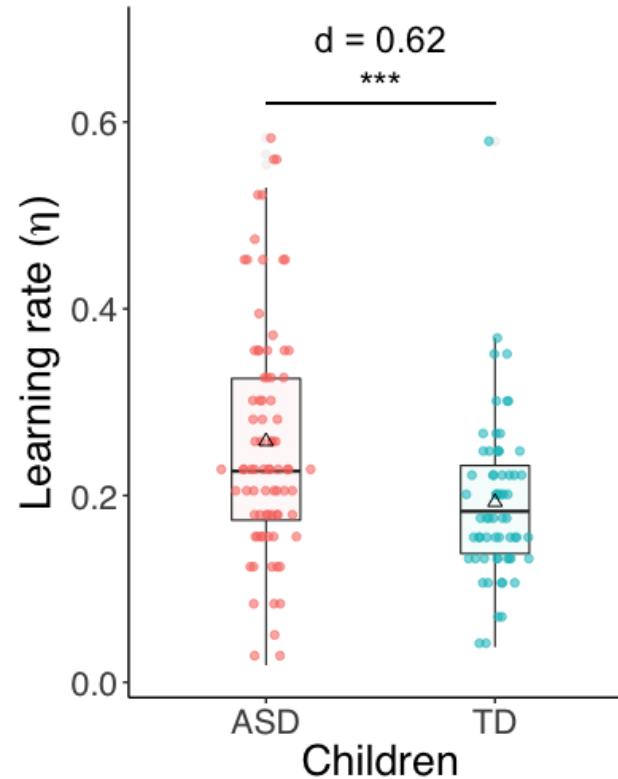
Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$

$$\begin{aligned}
 V_t &= (1 - \alpha) V_{t-1} + \alpha R_{t-1} \\
 &= (1 - \alpha) (V_{t-2} + \alpha (R_{t-2} - V_{t-2})) + \alpha R_{t-1} \\
 &= (1 - \alpha)^{t-1} V_1 + \sum_{i=1}^{t-1} (1 - \alpha)^{t-i-1} \alpha R_i
 \end{aligned}$$

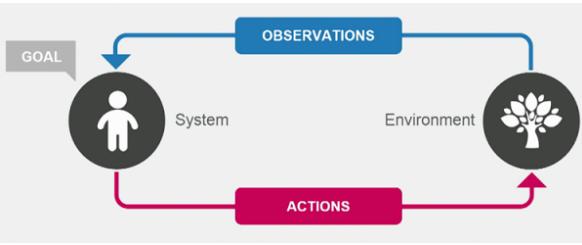


# Optimal learning rate?

cognitive model  
statistics  
computing



# Rescorla-Wagner Value Update



**Value update:**

$$V_{t+1} = V_t + \alpha * PE_t$$

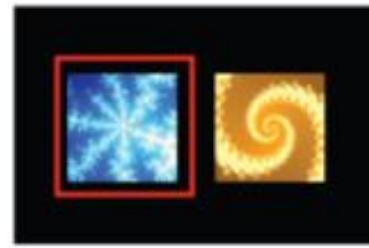
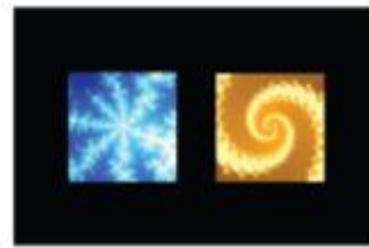
**Prediction error:**

$$PE_t = R_t - V_t$$

**choice rule:**

greedy /  $\epsilon$ -greedy / softmax

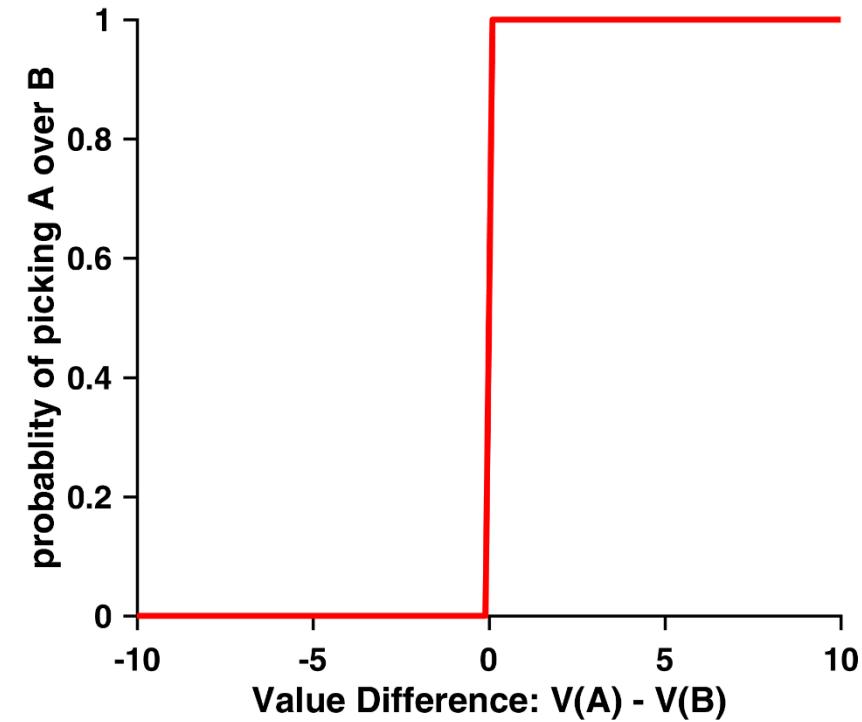
## Choice rule: greedy



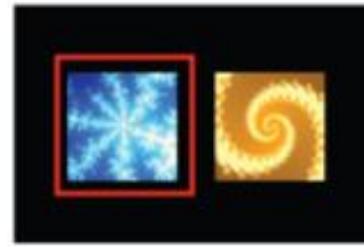
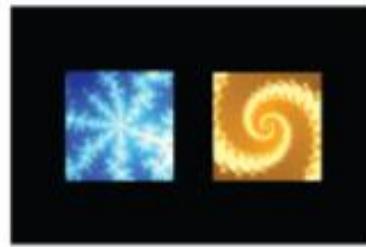
$$V(\text{orange spiral})$$

$$V(\text{blue snowflake})$$

$$p(C = a) = \begin{cases} 1, & V(a) > V(b) \\ 0, & V(a) < V(b) \end{cases}$$



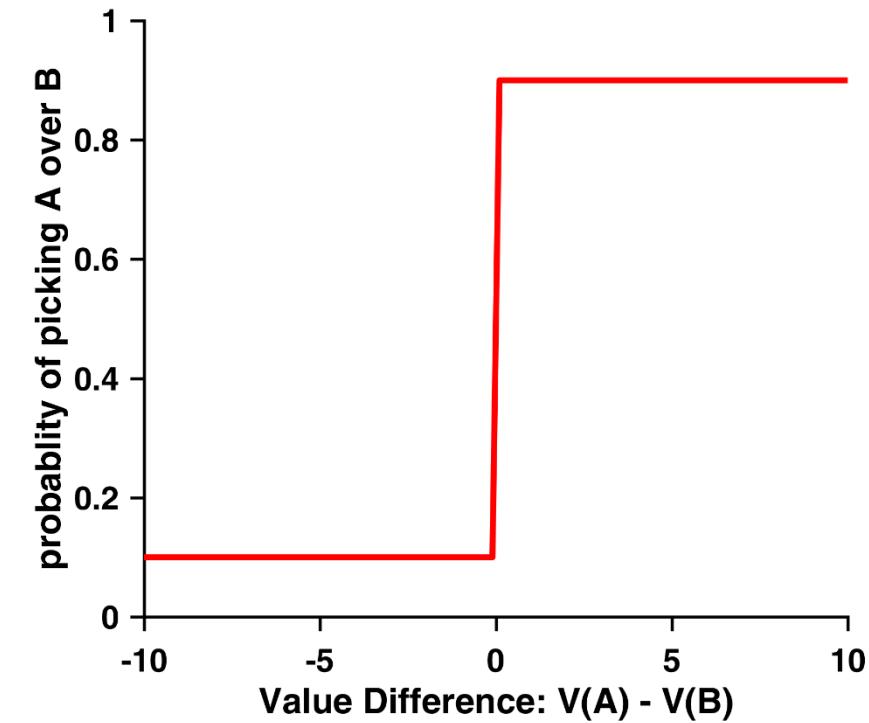
# Choice rule: $\epsilon$ -greedy



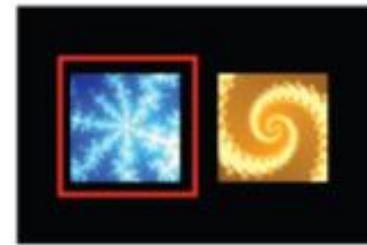
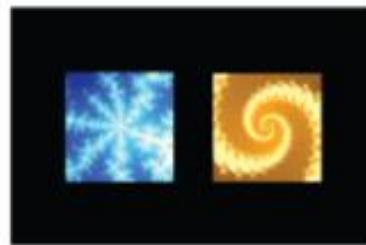
$$V(\text{swirl})$$

$$V(\text{snowflake})$$

$$p(C = a) = \begin{cases} 1 - \epsilon, & V(a) > V(b) \\ \epsilon, & V(a) < V(b) \end{cases}$$



# Choice rule: softmax

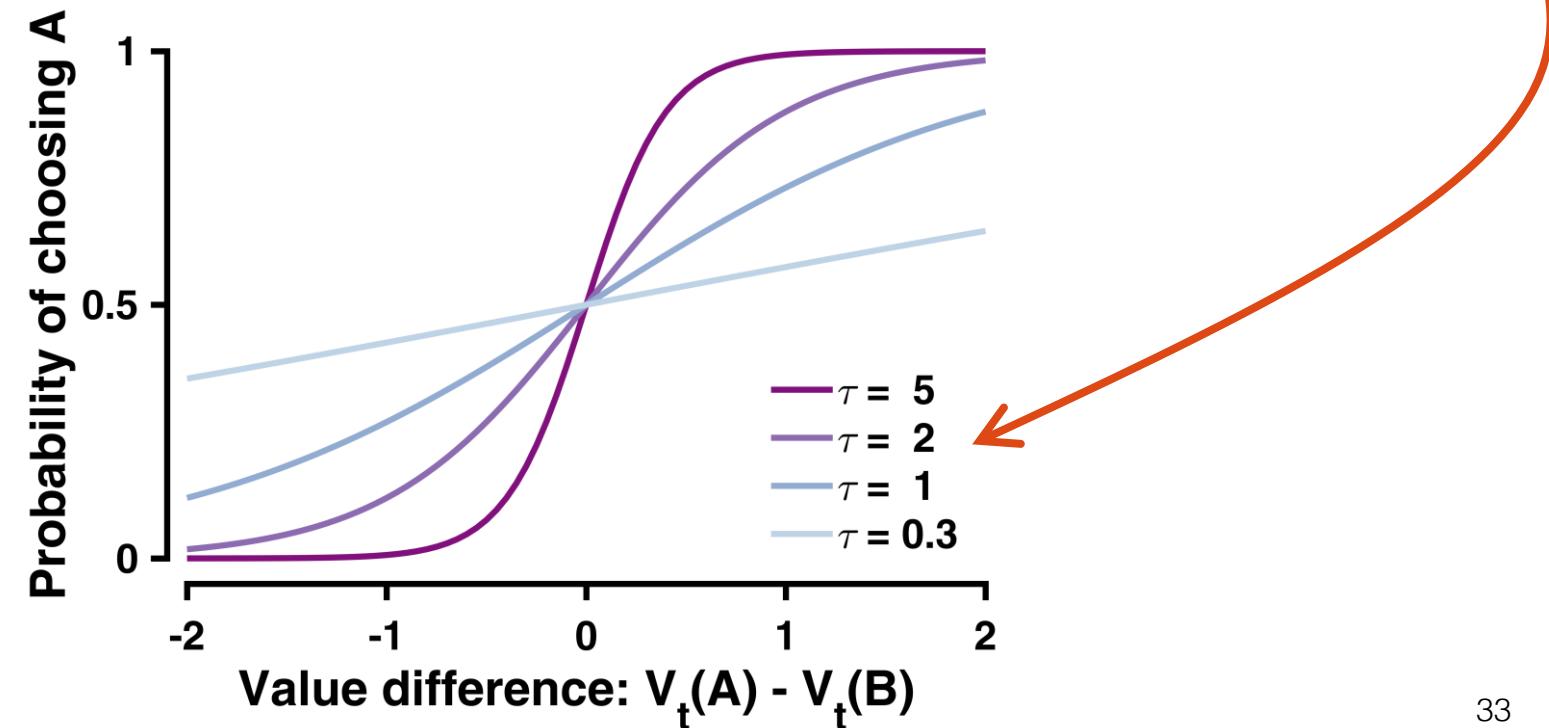


$$V(\text{orange})$$

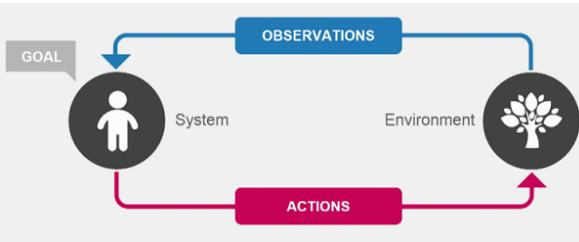
$$V(\text{blue})$$

$$p_t(A) = \frac{e^{\tau * V_t(A)}}{e^{\tau * V_t(A)} + e^{\tau * V_t(B)}}$$

$$= \frac{1}{1 + e^{-\tau * (V_t(A) - V_t(B))}}$$



# Rescorla-Wagner Value Update



**Value update:**

$$V_{t+1} = V_t + \alpha * PE_t$$

**Prediction error:**

$$PE_t = R_t - V_t$$

**choice rule (sigmoid /softmax):**

$$p(C=a) = \frac{1}{1+e^{\tau*(v(b)-v(a))}}$$

$\alpha$  - learning rate

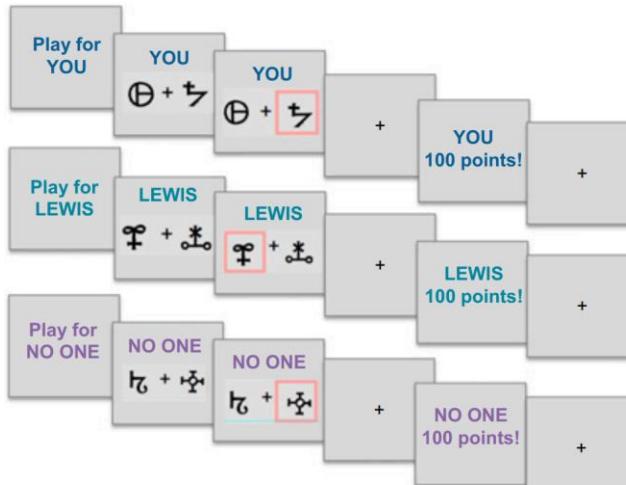
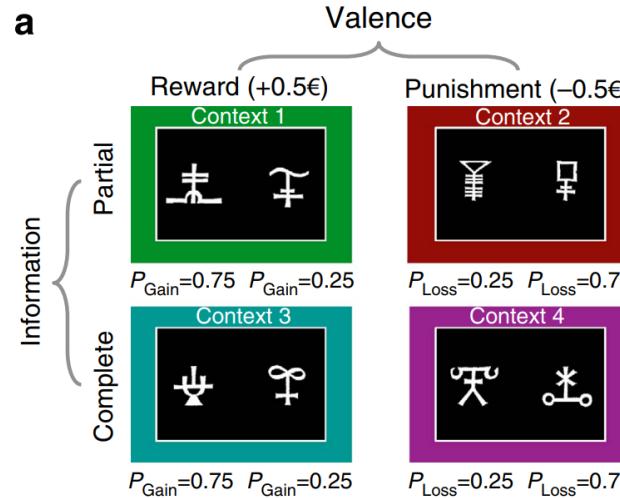
PE - reward prediction error

V - value

R - reward

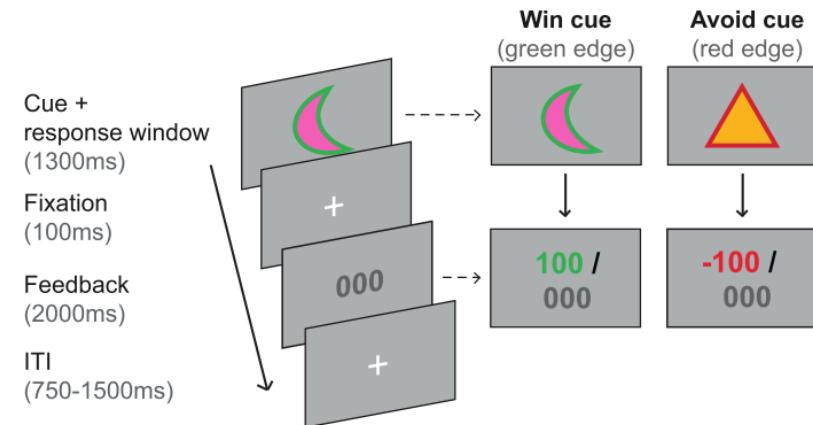
$\tau$  - softmax temperature

# Generalizing RL framework

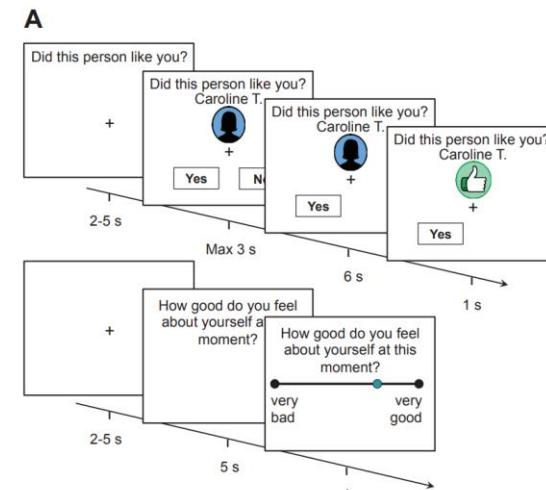


Lockwood et al. (2016)

A. Trial details



Swart et al. (2017)



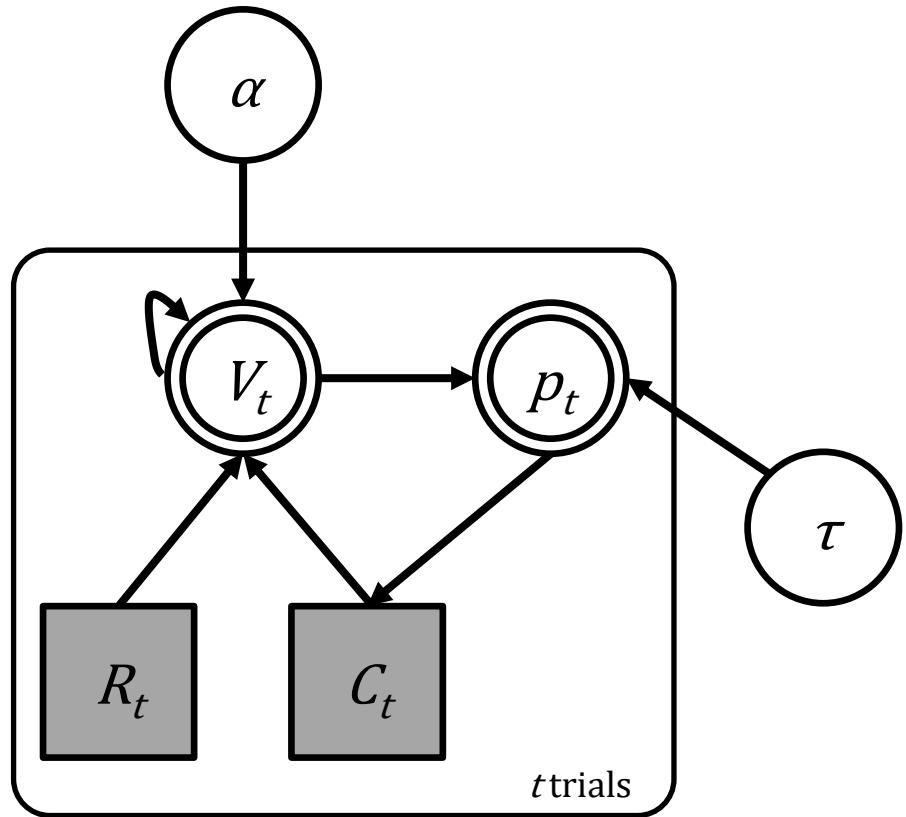
Will et al. (2017)

**B**

	Did this person like you?	How good do you feel about yourself at this moment?
Blue profile	85%	15%
Yellow profile	70%	30%
Purple profile	30%	70%
Orange profile	15%	85%

# RL – Implementation

cognitive model  
statistics  
computing



$$\alpha \sim Uniform(0, 1)$$

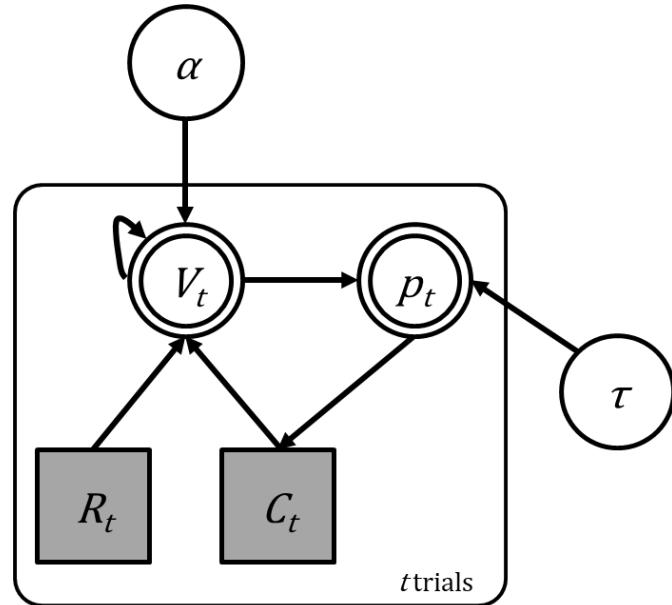
$$\tau \sim Uniform(0, 3)$$

$$p_t(C = A) = \frac{1}{1 + e^{\tau(V_t(B) - V_t(A))}}$$

$$V_{t+1}^c = V_t^C + \alpha (R_t - V_t^C)$$

# RL – Implementation

cognitive model  
statistics  
computing



$$\alpha \sim Uniform(0, 1)$$

$$\tau \sim Uniform(0, 3)$$

$$p_t(C = A) = \frac{1}{1 + e^{\tau(V_t(B) - V_t(A))}}$$

$$V_{t+1}^c = V_t^c + \alpha (R_t - V_t^c)$$

```

transformed data {
  vector[2] initV;
  initV = rep_vector(0.0, 2);
}

model {
  vector[2] v[nTrials+1];
  real pe[nTrials];

  v[1] = initV;

  for (t in 1:nTrials) {
    choice[t] ~ categorical_logit( tau * v[t] );

    pe[t] = reward[t] - v[t,choice[t]];

    v[t+1] = v[t];
    v[t+1, choice[t]] = v[t, choice[t]] + lr * pe[t];
  }
}

```

# RL - Implementation

cognitive model  
statistics  
computing

```
model {  
    vector[2] v[nTrials+1];  
    real pe[nTrials];  
  
    v[1] = initV;  
  
    for (t in 1:nTrials) {  
        choice[t] ~ categorical_logit( tau * v[t] );  
        pe[t] = reward[t] - v[t,choice[t]];  
  
        v[t+1] = v[t];  
        v[t+1, choice[t]] = v[t, choice[t]] + lr * pe[t];  
    }  
}
```

```
model {  
    vector[2] v;  
    real pe;  
  
    v = initV;  
  
    for (t in 1:nTrials) {  
        choice[t] ~ categorical_logit( tau * v );  
        pe = reward[t] - v[choice[t]];  
  
        v[choice[t]] = v[choice[t]] + lr * pe;  
    }  
}
```

# RL – Fitting with Stan

cognitive model  
statistics  
computing

.../06.reinforcement\_learning/\_scripts/reinforcement\_learning\_single\_parm\_main.R

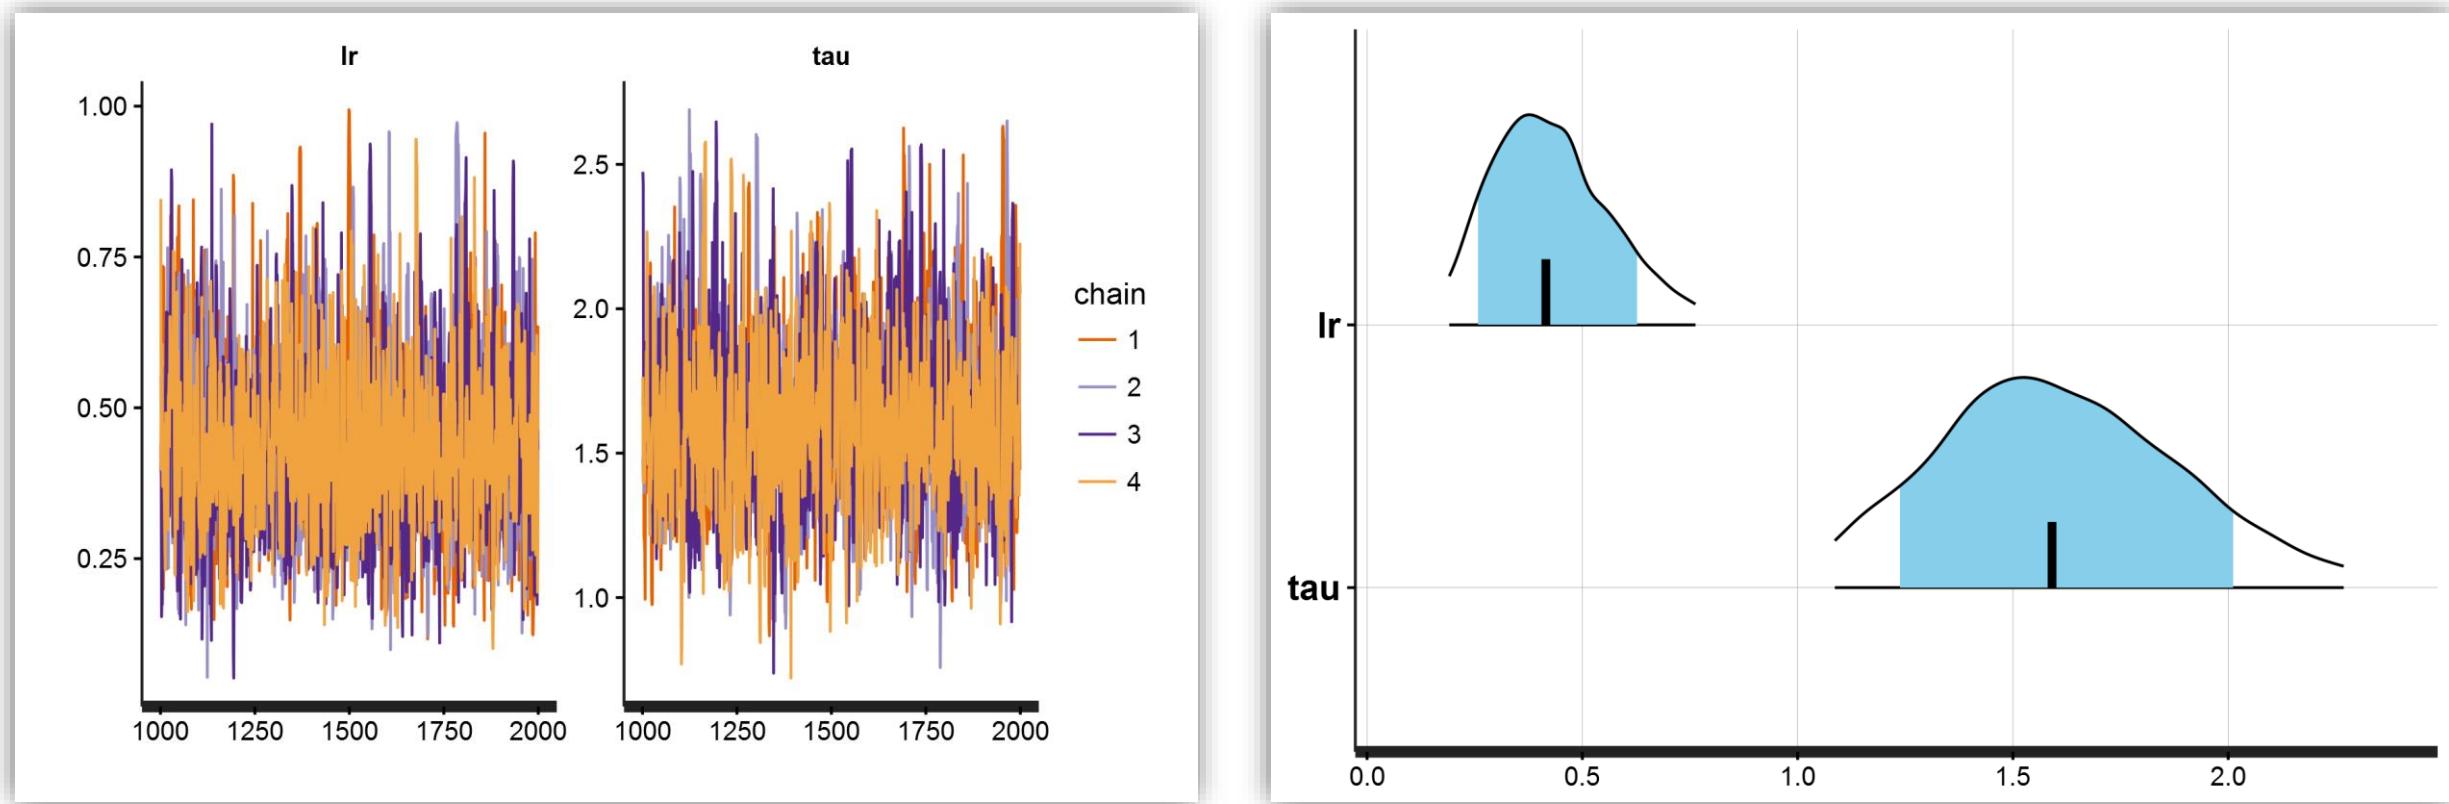
TASK: fit the model for single participants

```
> source('_scripts/reinforcement_learning_single_parm_main.R') # a function  
  
> fit_rl1 <- run_rl_sp(multiSubj = FALSE)
```

```
> load('_data/rl_sp_ss.RData')  
> head(rl_ss)  
     [,1] [,2]  
[1,]    2   -1  
[2,]    1    1  
[3,]    1    1  
[4,]    1    1  
[5,]    2   -1  
[6,]    1    1
```

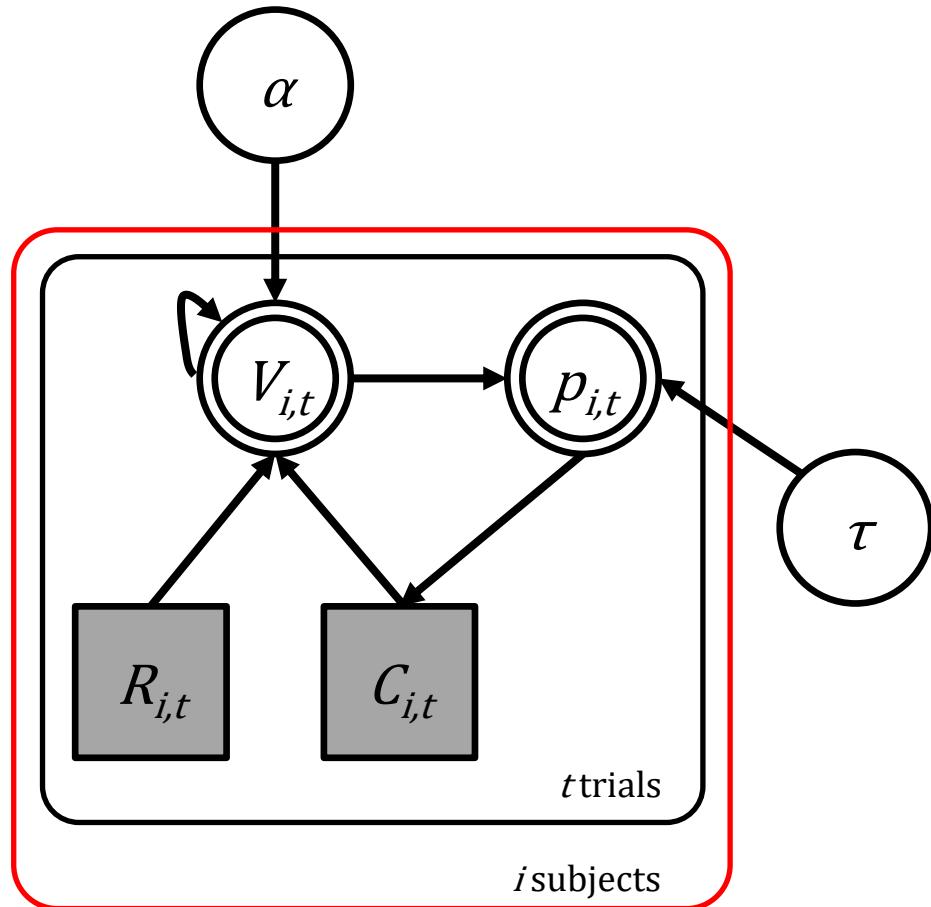
# RL - MCMC Output

cognitive model  
statistics  
computing



# Fitting Multiple Participants as ONE

cognitive model  
statistics  
computing



```
model {  
    for (s in 1:nSubjects) {  
        vector[2] v;  
        real pe;  
        v = initV;  
  
        for (t in 1:nTrials) {  
            choice[s,t] ~ categorical_logit( tau * v );  
            pe = reward[s,t] - v[choice[s,t]];  
            v[choice[s,t]] = v[choice[s,t]] + lr * pe;  
        }  
    }  
}
```

# Exercise X

cognitive model  
statistics  
computing

.../06.reinforcement\_learning/\_scripts/reinforcement\_learning\_single\_parm\_main.R

TASK:

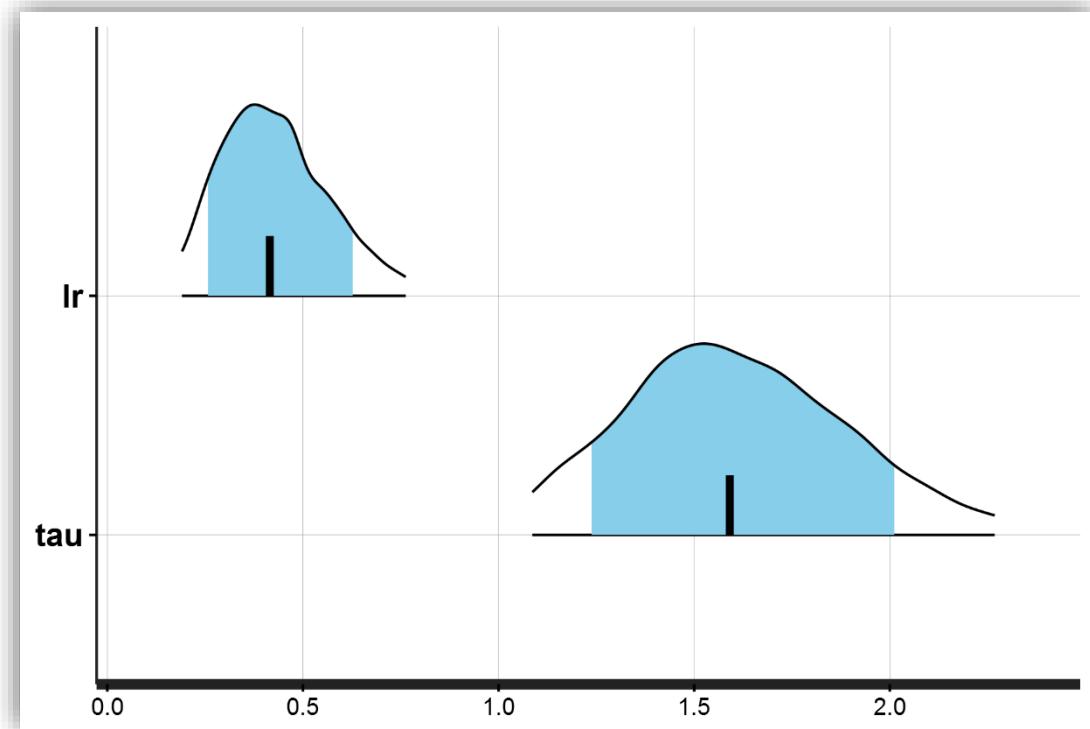
- (1) complete the model (Tip: the for-loop)
- (2) fit the model for multiple participants (assuming same parameters)

```
> source('_scripts/reinforcement_learning_single_parm_main.R')  
  
> fit_rl2 <- run_rl_sp(multiSubj = TRUE)
```

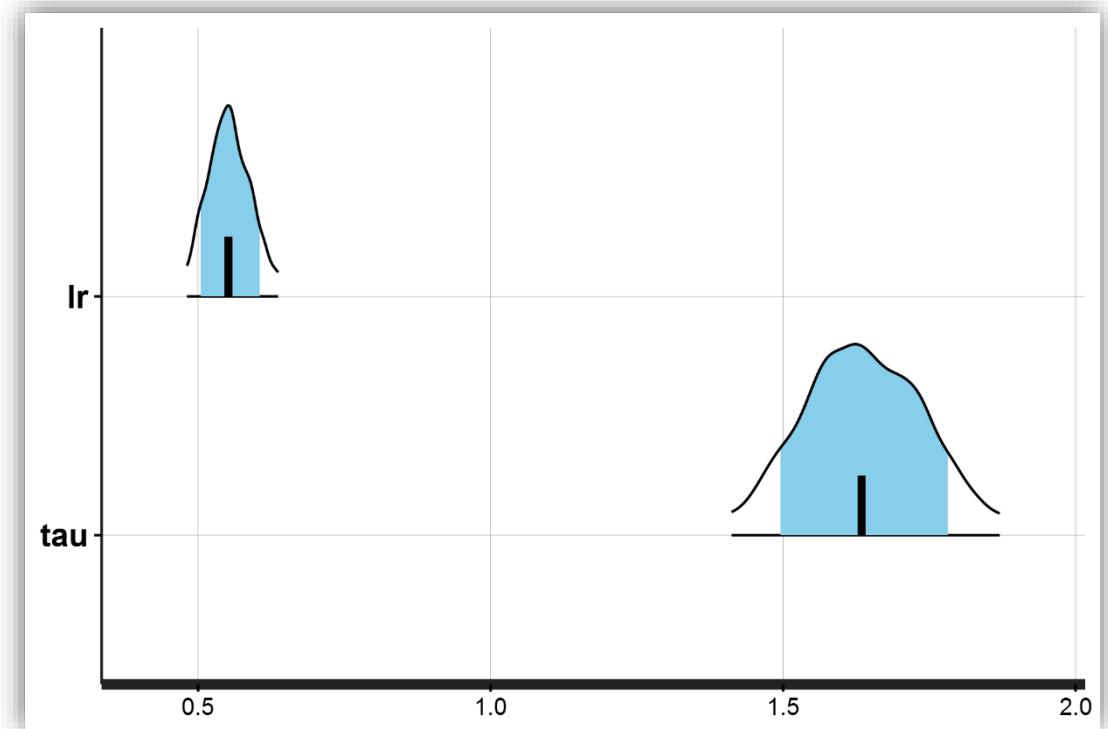
# Exercise X

cognitive model  
statistics  
computing

$N = 1$

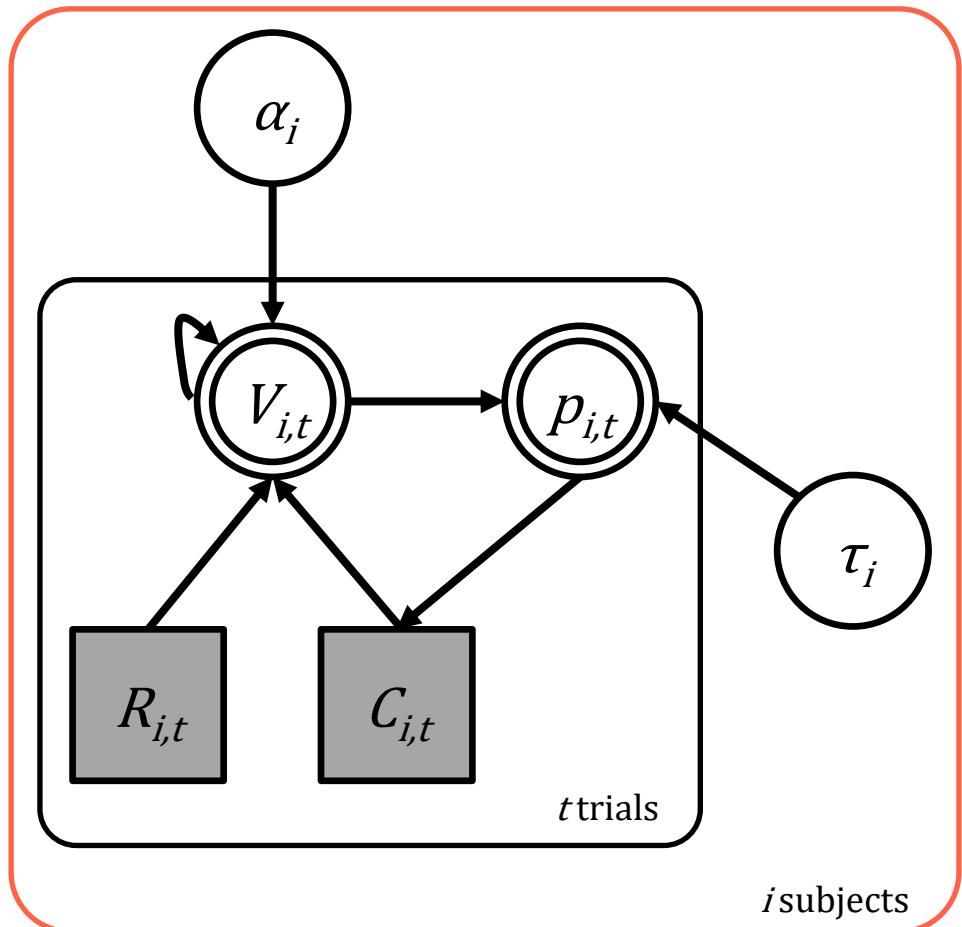


$N = 10$



# Fitting Multiple Participants Independently

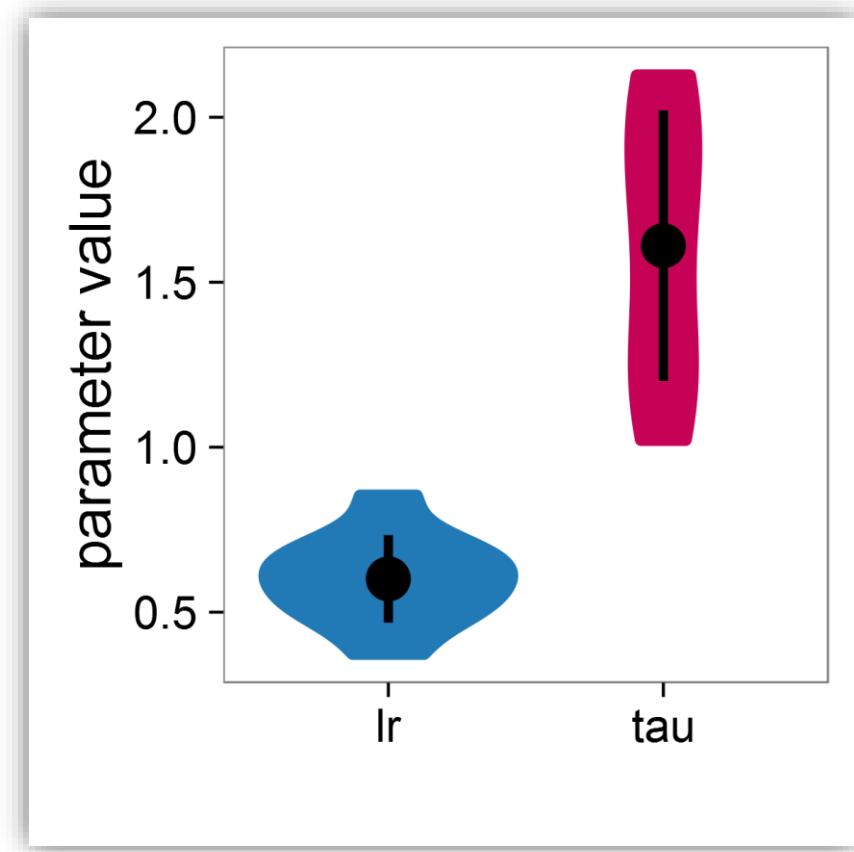
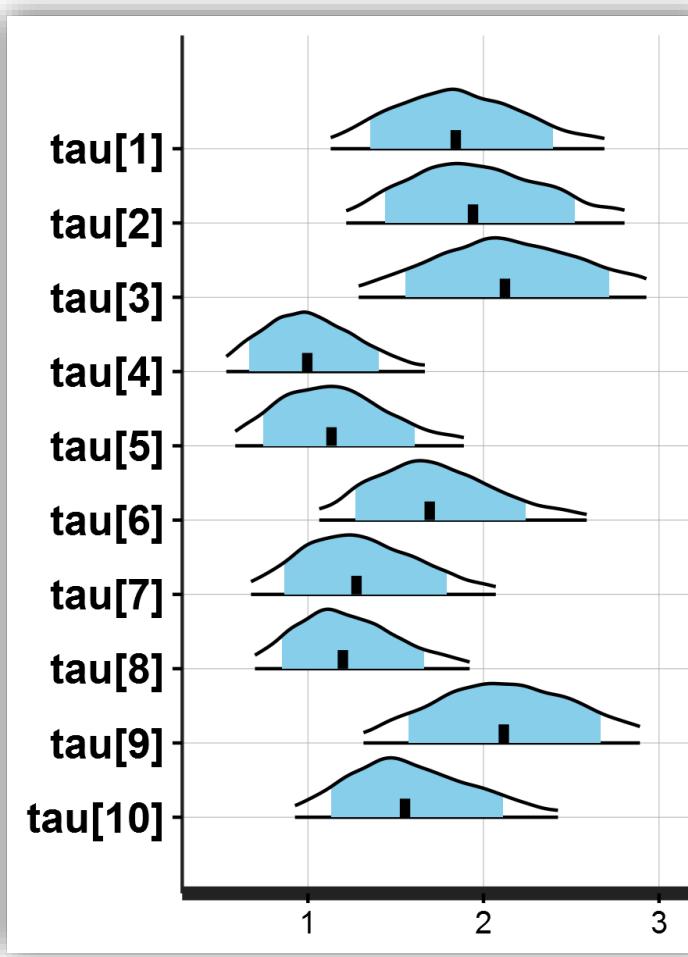
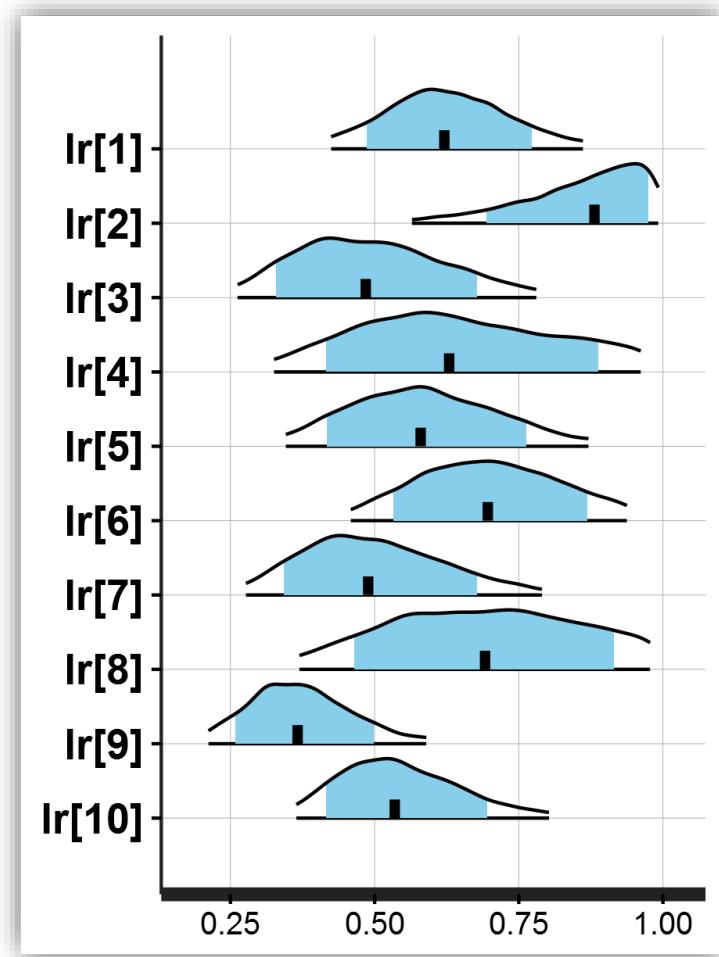
cognitive model  
statistics  
computing



```
model {  
    for (s in 1:nSubjects) {  
        vector[2] v;  
        real pe;  
        v = initV;  
  
        for (t in 1:nTrials) {  
            choice[s,t] ~ categorical_logit( tau[s] * v );  
            pe = reward[s,t] - v[choice[s,t]];  
            v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
        }  
    }  
}
```

# Individual Fitting

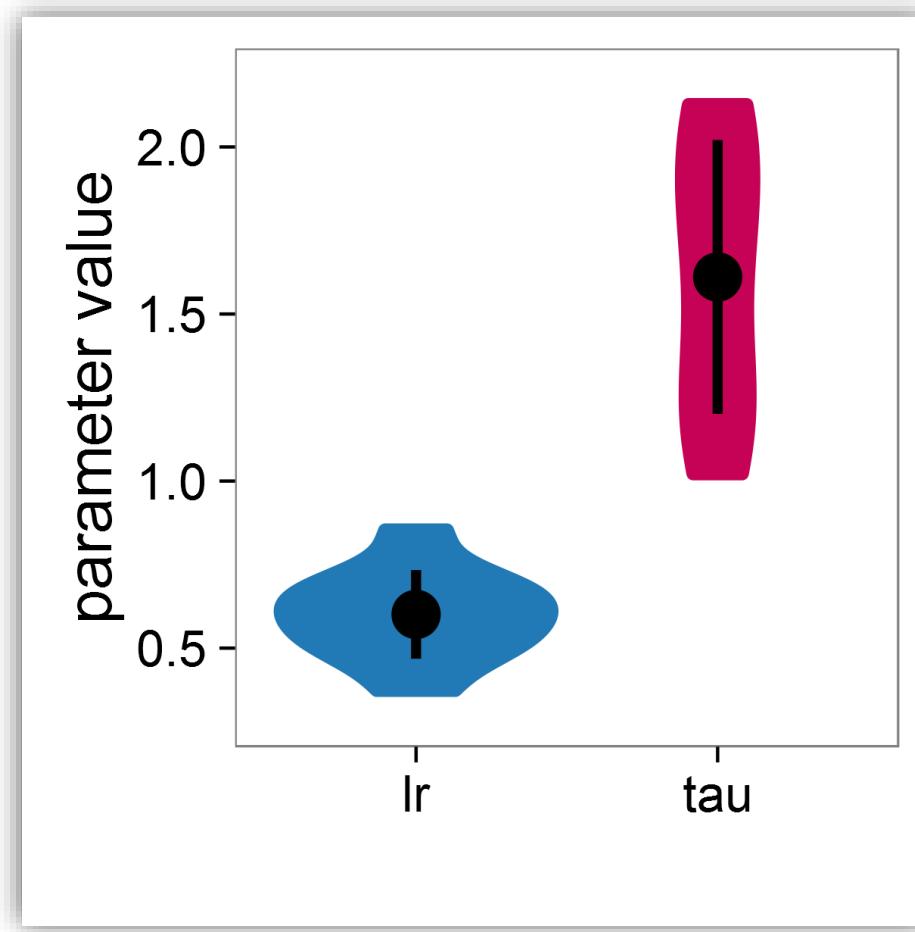
cognitive model  
statistics  
computing



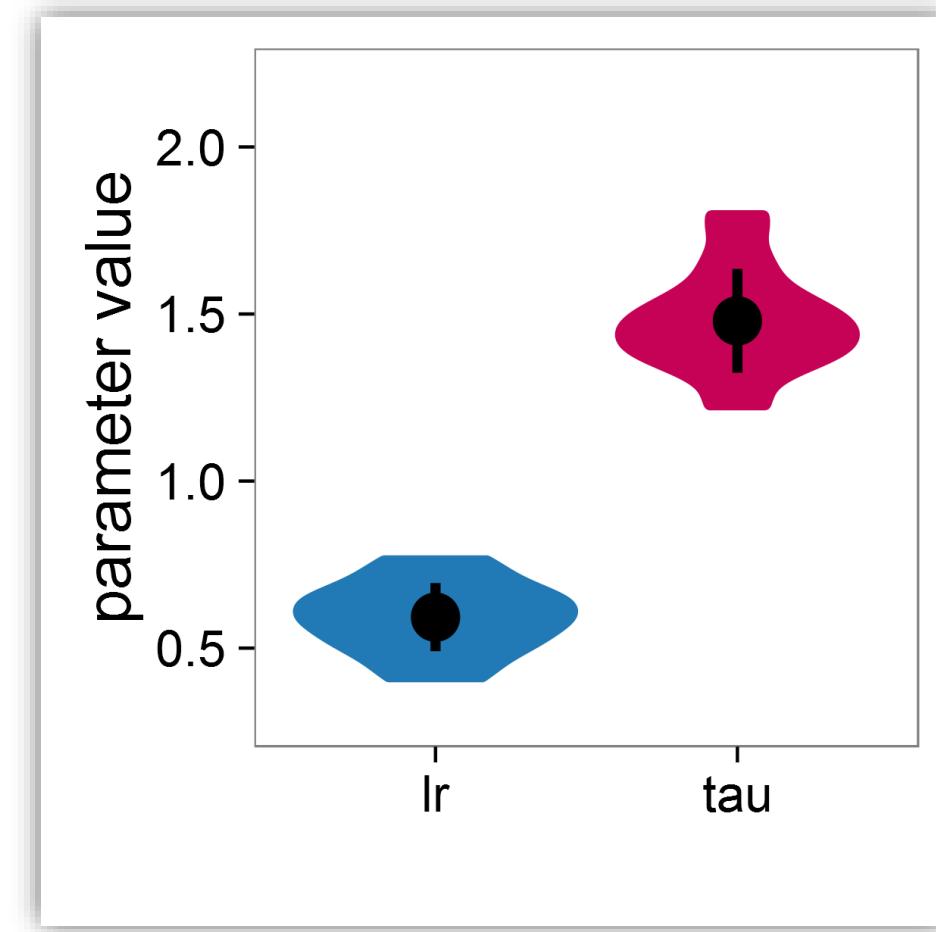
# Comparing with True Parameters

cognitive model  
statistics  
computing

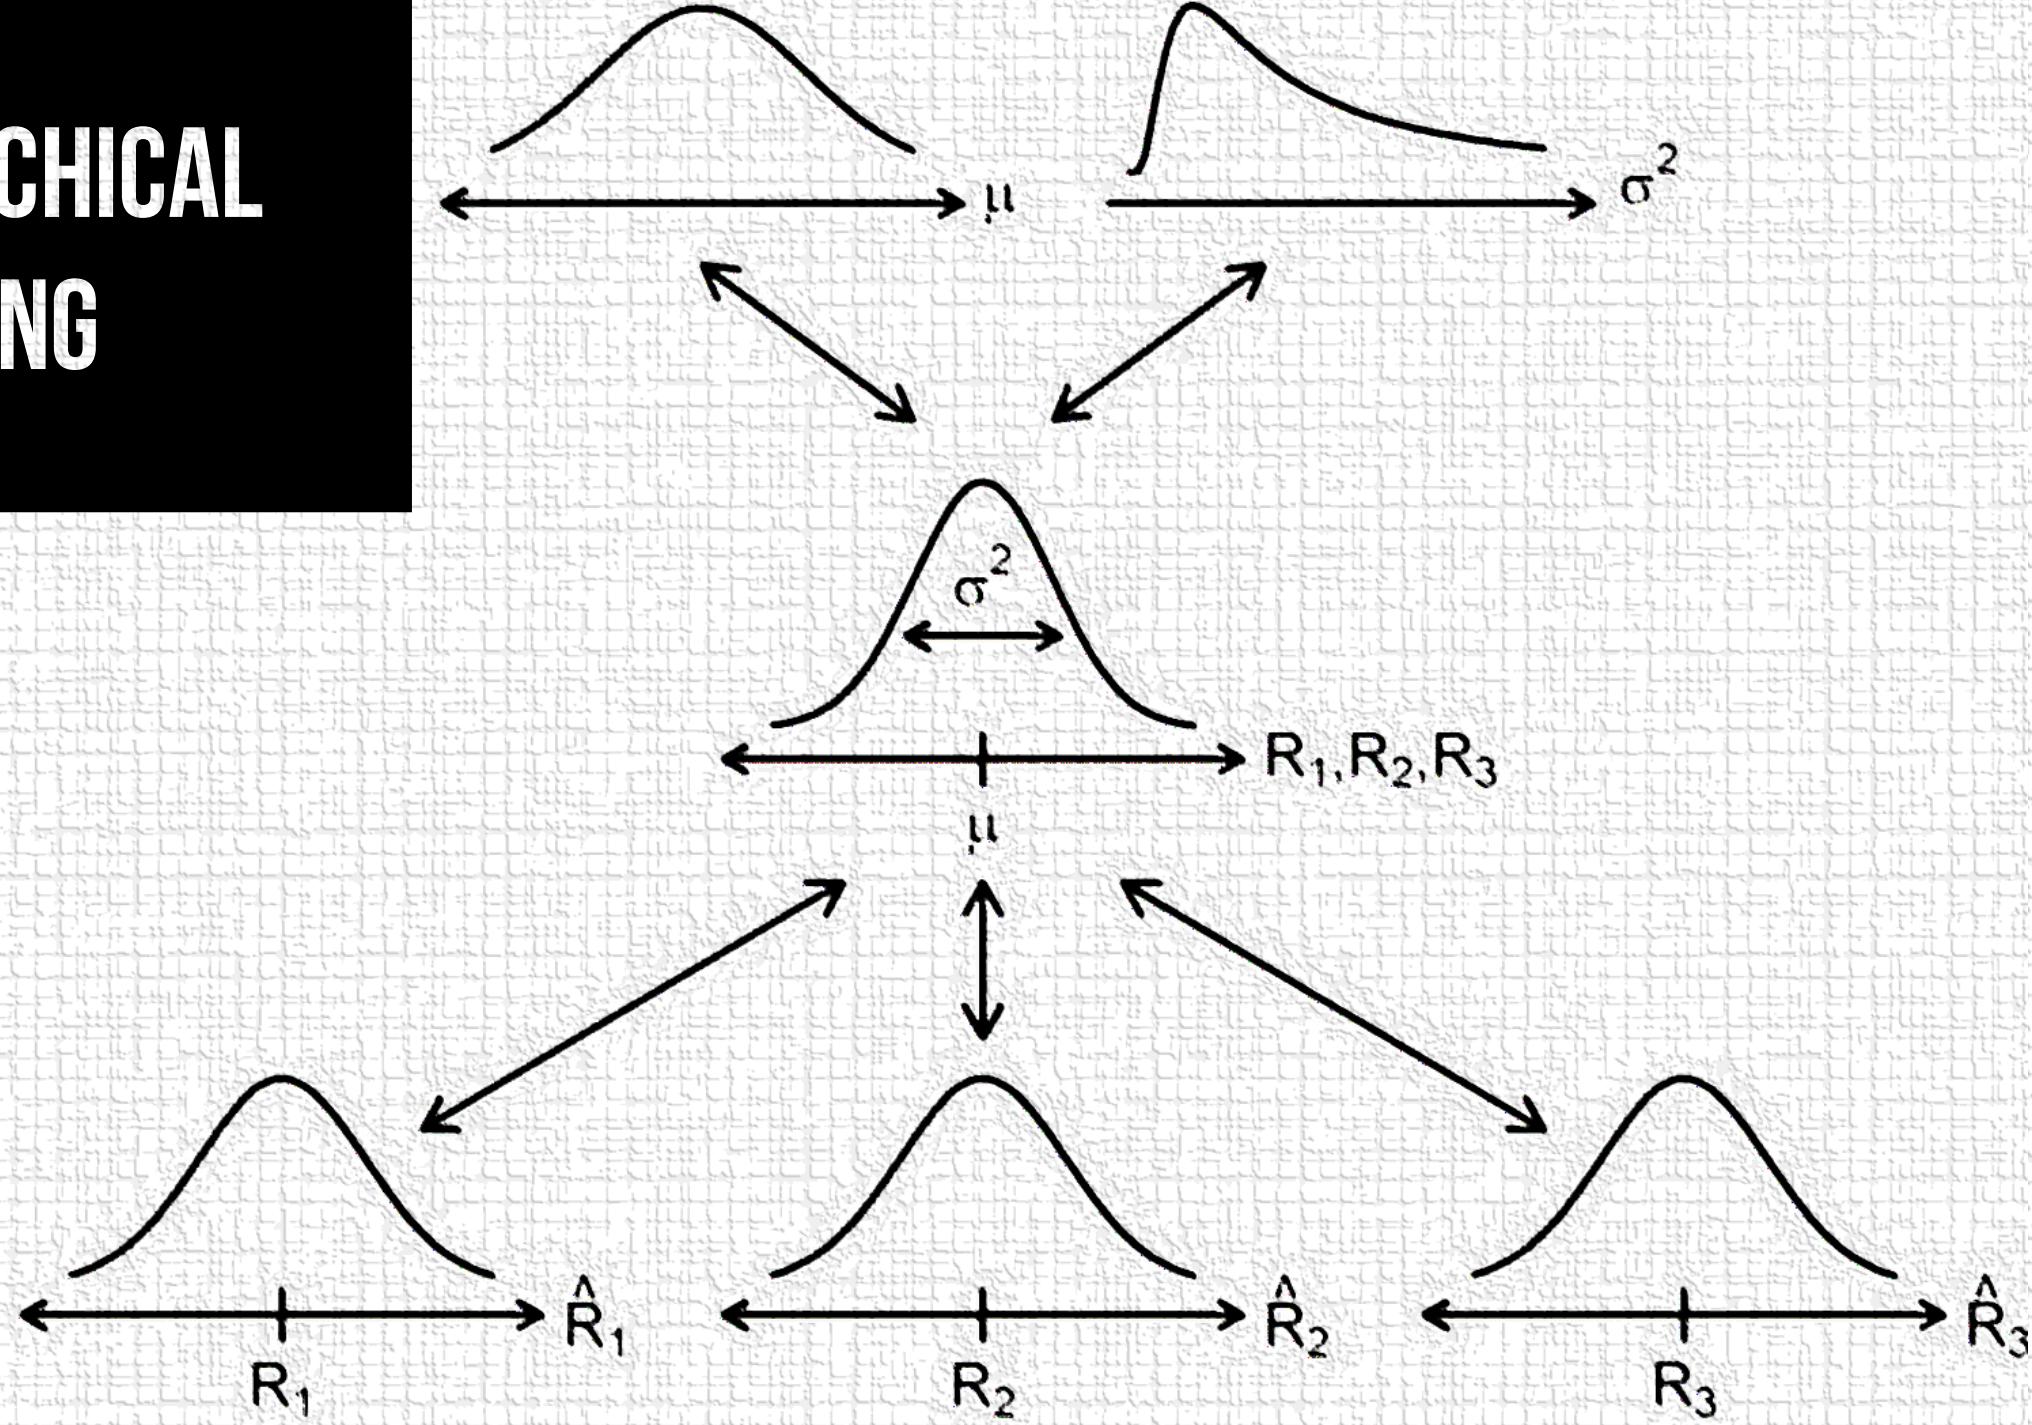
Posterior Means

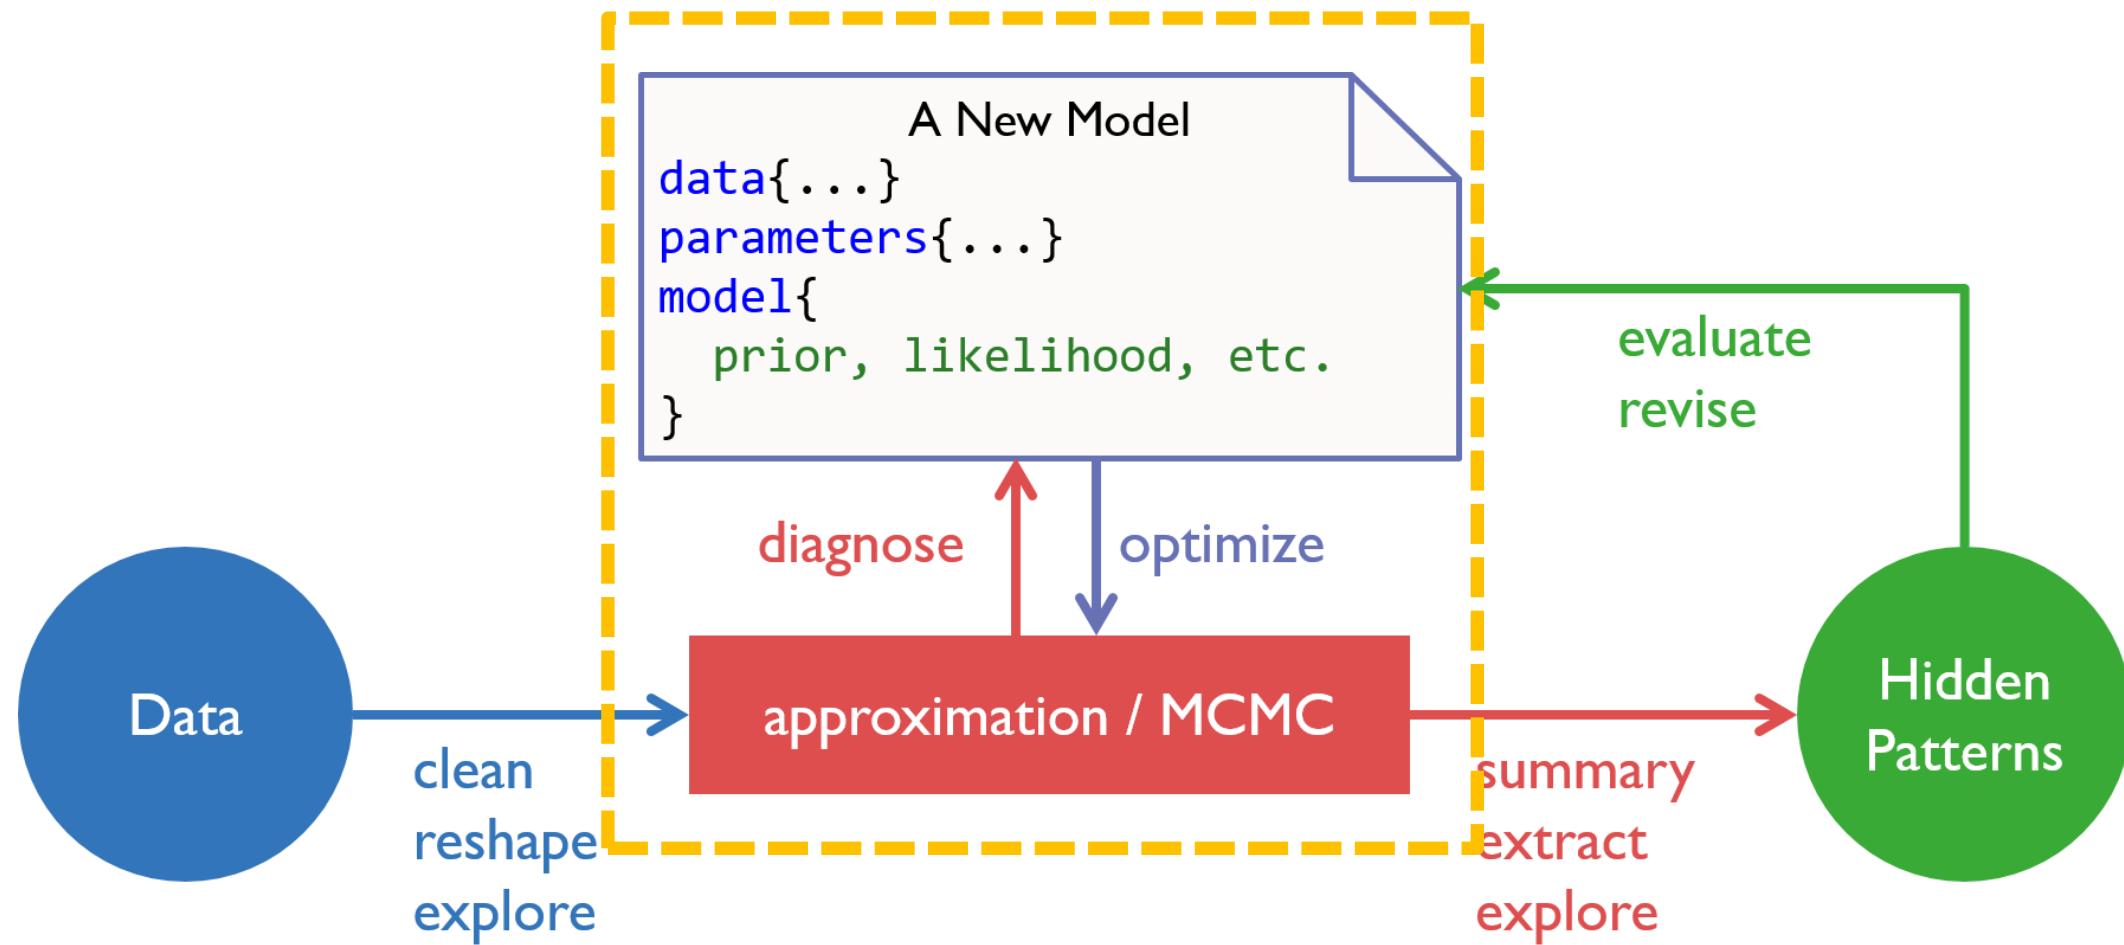


True Parameters

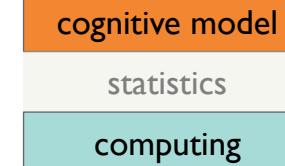


# HIERARCHICAL MODELING



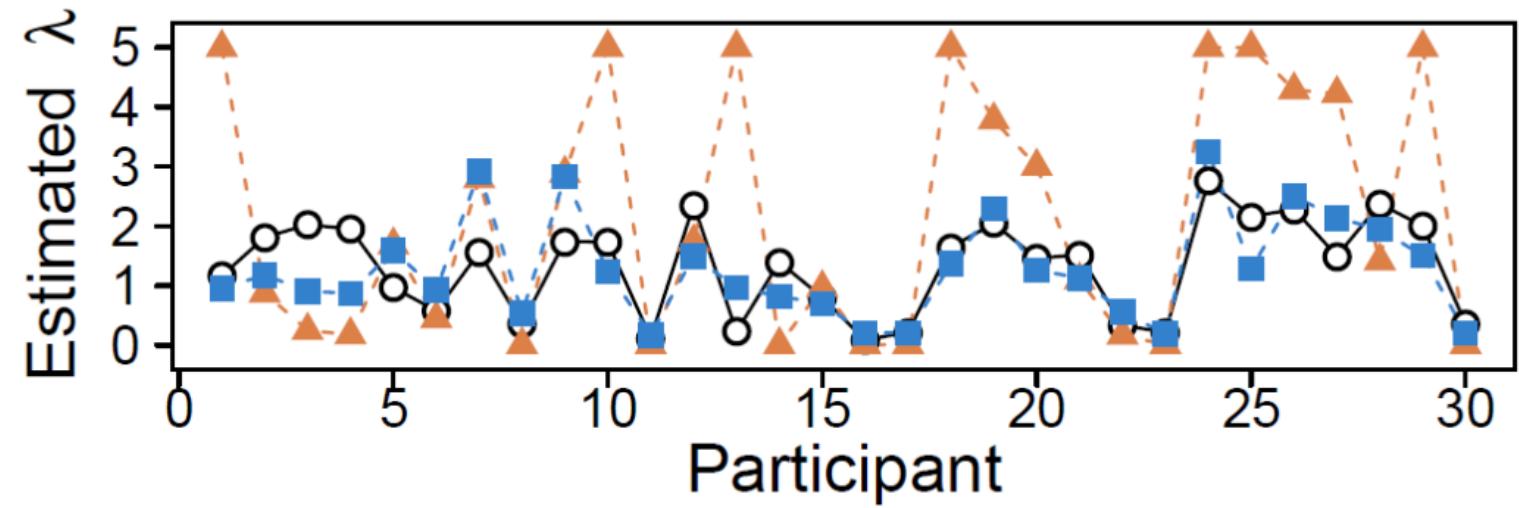


# Why Hierarchical Bayesian Cognitive Modeling?



## Simulation study

Hierarchical Bayesian ■  
Maximum likelihood ▲  
Actual values ○



# Why Hierarchical Bayesian Cognitive Modeling?

cognitive model  
statistics  
computing

## Fixed effects

- all subjects are fitted with the same set of parameters
- worse model fit than “random effects”

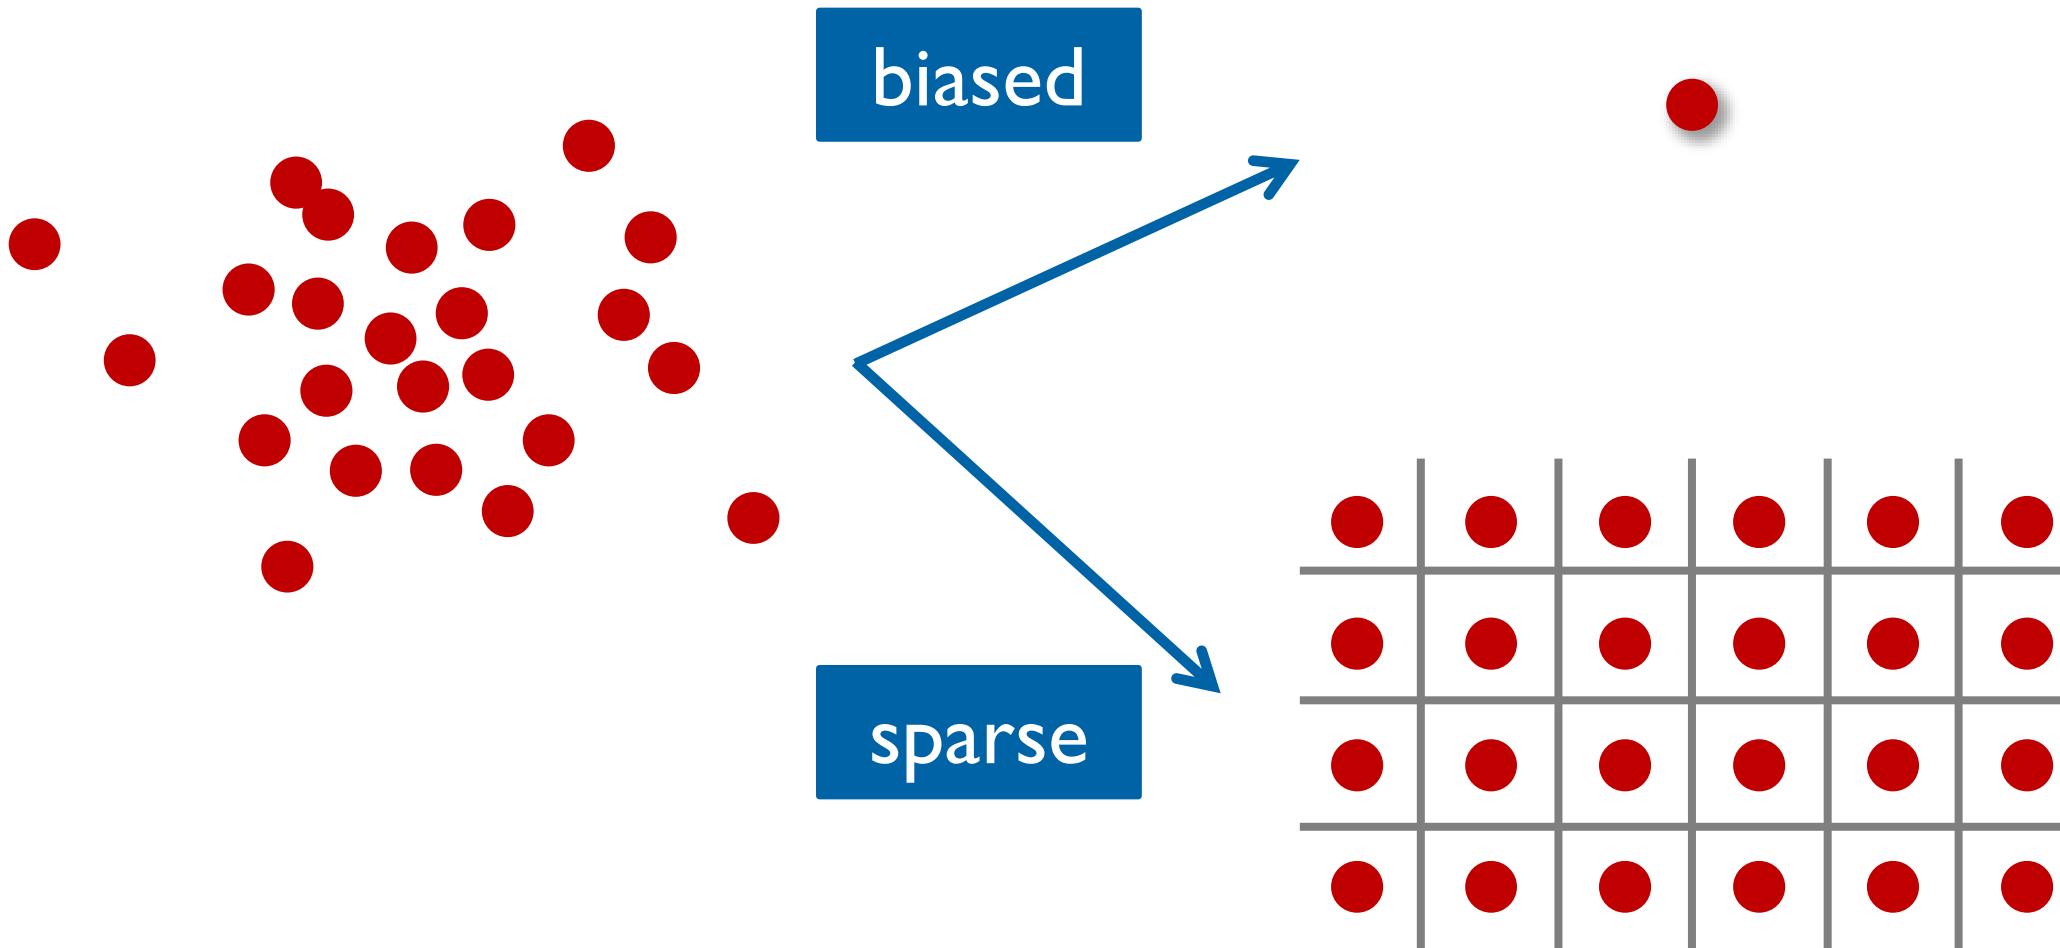
## Random effects

- each subject is fitted independently of the others
- best model fit for each subject
- parameter estimates can be noisy

Adapted from Jan Gläscher's workshop

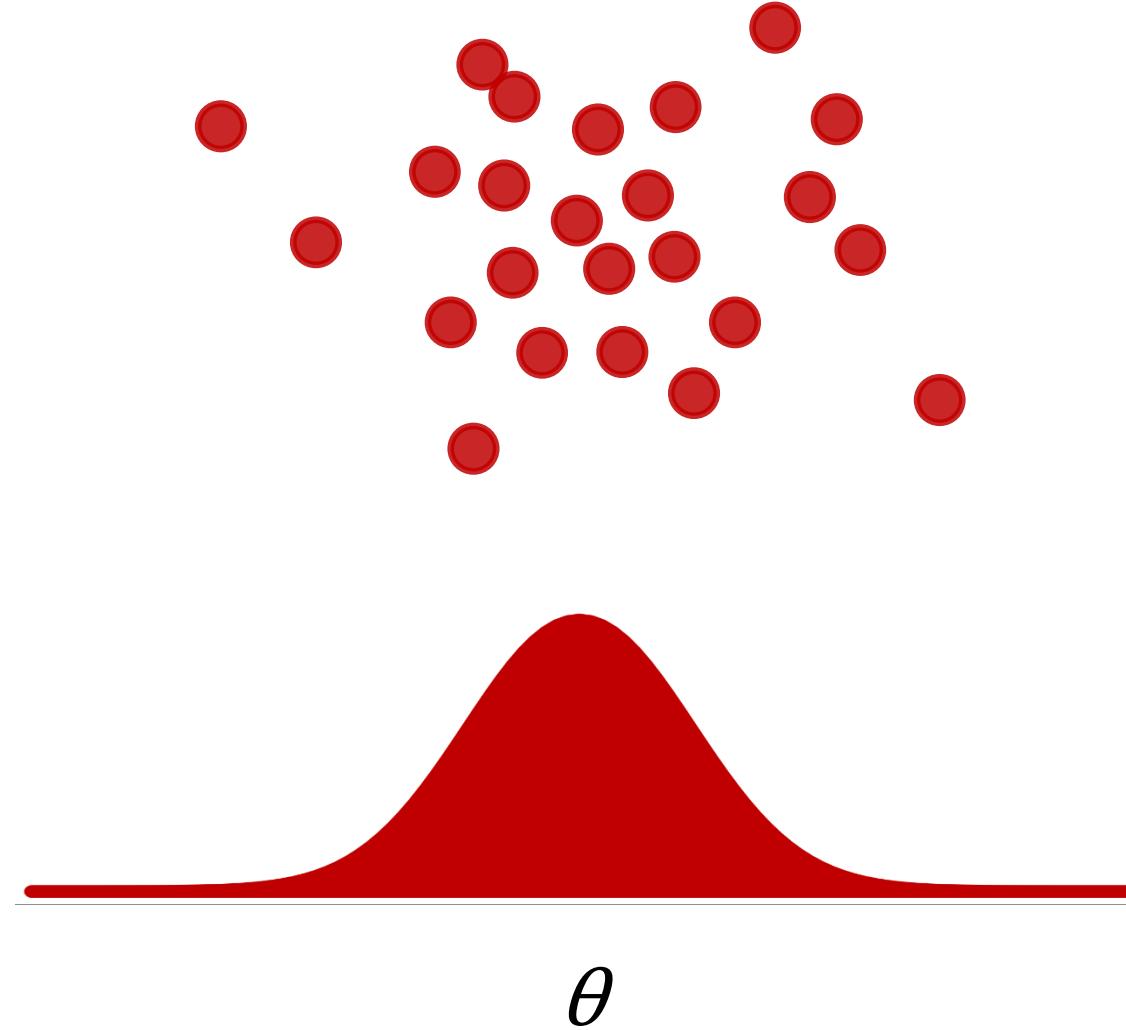
# Fitting Multiple Participants

cognitive model  
statistics  
computing



# Fitting Multiple Participants

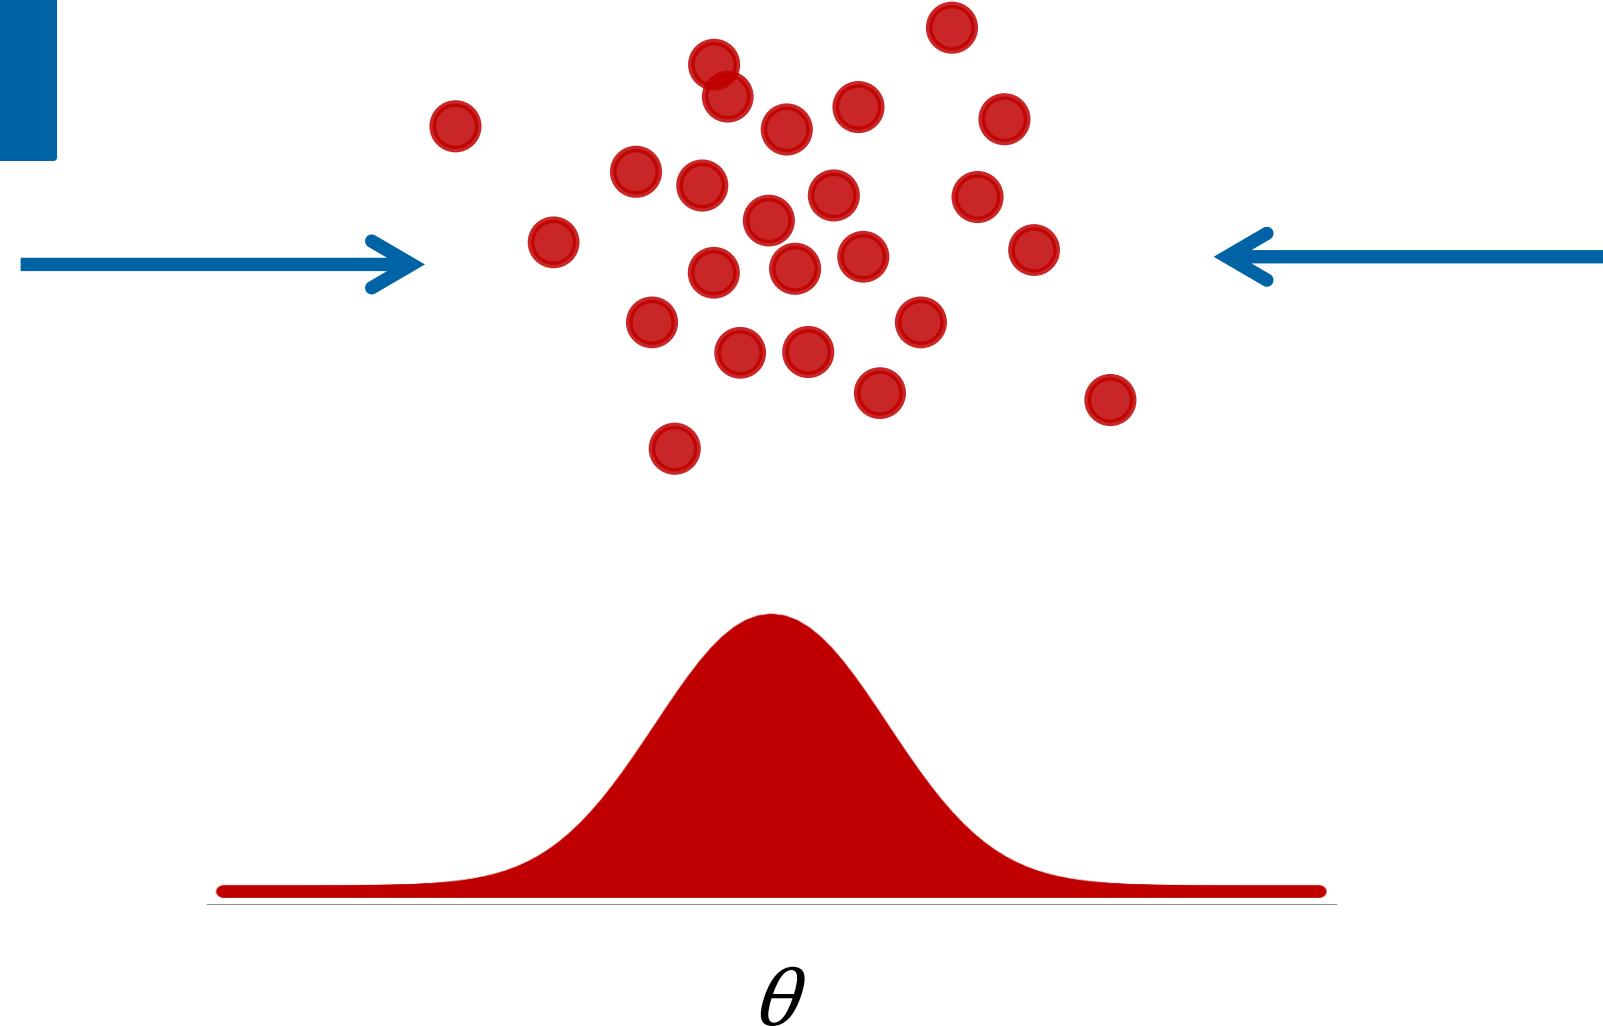
cognitive model  
statistics  
computing



# Fitting Multiple Participants

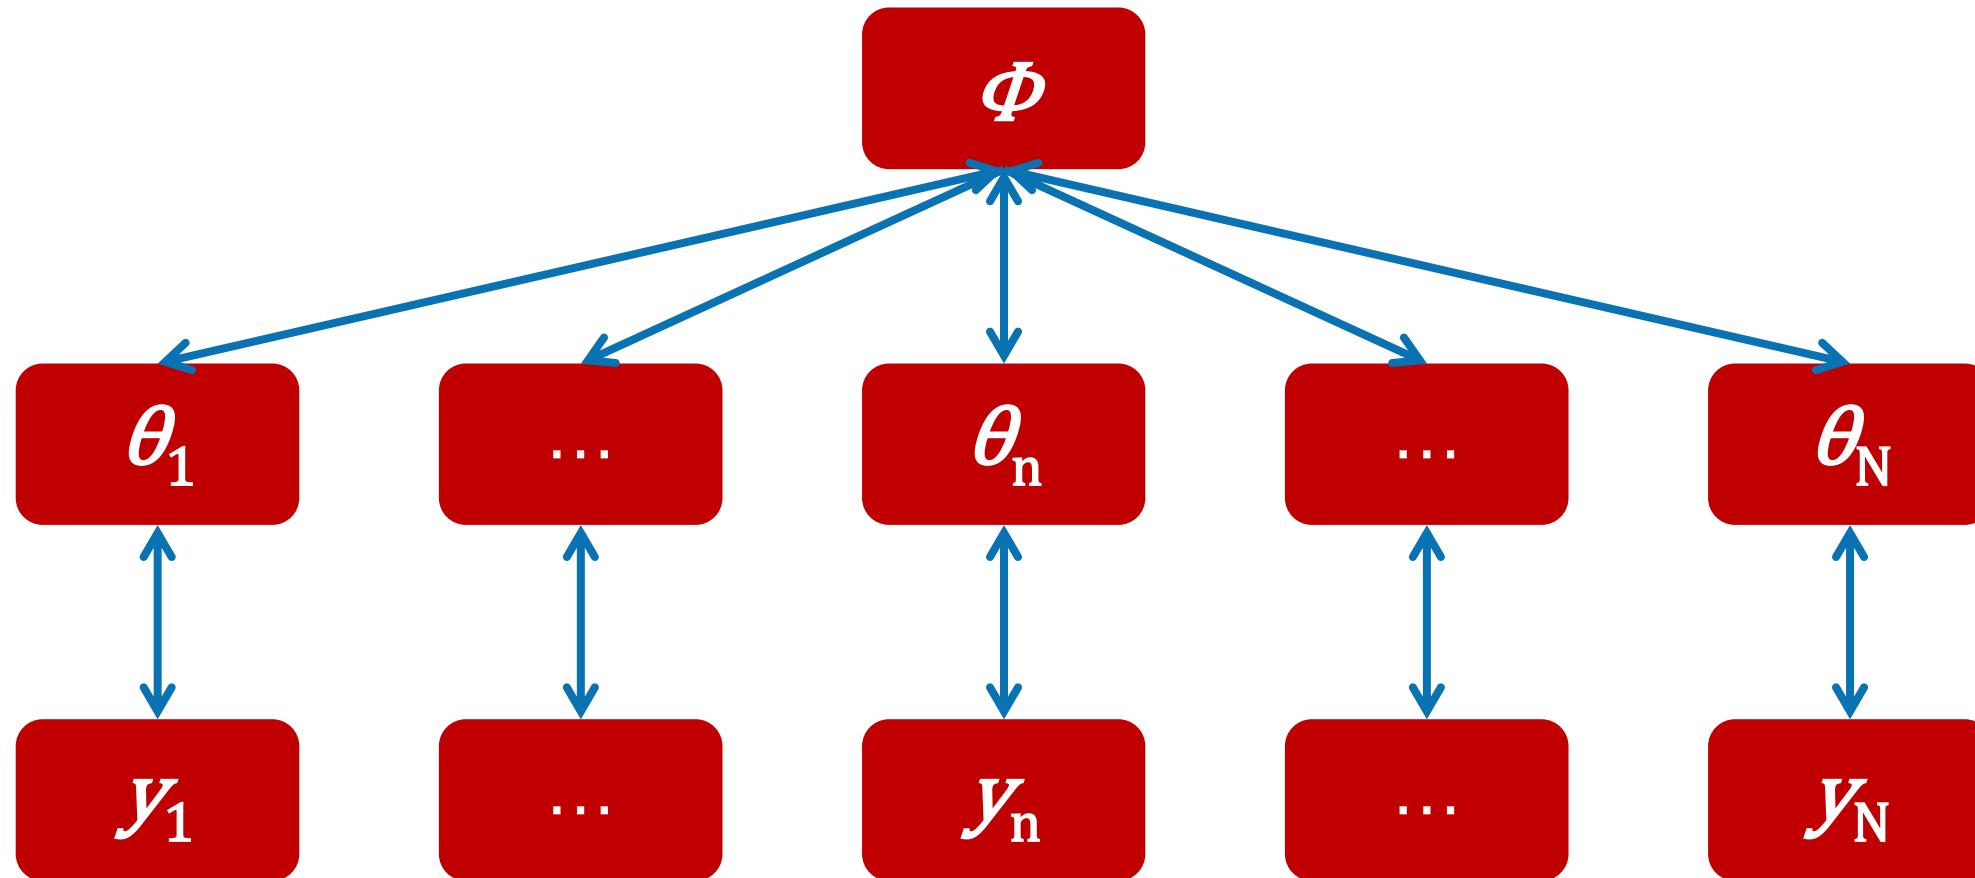
cognitive model  
statistics  
computing

shrinkage effect



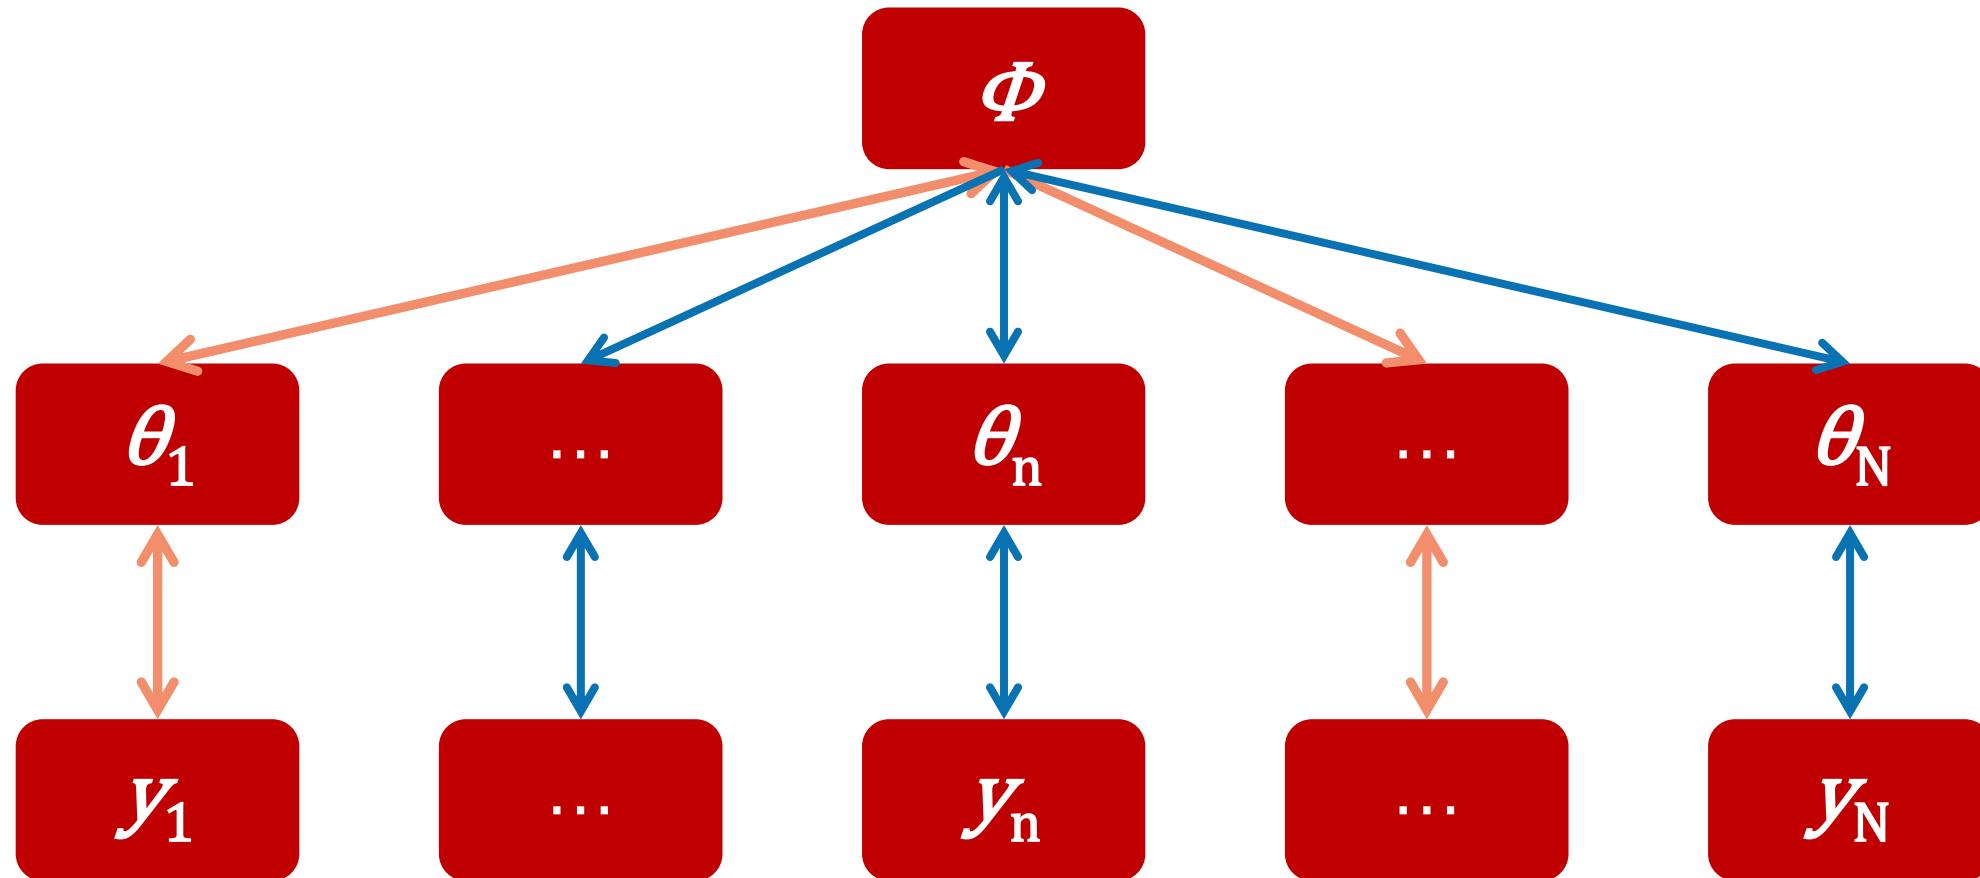
# Hierarchical Structure

cognitive model  
statistics  
computing



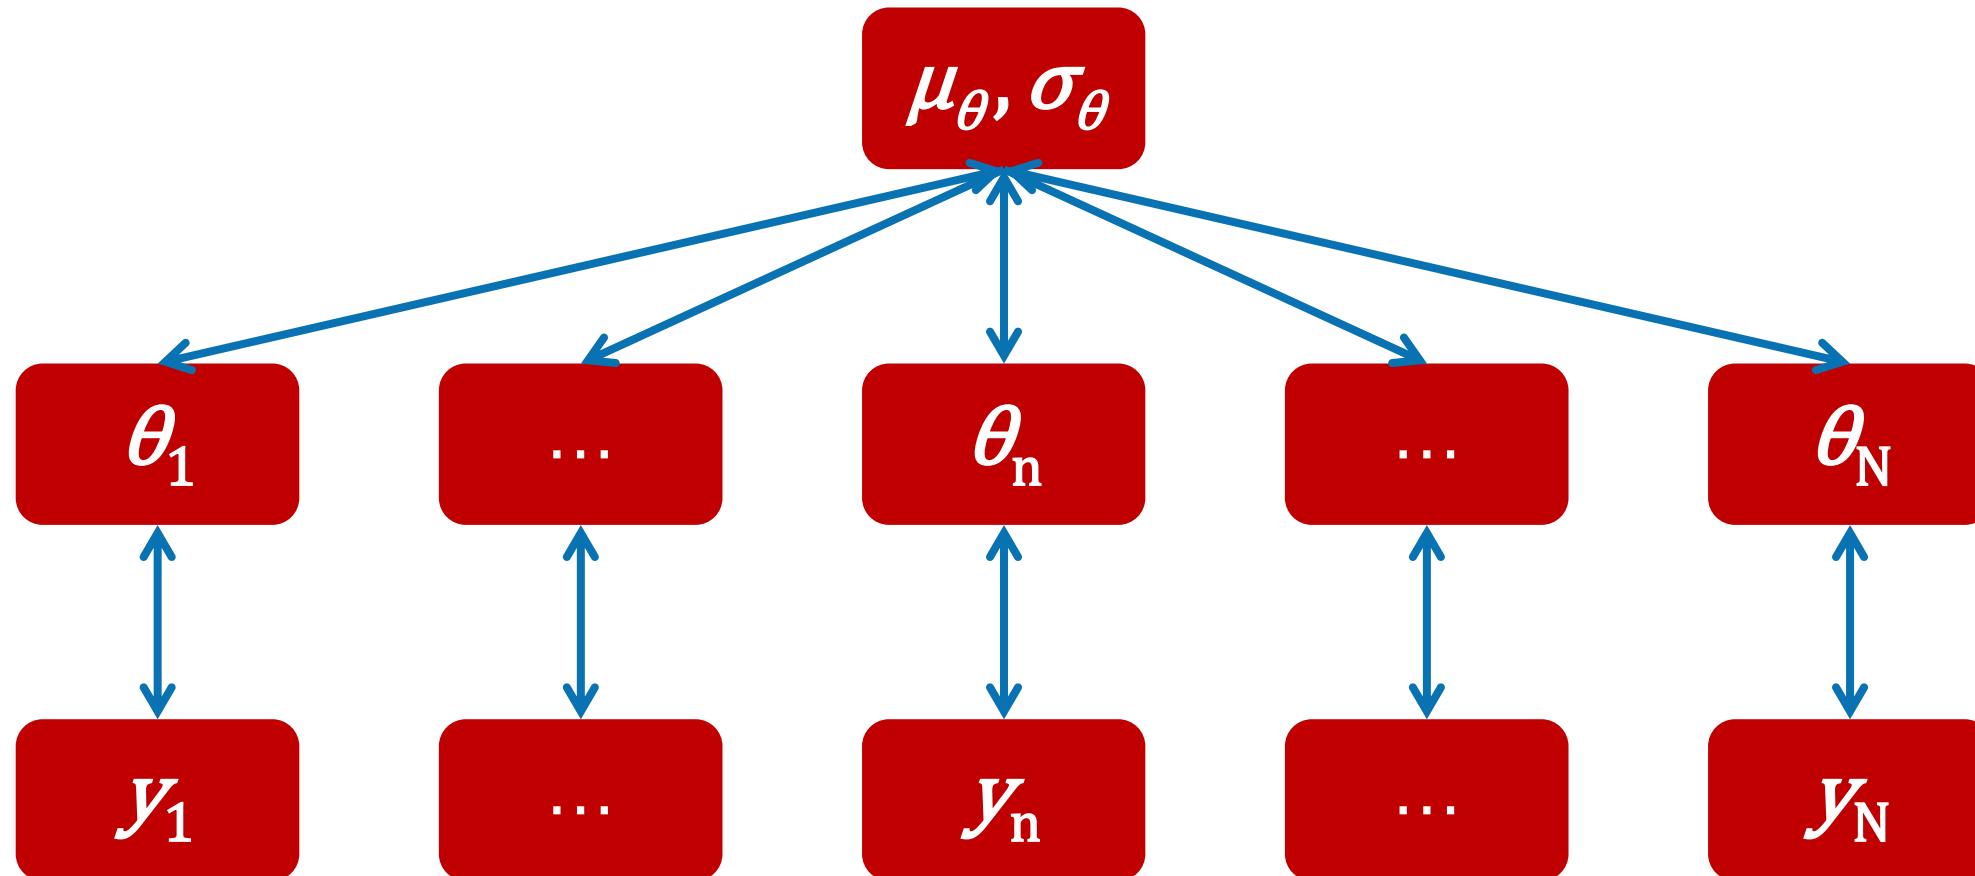
# Hierarchical Structure

cognitive model  
statistics  
computing

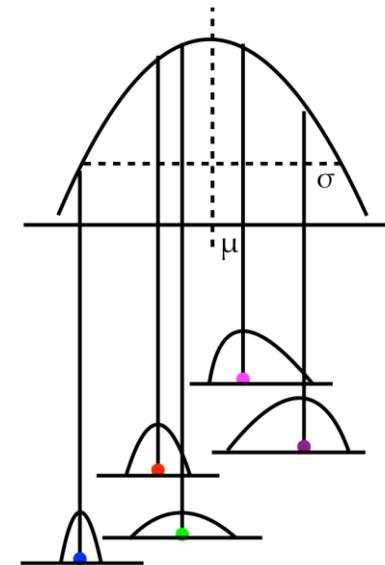
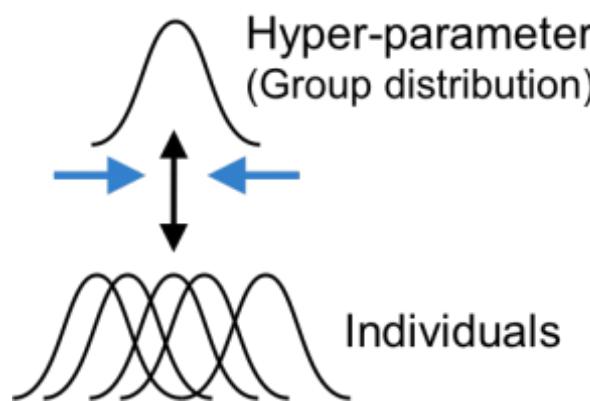
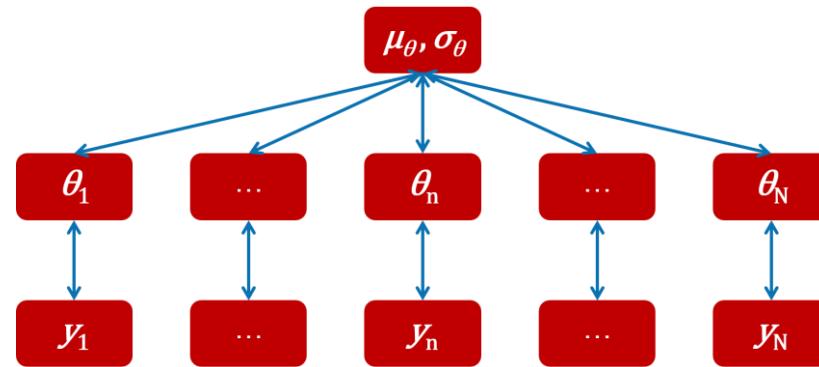


# Hierarchical Structure

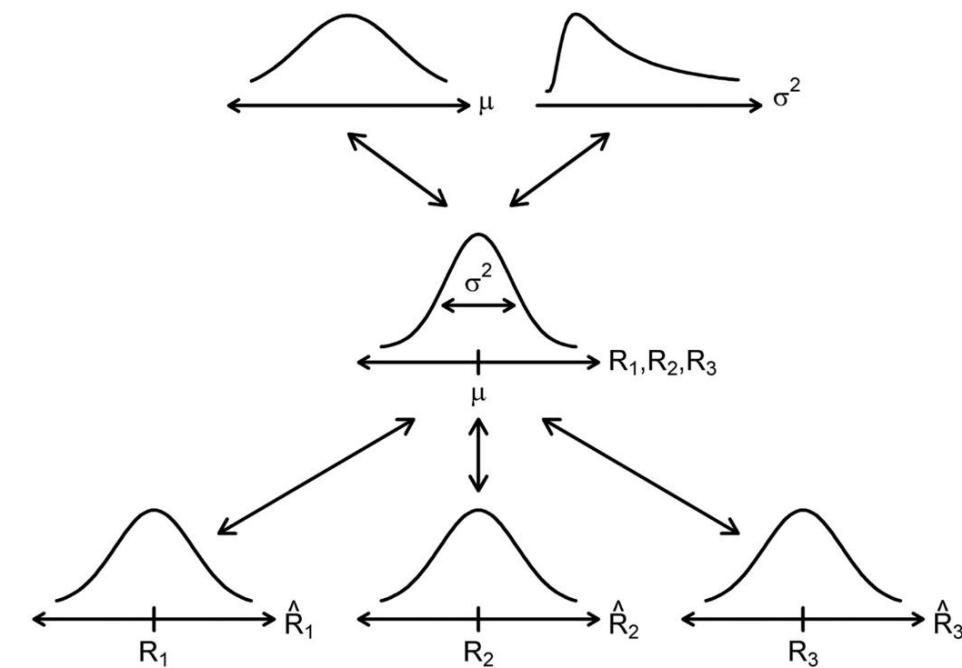
cognitive model  
statistics  
computing



# Hierarchical Structure



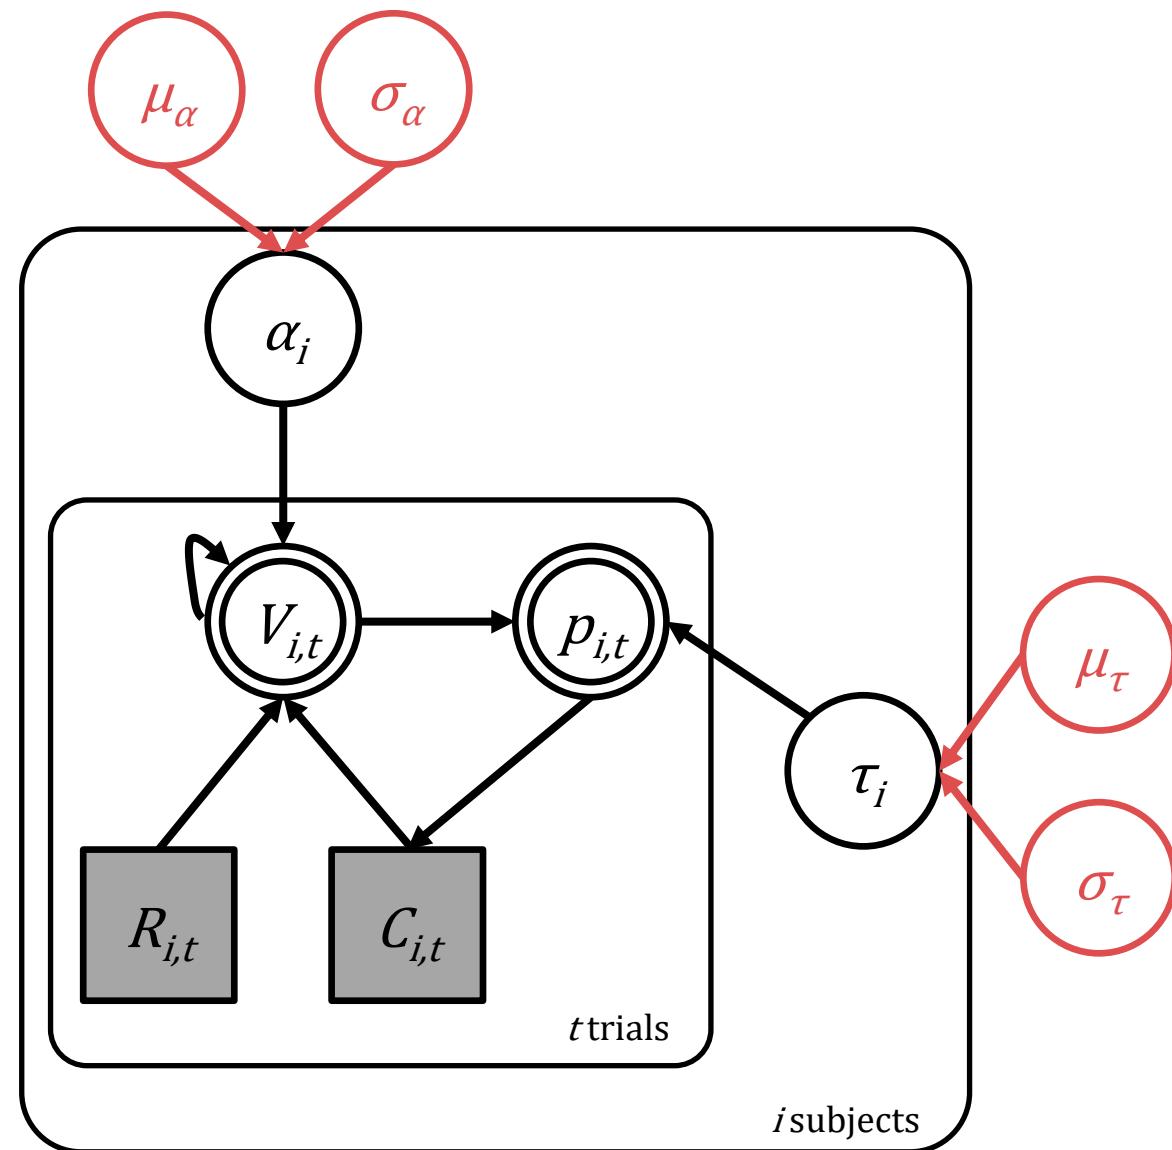
$$P(\Theta, \Phi | D) = \frac{P(D | \Theta, \Phi) P(\Theta, \Phi)}{P(D)} \propto P(D | \Theta) P(\Theta | \Phi) P(\Phi)$$



# Hierarchical RL Model

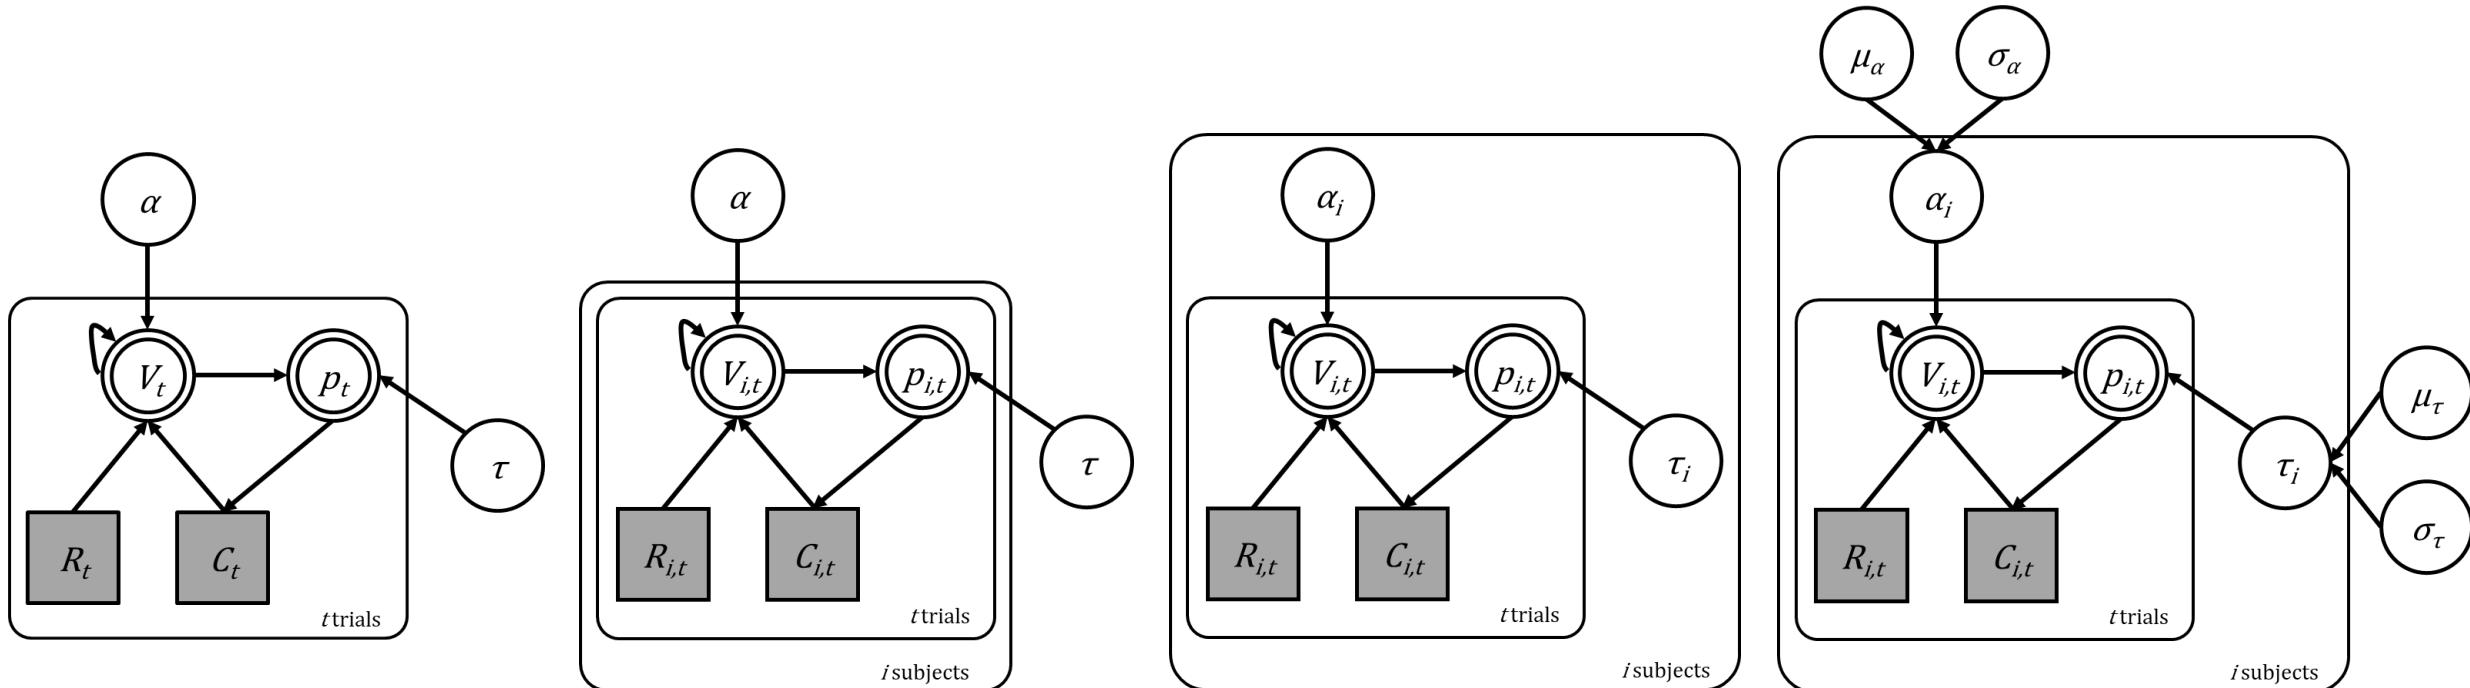
cognitive model  
statistics  
computing

Voilà!



# HOW DID WE GET HERE?

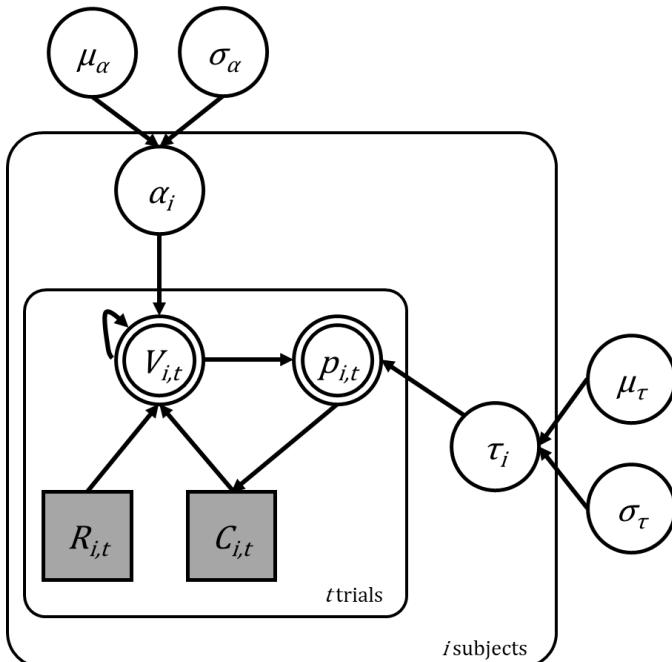
cognitive model  
statistics  
computing



The cognitive model *per se* is the same!

# Implementing Hierarchical RL Model

cognitive model  
statistics  
computing



$$\begin{aligned}\mu_\alpha &\sim Uniform(0,1) \\ \sigma_\alpha &\sim halfCauchy(0,1) \\ \mu_\tau &\sim Uniform(0,3) \\ \sigma_\tau &\sim halfCauchy(0,3) \\ \alpha_i &\sim Normal(\mu_\alpha, \sigma_\alpha)_{\mathcal{T}(0,1)} \\ \tau_i &\sim Normal(\mu_\tau, \sigma_\tau)_{\mathcal{T}(0,3)}\end{aligned}$$

$$p_{i,t}(C = A) = \frac{1}{1 + e^{\tau_i(V_{i,t}(B) - V_{i,t}(A))}}$$

$$V_{i,t+1}^c = V_{i,t}^C + \alpha_i(R_{i,t} - V_{i,t}^C)$$

```
parameters {
    real<lower=0,upper=1> lr_mu;
    real<lower=0,upper=3> tau_mu;
    real<lower=0> lr_sd;
    real<lower=0> tau_sd;
    real<lower=0,upper=1> lr[nSubjects];
    real<lower=0,upper=3> tau[nSubjects];
}

model {
    lr_sd ~ cauchy(0,1);
    tau_sd ~ cauchy(0,3);
    lr ~ normal(lr_mu, lr_sd) ;
    tau ~ normal(tau_mu, tau_sd) ;

    for (s in 1:nSubjects) {
        vector[2] v;
        real pe;
        v = initV;

        for (t in 1:nTrials) {
            choice[s,t] ~ categorical_logit( tau[s] * v );
            pe = reward[s,t] - v[choice[s,t]];
            v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;
        }
    }
}
```

# Exercise XI

cognitive model  
statistics  
computing

.../06.reinforcement\_learning/\_scripts/reinforcement\_learning\_multi\_parm\_main.R

TASK: (1) complete the model (TIP: individual ~ group)  
(2) fit the hierarchical RL model

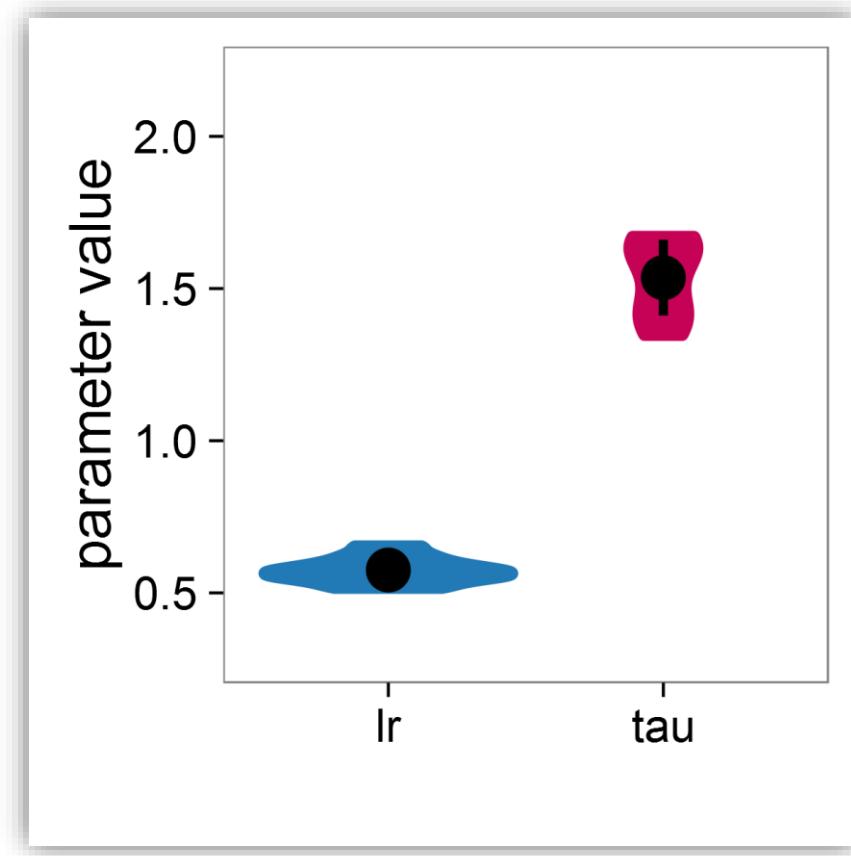
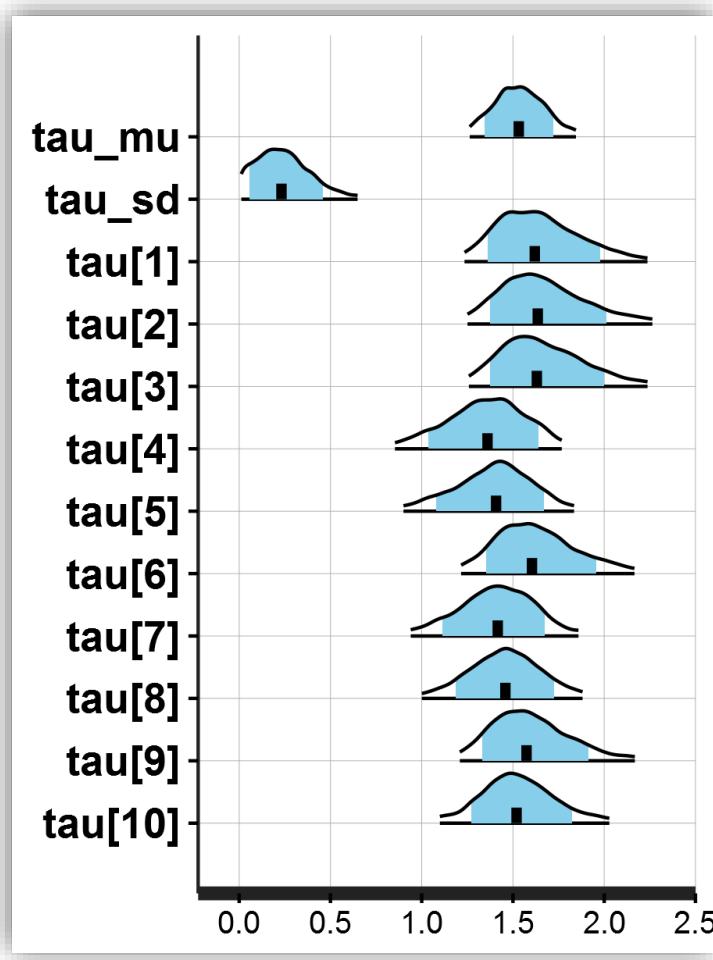
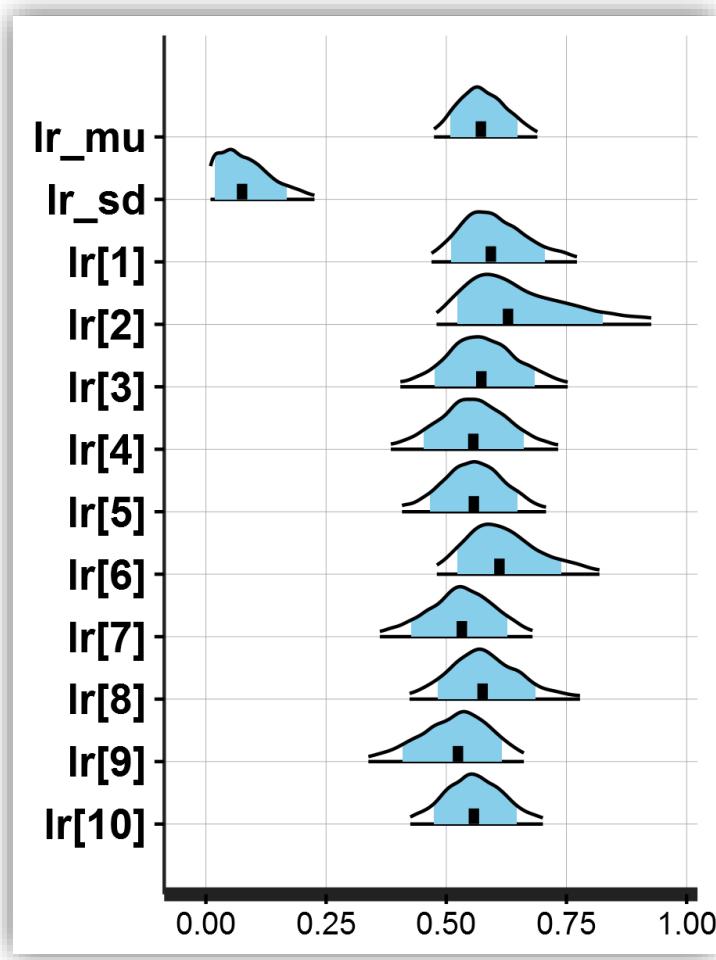
```
> source('_scripts/reinforcement_learning_multi_parm_main.R')  
  
> fit_rl3 <- run_rl_mp( modelType = 'hrch' )
```

In addition: Warning messages:

1: There were 97 divergent transitions after warmup. Increasing adapt\_delta above 0.8 may help. See <http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup>  
2: Examine the pairs() plot to diagnose sampling problems

# Hierarchical Fitting\*

cognitive model  
statistics  
computing

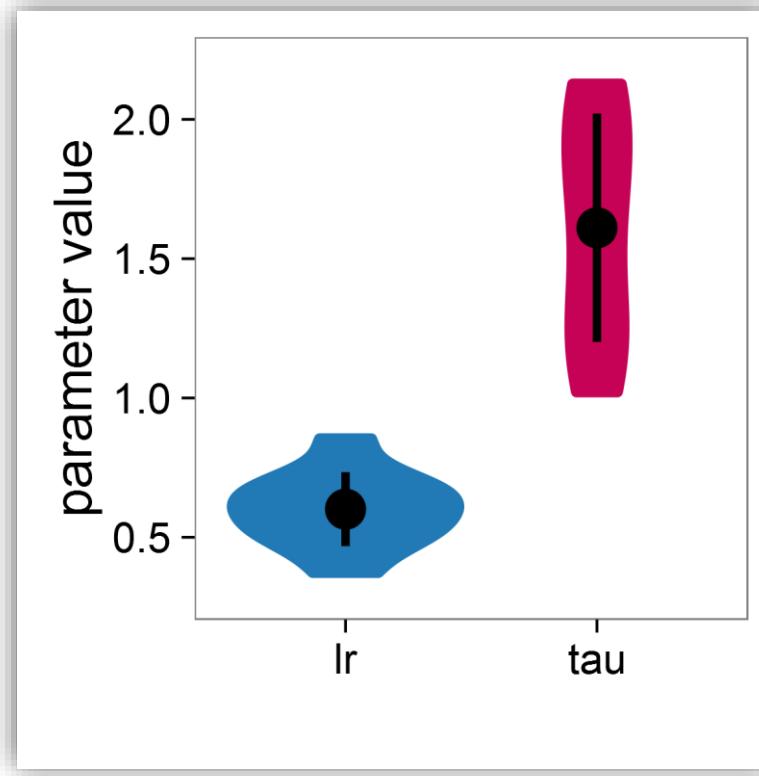


\*: adapt\_delta=0.999, max\_treedepth=100

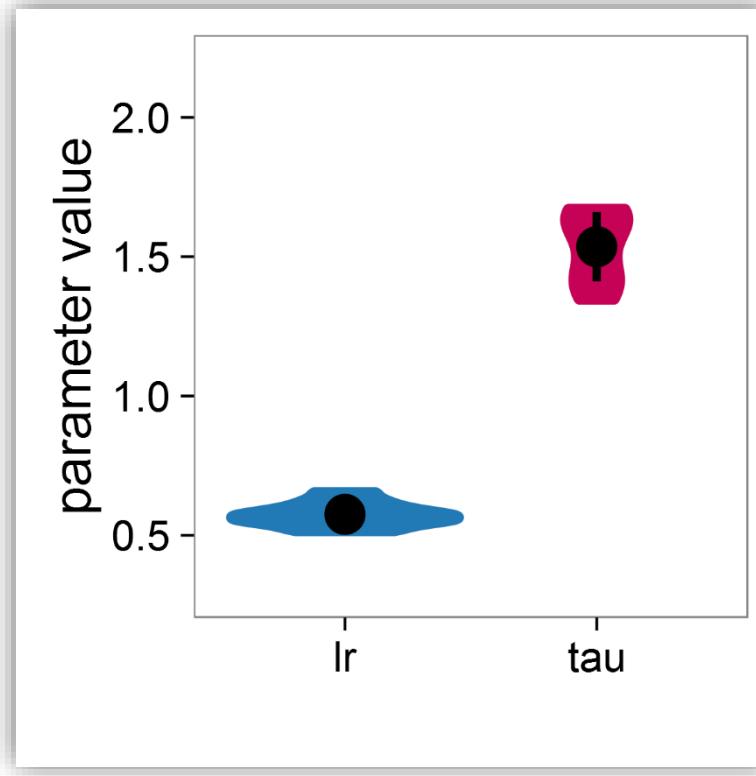
# Comparing with True Parameters

cognitive model  
statistics  
computing

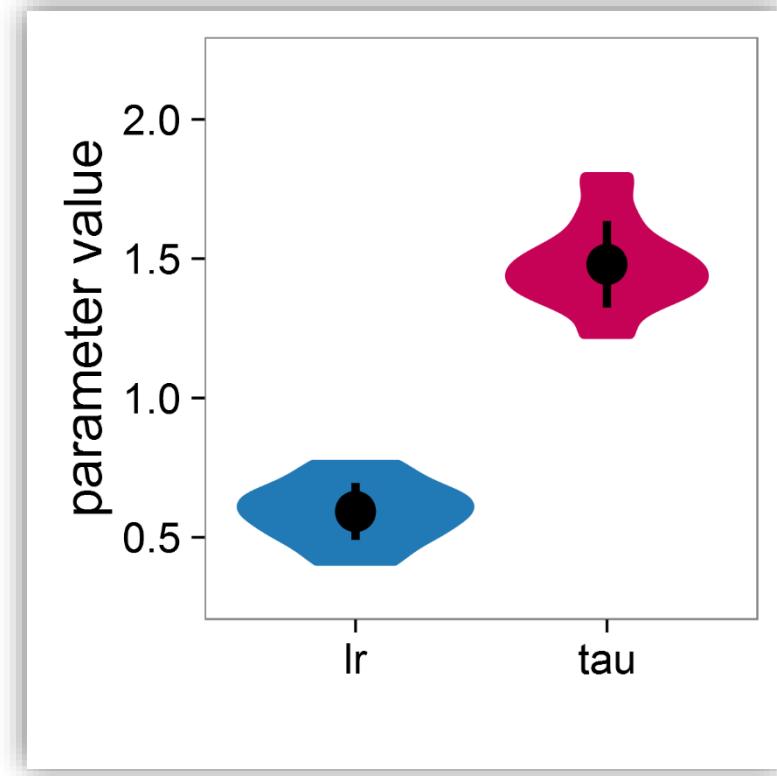
Posterior Means (indv)



Posterior Means (hrch)\*



True Parameters



\*: adapt\_delta=0.999, max\_treedepth=100

# Group-level Parameters

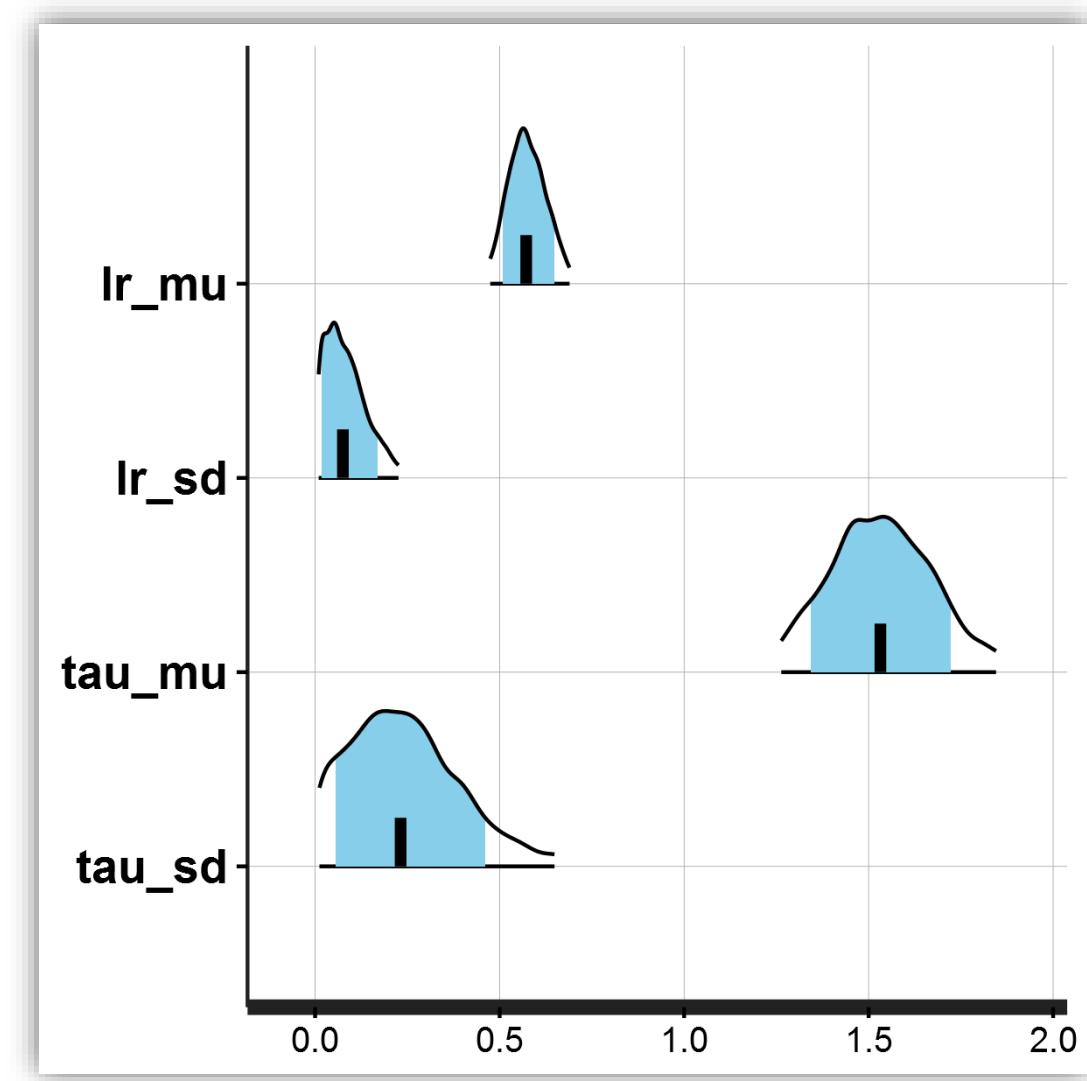
cognitive model  
statistics  
computing

## True group parameters

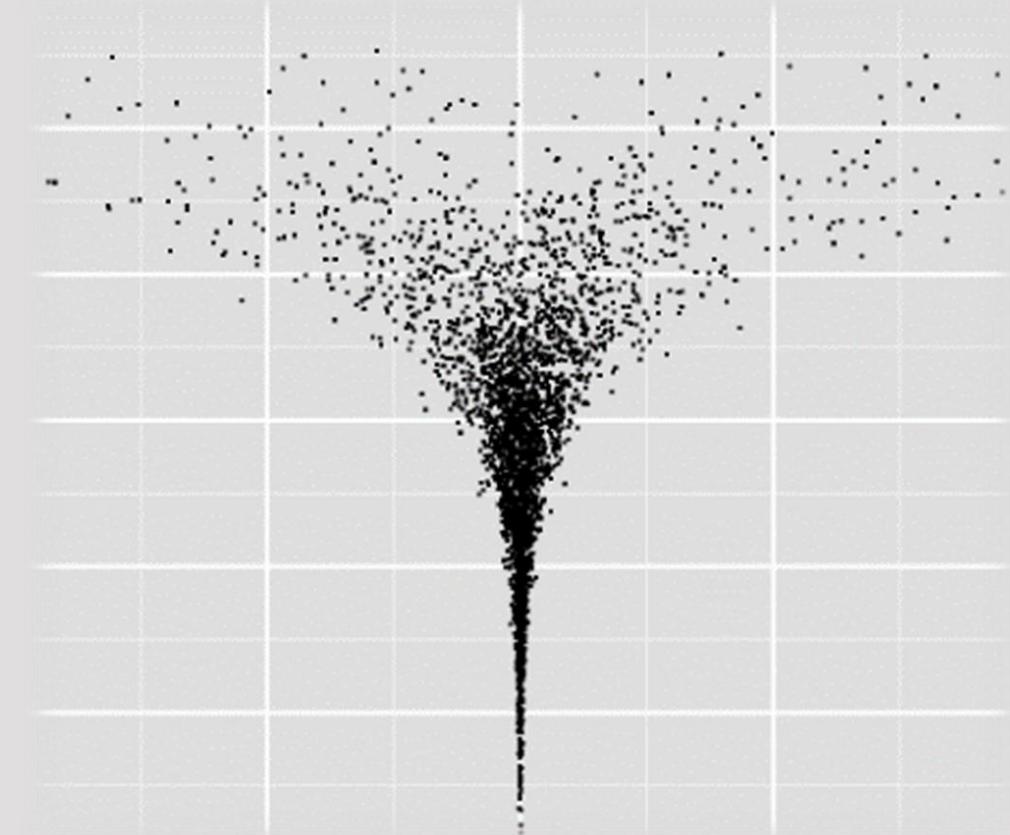
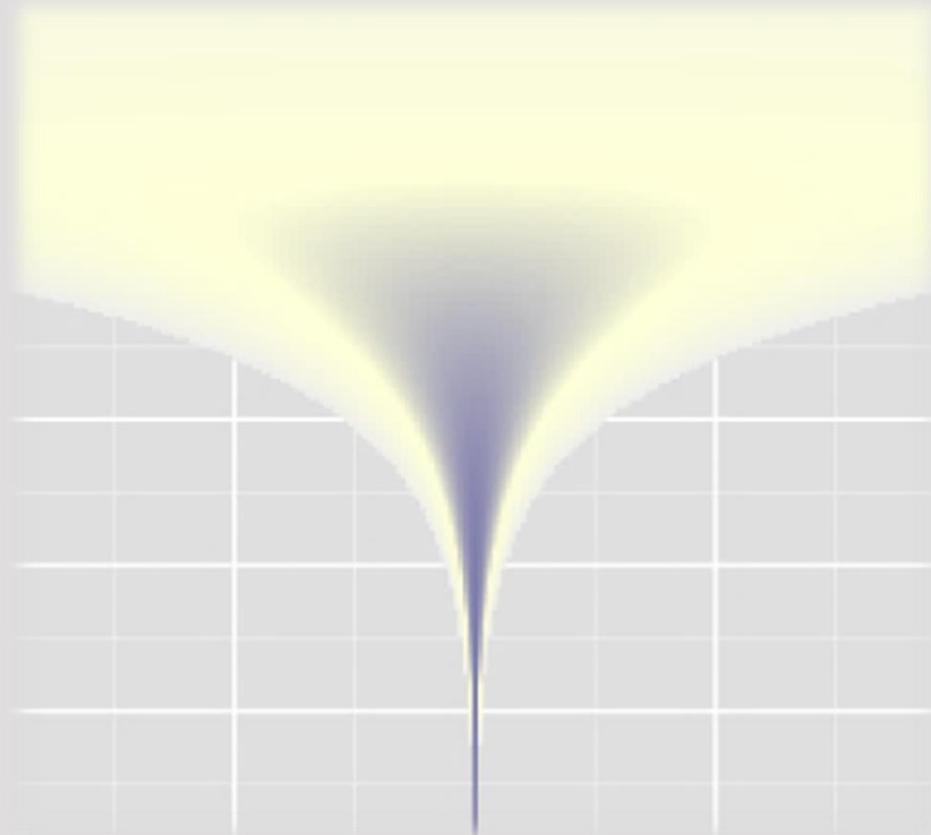
```
lr = rnorm(10, mean=0.6, sd=0.12)  
tau = rnorm(10, mean=1.5, sd=0.2)
```

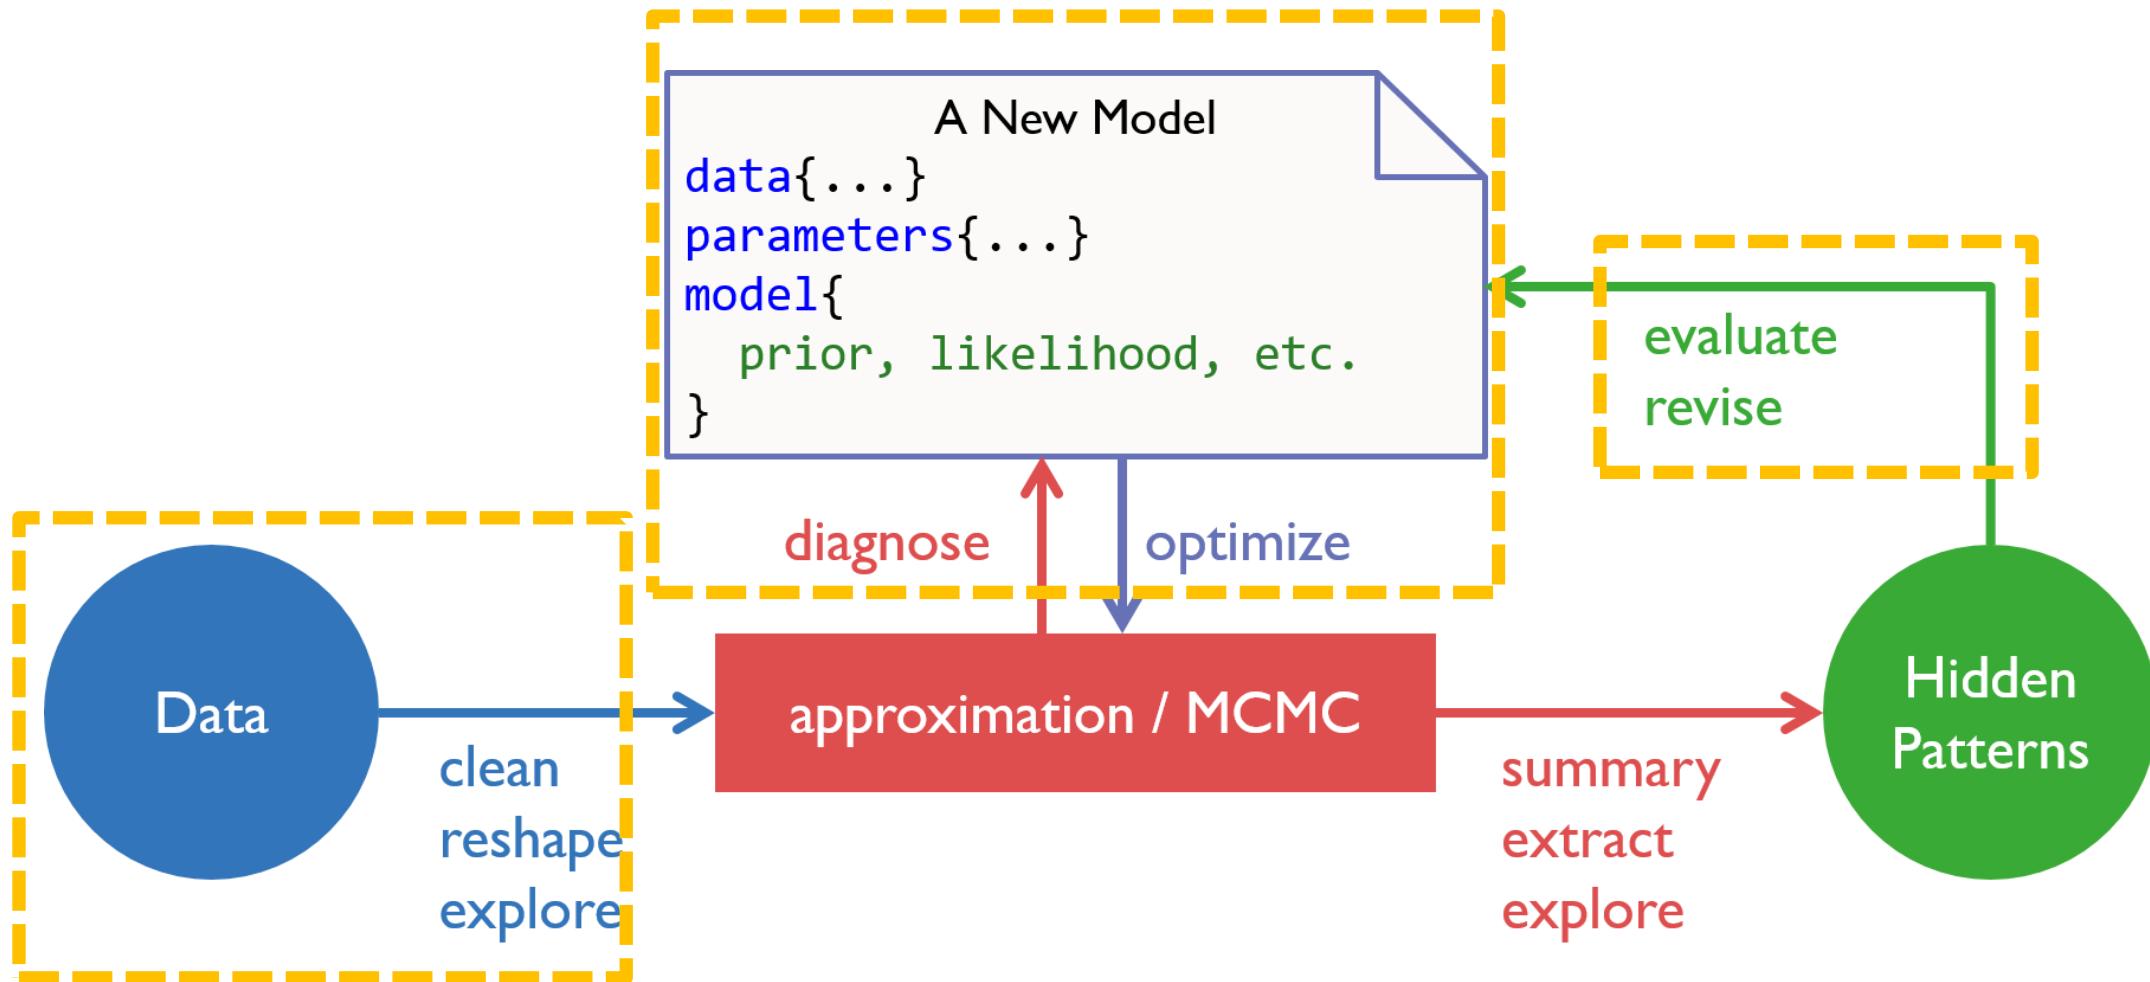
## Estimated group parameters

	mean	2.5%	25%	50%	75%	97.5%
lr_mu	0.58	0.47	0.54	0.57	0.61	0.69
lr_sd	0.09	0.01	0.04	0.08	0.12	0.23
tau_mu	1.54	1.26	1.43	1.53	1.63	1.85
tau_sd	0.25	0.01	0.13	0.23	0.34	0.65



# OPTIMIZING STAN CODES







# Optimizing Stan Code

cognitive model  
statistics  
computing

## Preprocess data

run as many calculations as you can outside Stan

## Specify a proper model

follow literature, supervision, experience, etc.

## Vectorizing

vectorize Stan code whenever you can

## Reparameterizing

reparameterize target parameter to simple distributions

# Preprocess Data

$$\overline{\text{height}} = \alpha + \beta_1 * \text{weight} + \beta_2 * \text{weight}^2$$

```
d$weight_sq <- d$weight^2
```

```
data {
  int<lower=0> N;
  vector<lower=0>[N] height;
  vector<lower=0>[N] weight;
  vector<lower=0>[N] weight_sq;
}
```

# Specify a Proper Model

- Visualize your data
- Follow Literatures
- Start from simple and then build complexities
- Simulate data and run model recovery

A New Model

```
data{...}  
parameters{...}  
model{  
    prior, likelihood, etc.  
}
```

# Vectorization

cognitive model  
statistics  
computing

```
model {  
    for (n in 1:N) {  
        flip[n] ~ bernoulli(theta);  
    }  
}
```

```
model {  
    flip ~ bernoulli(theta);  
}
```

```
parameters {  
    ...  
    real<lower=0,upper=1> lr[nSubjects];  
    real<lower=0,upper=3> tau[nSubjects];  
}  
  
model {  
    ...  
    lr      ~ normal(lr_mu, lr_sd) ;  
    tau    ~ normal(tau_mu, tau_sd) ;  
    ...  
}
```

```
model {  
    vector[N] mu;  
    for (i in 1:N) {  
        mu[i] = alpha + beta * weight[i];  
        height[i] ~ normal(mu[i], sigma)  
    }  
}
```

```
model {  
    vector[N] mu;  
    mu = alpha + beta * weight;  
    height ~ normal(mu, sigma);  
}
```

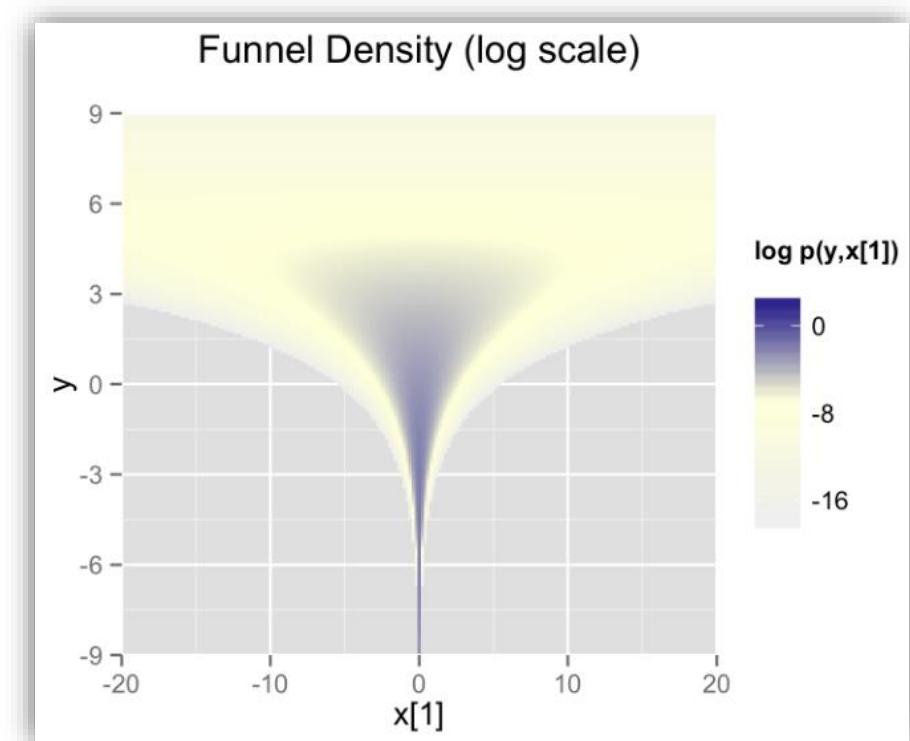
```
model {  
    height ~ normal(alpha + beta * weight, sigma);  
}
```

# Reparameterization

## Neal's Funnel

$$p(y, \mathbf{x}) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {
  real y;
  vector[9] x;
}
model {
  y ~ normal(0,3);
  x ~ normal(0,exp(y/2));
}
```



# Non-centered Reparameterization\*

cognitive model
statistics
computing

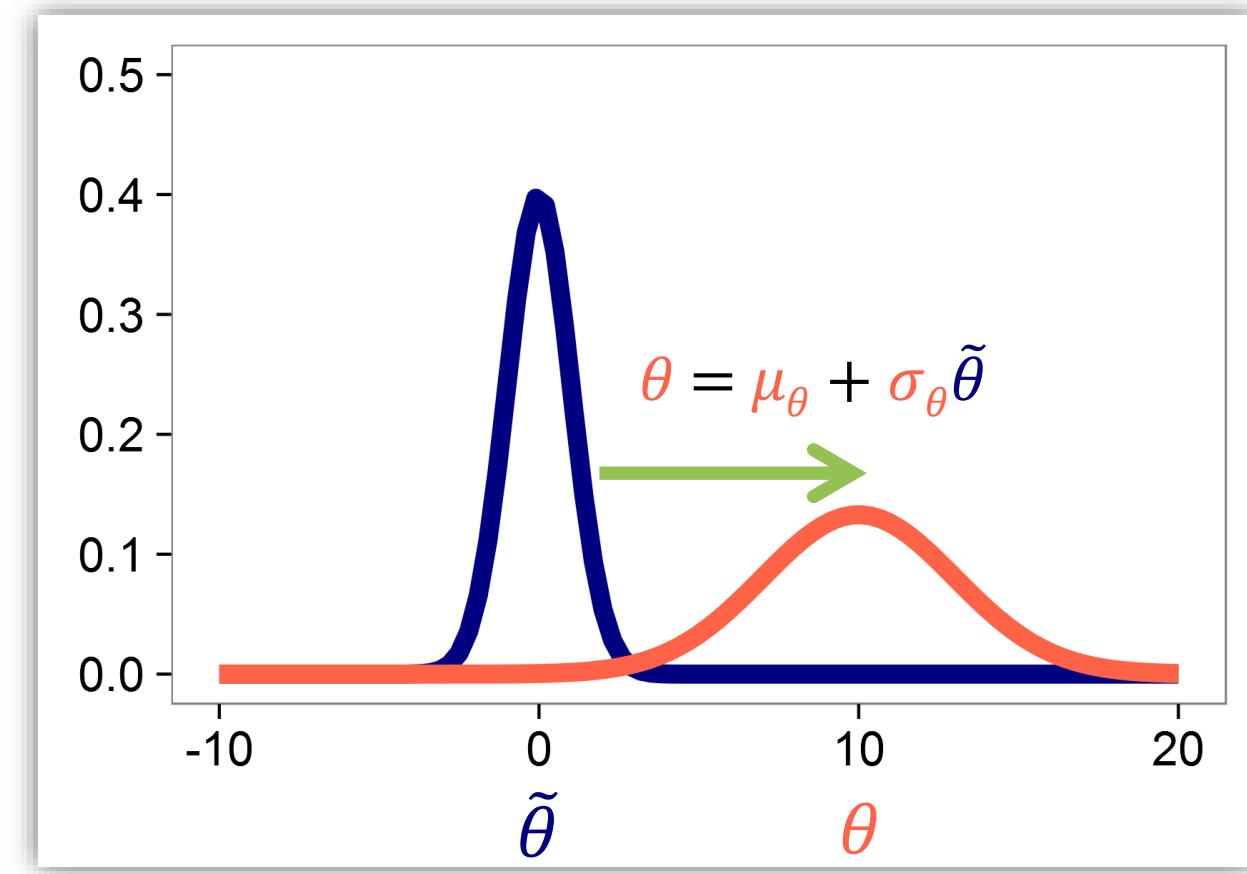
$$\theta \sim Normal(\mu_\theta, \sigma_\theta)$$



$$\tilde{\theta} \sim Normal(0, 1)$$

$$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$$

Stan likes **simple** distributions!

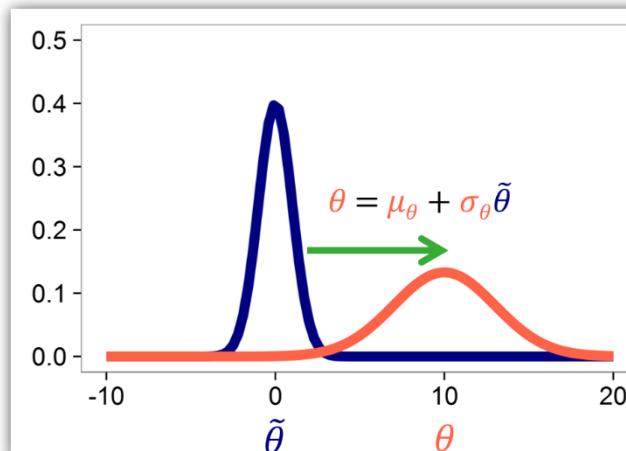


# Reparameterization

## Neal's Funnel

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {
  real y;
  vector[9] x;
}
model {
  y ~ normal(0,3);
  x ~ normal(0,exp(y/2));
}
```



```
parameters {
  real y_raw;
  vector[9] x_raw;
}
transformed parameters {
  real y;
  vector[9] x;
}

y = 3.0 * y_raw;
x = exp(y/2) * x_raw;

model {
  y_raw ~ normal(0,1);
  x_raw ~ normal(0,1);
```

# Stan Sampling Parameters

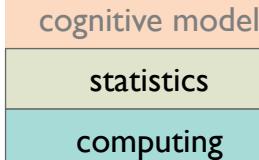
parameter	description	constraint	default
iterations	number of MCMC samples (per chain)	int, $> 0$	2000
delta: $\delta$	target Metropolis acceptance rate	$\delta \in [0, 1]$	0.80
stepsize: $\varepsilon$	initial HMC step size	real, $\varepsilon > 0$	2.0
max_treedepth: $L$	maximum HMC steps per iteration	int, $L > 0$	10

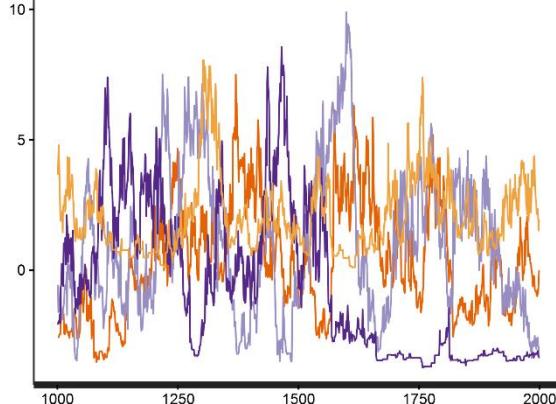
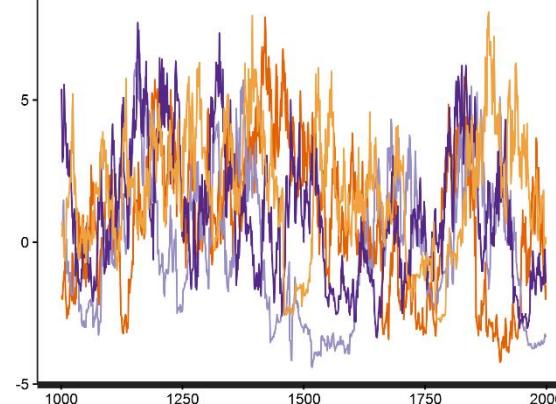
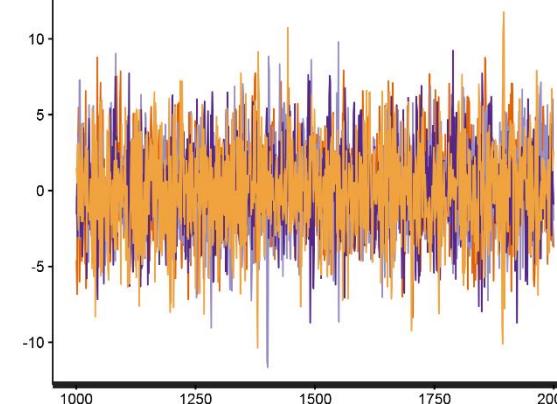
## Typical adjustments

- Increase iterations
- Increase delta
- Decrease stepsize
- Might have to increase max\_treedepth

```
funnel_fit2 <- stan("_scripts/funnel.stan",
  iter = 4000,
  control = list(adapt_delta = 0.999,
    stepsize = 1.0,
    max_treedepth = 20))
```

# Neal's Funnel: Comparing Performance



	<b>direct model</b>	<b>adjusted direct model</b>	<b>reparameterized model</b>
Rhat ( $y$ )	1.22	1.1	1.0
n_eff ( $y$ )	18	42	3886
runtime*	48.50 sec	50.76 sec	50.12 sec
n_eff ( $y$ ) / runtime	0.37 / sec	0.82 / sec	77.53 / sec
n_divergent	53	0	0
traceplot ( $y$ )			

\*: 2 cores in parallel, including compiling time

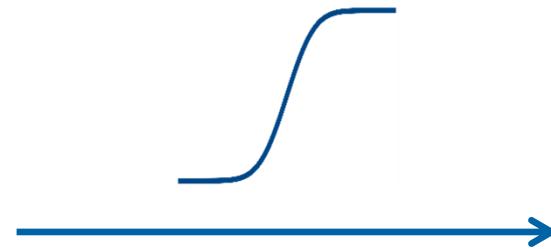
# How about Bounded Parameters?

cognitive model  
statistics  
computing

$$\tilde{\theta} \sim Normal(0, 1)$$

$$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$$

$$\theta \in (-\infty, +\infty)$$



$$\tilde{\theta} \sim Normal(0, 1)$$

$$\theta = Probit^{-1}(\mu_\theta + \sigma_\theta \tilde{\theta})$$

$$\theta \in [0, 1]$$

constraint	reparameterization
$\theta \in (-\infty, +\infty)$	$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$
$\theta \in [0, N]$	$\theta = Probit^{-1}(\mu_\theta + \sigma_\theta \tilde{\theta}) \times N$
$\theta \in [M, N]$	$\theta = Probit^{-1}(\mu_\theta + \sigma_\theta \tilde{\theta}) \times (N-M) + M$
$\theta \in (0, +\infty)$	$\theta = exp(\mu_\theta + \sigma_\theta \tilde{\theta}); \theta = log(1+exp(\mu_\theta + \sigma_\theta \tilde{\theta}))$

\* Probit<sup>-1</sup>: Normal cumulative distribution function (normcdf)

# Apply to Our Hierarchical RL Model

cognitive model  
statistics  
computing

```
parameters {  
    real<lower=0,upper=1> lr_mu;  
    real<lower=0,upper=3> tau_mu;  
  
    real<lower=0> lr_sd;  
    real<lower=0> tau_sd;  
  
    real<lower=0,upper=1> lr[nSubjects];  
    real<lower=0,upper=3> tau[nSubjects];  
}
```



```
parameters {  
    # group-Level parameters  
    real lr_mu_raw;  
    real tau_mu_raw;  
    real<lower=0> lr_sd_raw;  
    real<lower=0> tau_sd_raw;  
  
    # subject-Level raw parameters  
    vector[nSubjects] lr_raw;  
    vector[nSubjects] tau_raw;  
}  
  
transformed parameters {  
    vector<lower=0,upper=1>[nSubjects] lr;  
    vector<lower=0,upper=3>[nSubjects] tau;  
  
    for (s in 1:nSubjects) {  
        lr[s] = Phi_approx( lr_mu_raw + lr_sd_raw * lr_raw[s] );  
        tau[s] = Phi_approx( tau_mu_raw + tau_sd_raw * tau_raw[s] ) * 3;  
    }  
}
```

# Apply to Our Hierarchical RL Model

cognitive model  
statistics  
computing

```
model {  
    lr_sd ~ cauchy(0,1);  
    tau_sd ~ cauchy(0,3);  
    lr ~ normal(lr_mu, lr_sd) ;  
    tau ~ normal(tau_mu, tau_sd) ;  
  
    for (s in 1:nSubjects) {  
        vector[2] v;  
        real pe;  
        v = initV;  
  
        for (t in 1:nTrials) {  
            choice[s,t] ~ categorical_logit( tau[s] * v );  
            pe = reward[s,t] - v[choice[s,t]];  
            v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
        }  
    }  
}
```



```
model {  
    lr_mu_raw ~ normal(0,1);  
    tau_mu_raw ~ normal(0,1);  
    lr_sd_raw ~ cauchy(0,3);  
    tau_sd_raw ~ cauchy(0,3);  
  
    lr_raw ~ normal(0,1);  
    tau_raw ~ normal(0,1);  
  
    for (s in 1:nSubjects) {  
        ...  
  
generated quantities {  
    real<lower=0,upper=1> lr_mu;  
    real<lower=0,upper=3> tau_mu;  
  
    lr_mu = Phi_approx(lr_mu_raw);  
    tau_mu = Phi_approx(tau_mu_raw) * 3;  
}
```

# Exercise XII

cognitive model  
statistics  
computing

.../07.optm\_rl/\_scripts/reinforcement\_learning\_hrch\_main.R

TASK: (1) Complete the Matt Trick

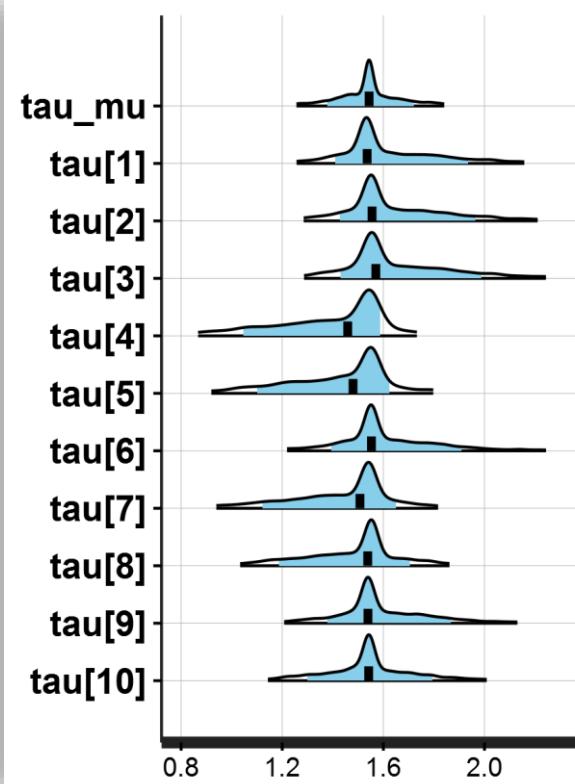
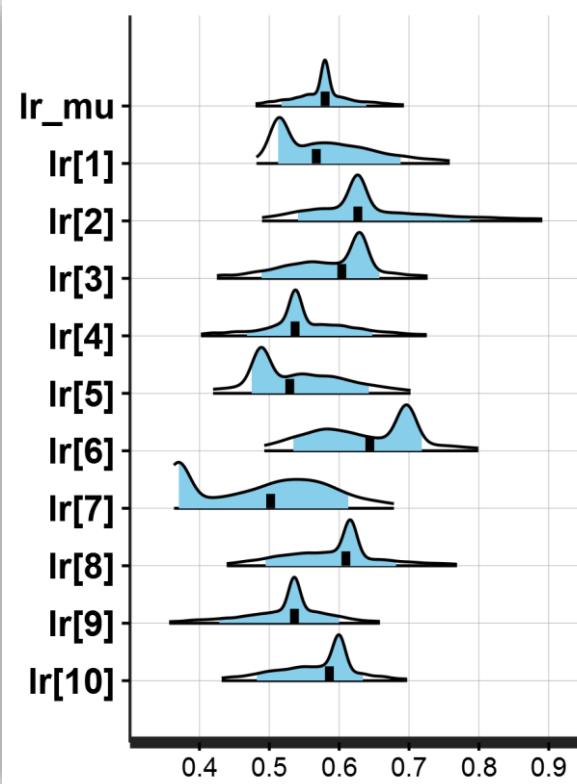
(2) fit the optimized hierarchical RL model

```
> source('_scripts/reinforcement_learning_hrch_main.R')  
  
> fit_rl4 <- run_rl_mp2(optimized = TRUE)
```

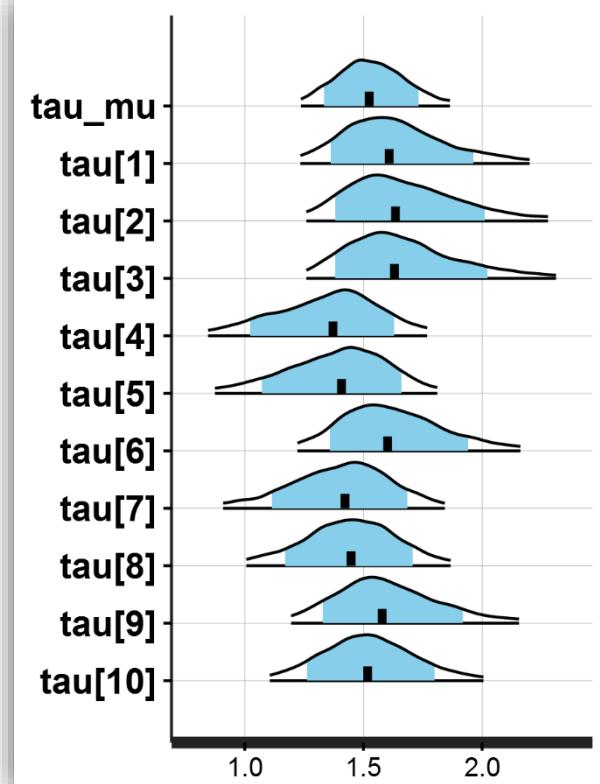
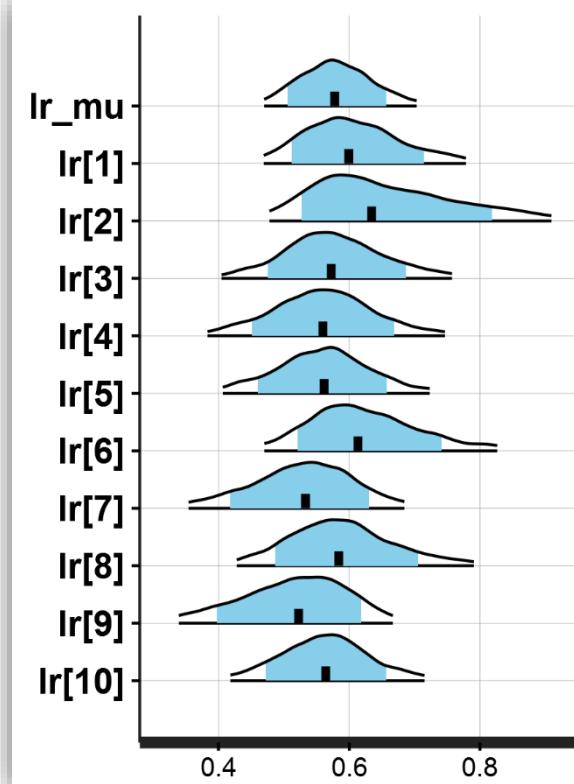
# Hierarchical Fitting – Optimized

cognitive model  
statistics  
computing

Posterior Means (hrch)



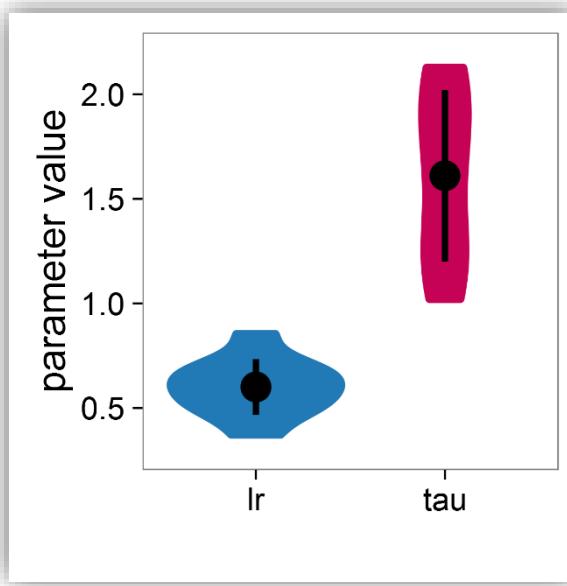
Posterior Means (hrch + optim)



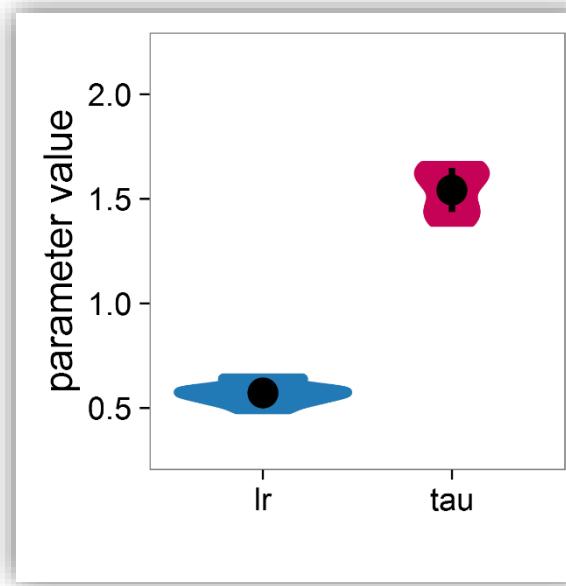
# Comparing with True Parameters

cognitive model  
statistics  
computing

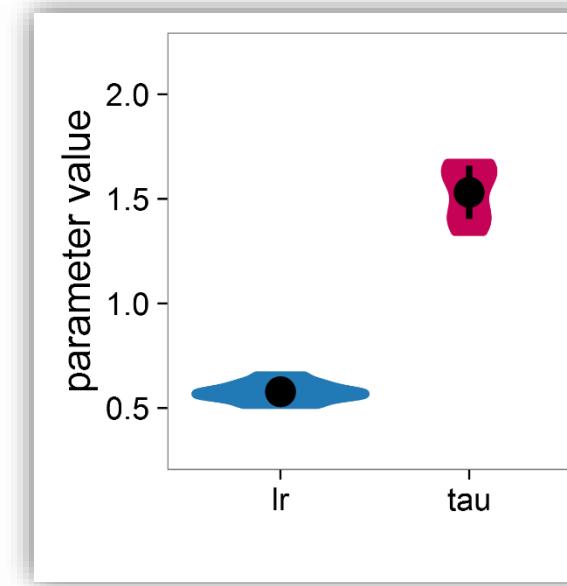
Posterior Means (indv)



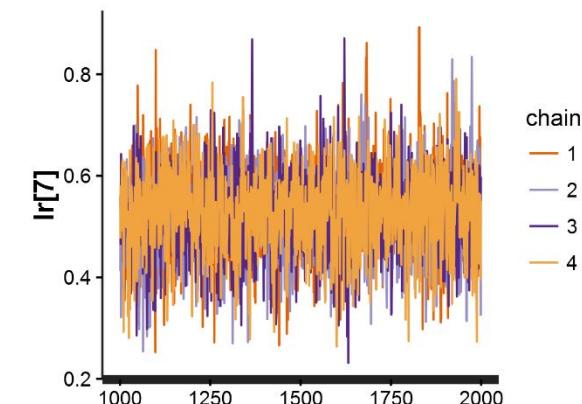
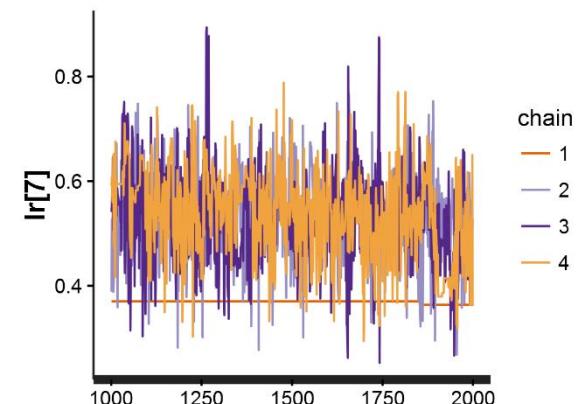
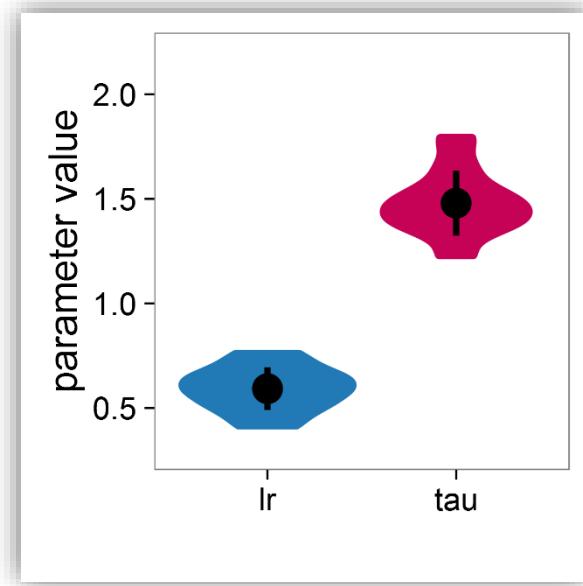
Posterior Means (hrch)



Posterior Means (hrch+optm)

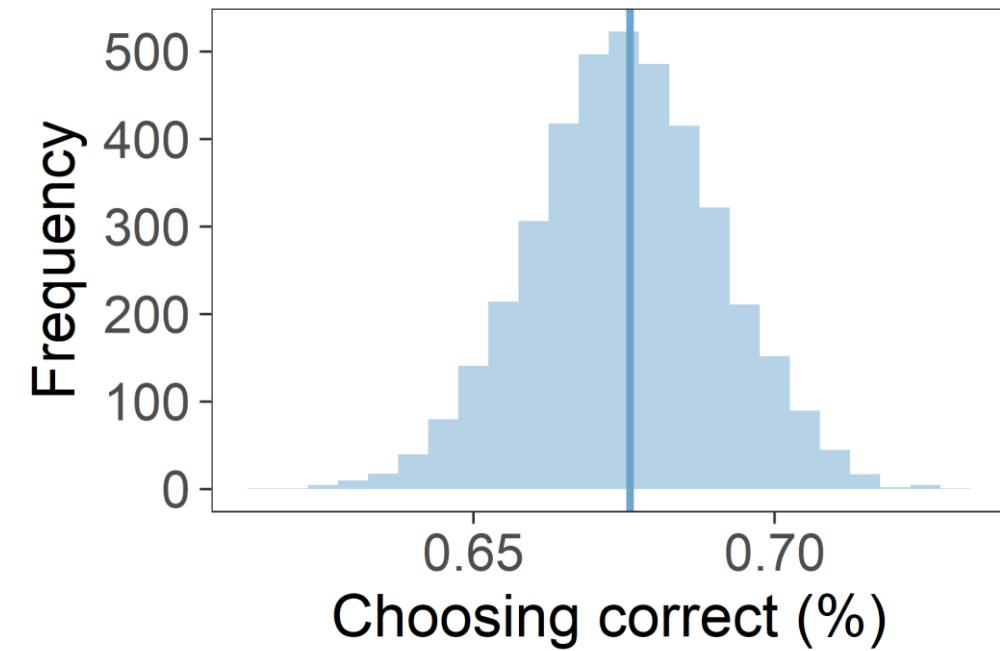
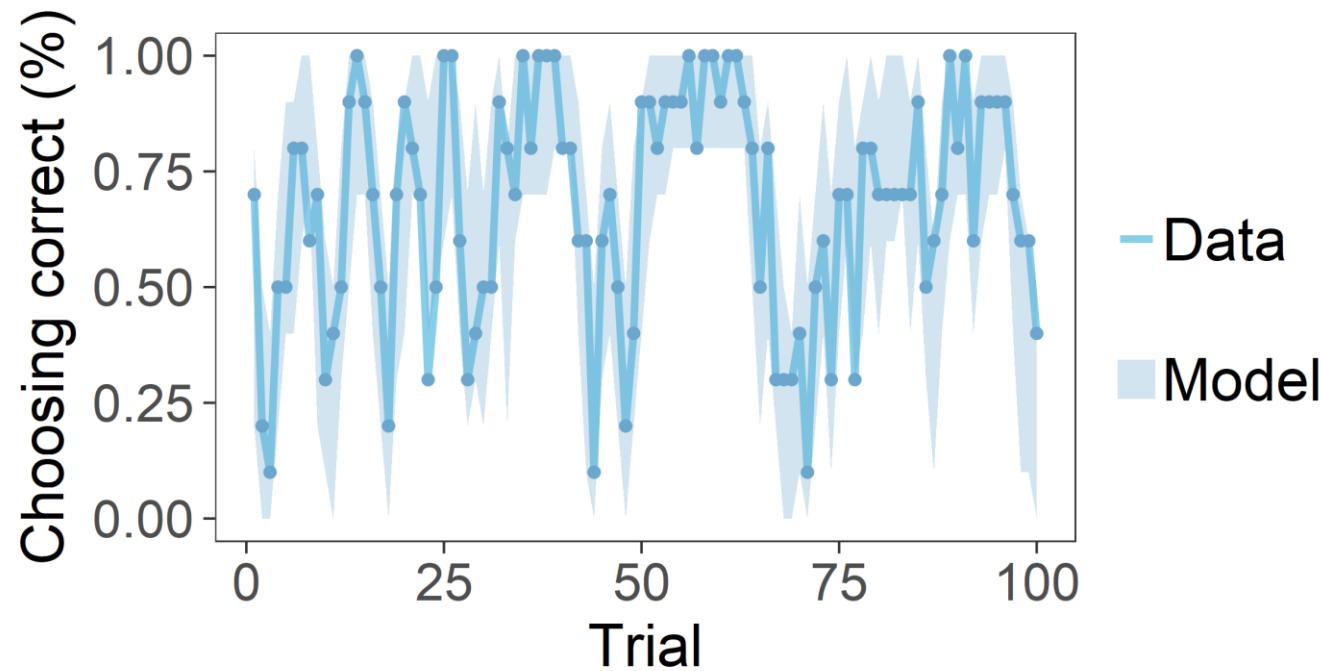


True Parameters

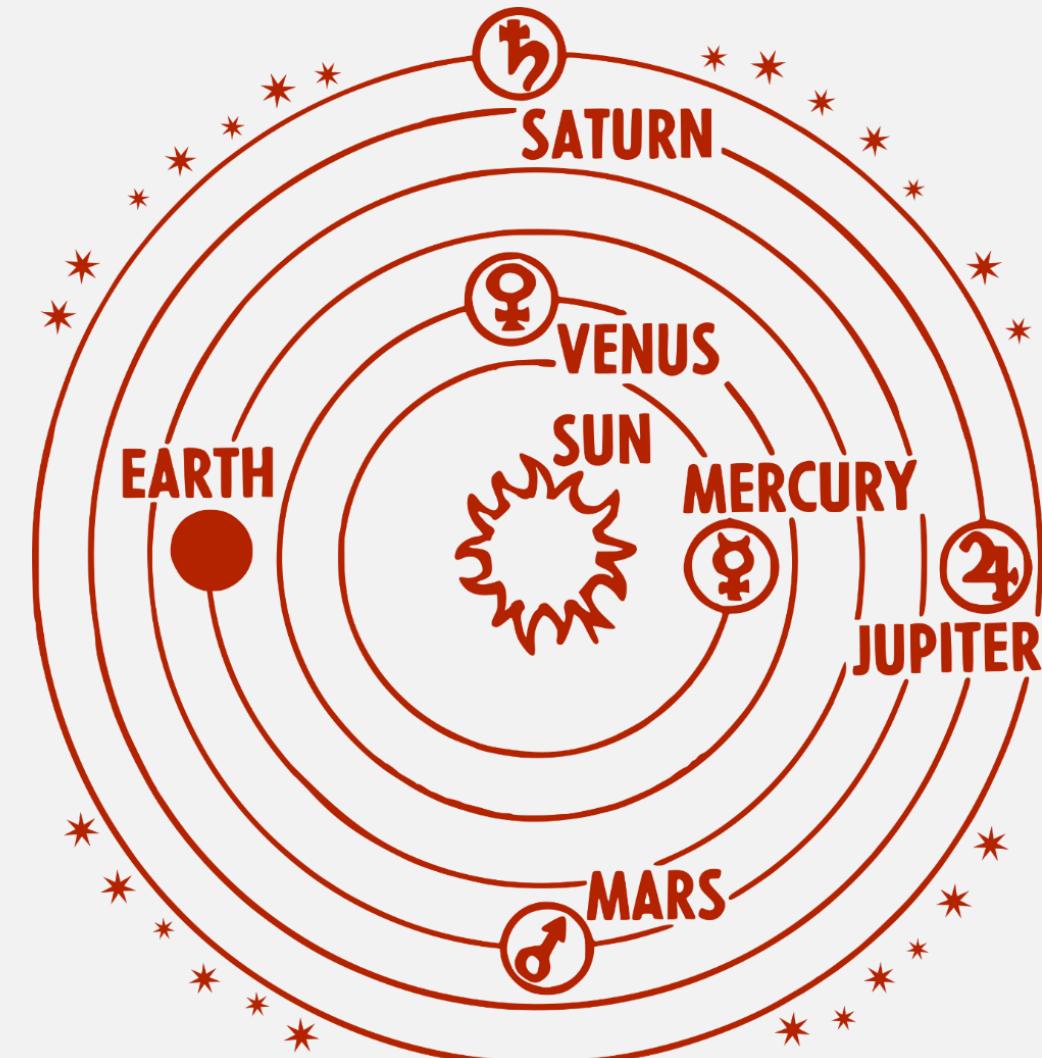
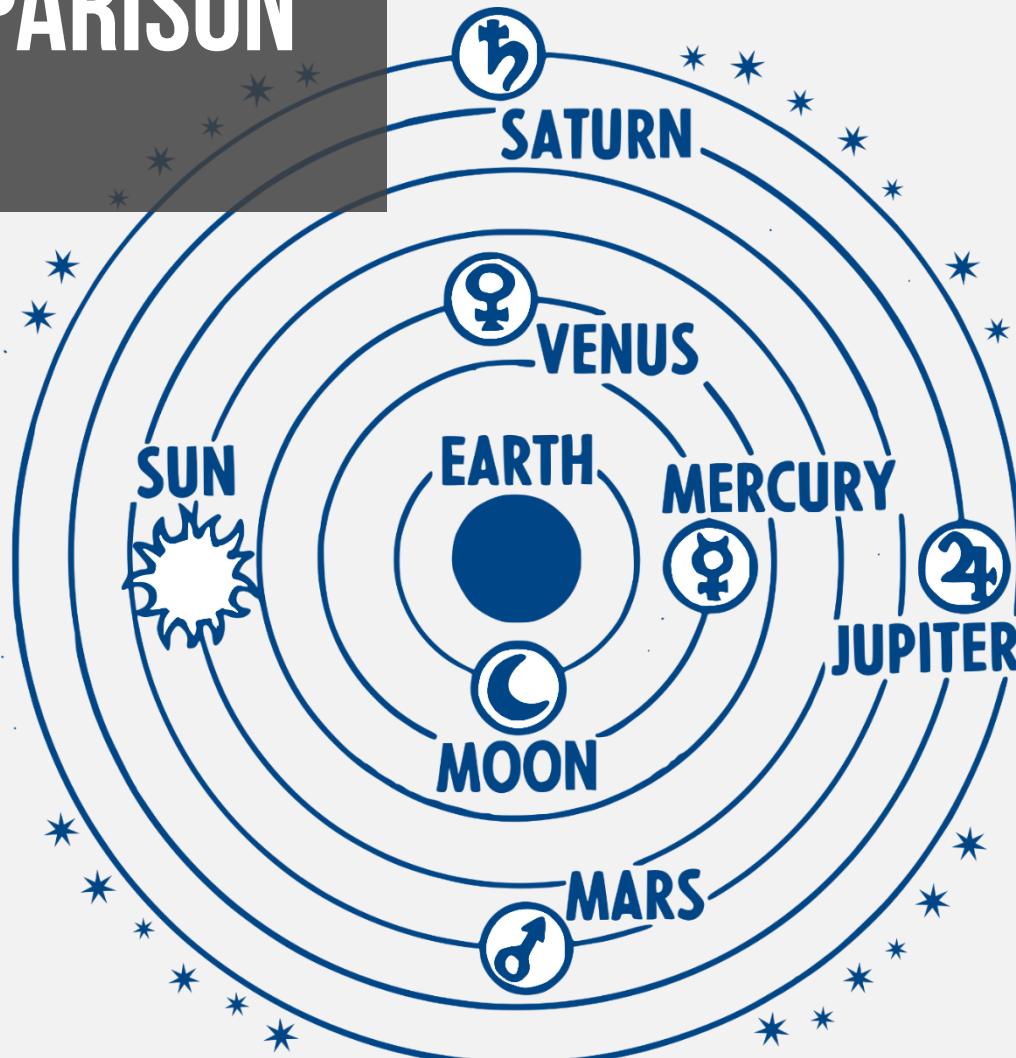


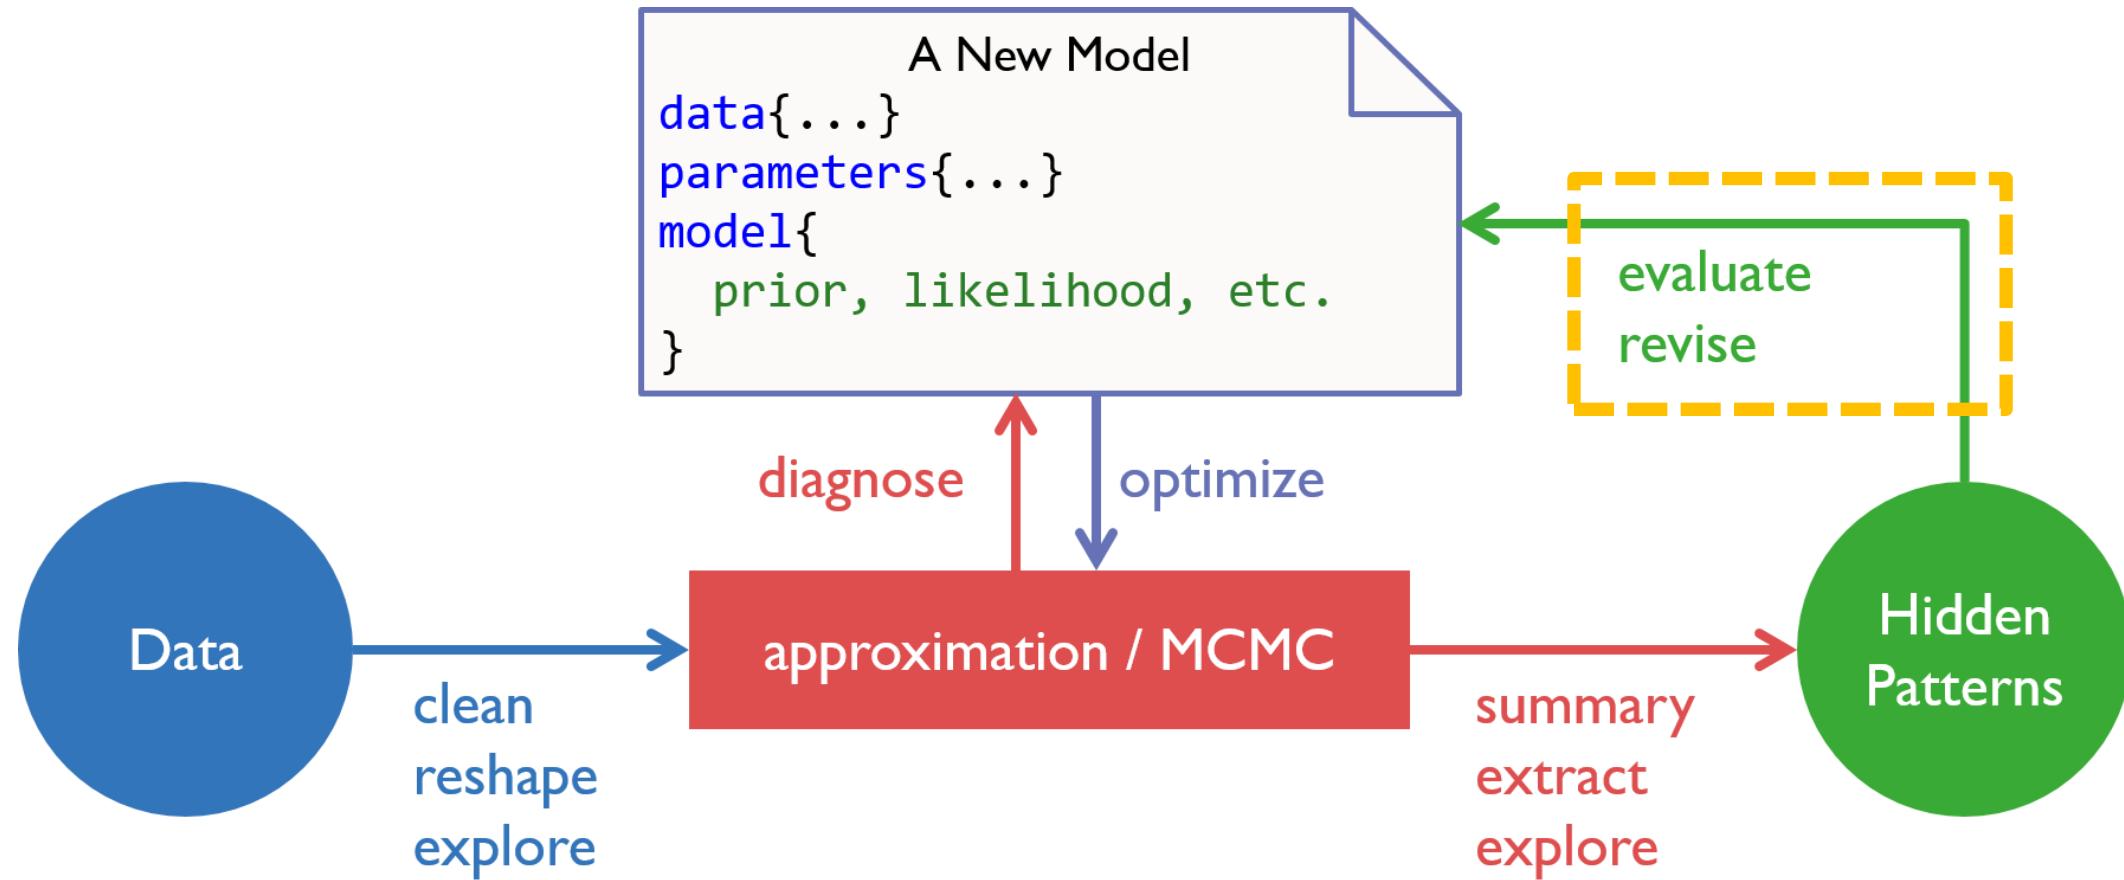
# Posterior Predictive Check

cognitive model  
statistics  
computing



# MODEL COMPARISON





# Model Comparison

cognitive model  
statistics  
computing

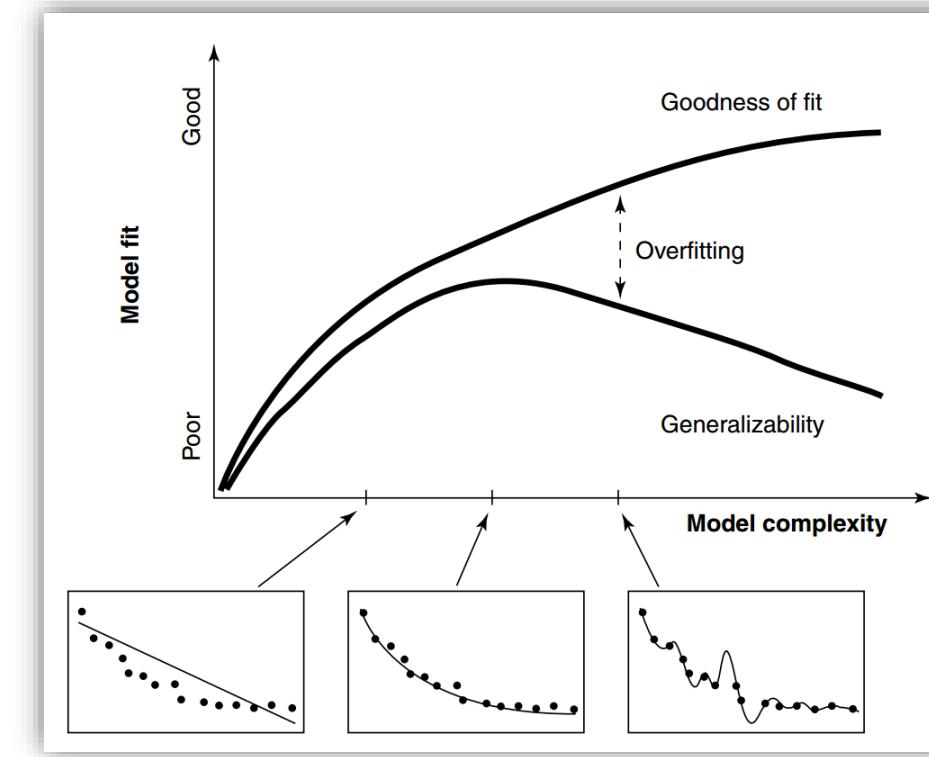
Which model provides the best **fit**?



Which model represents the best **balance** between model fit and model complexity?

Ockham's razor:

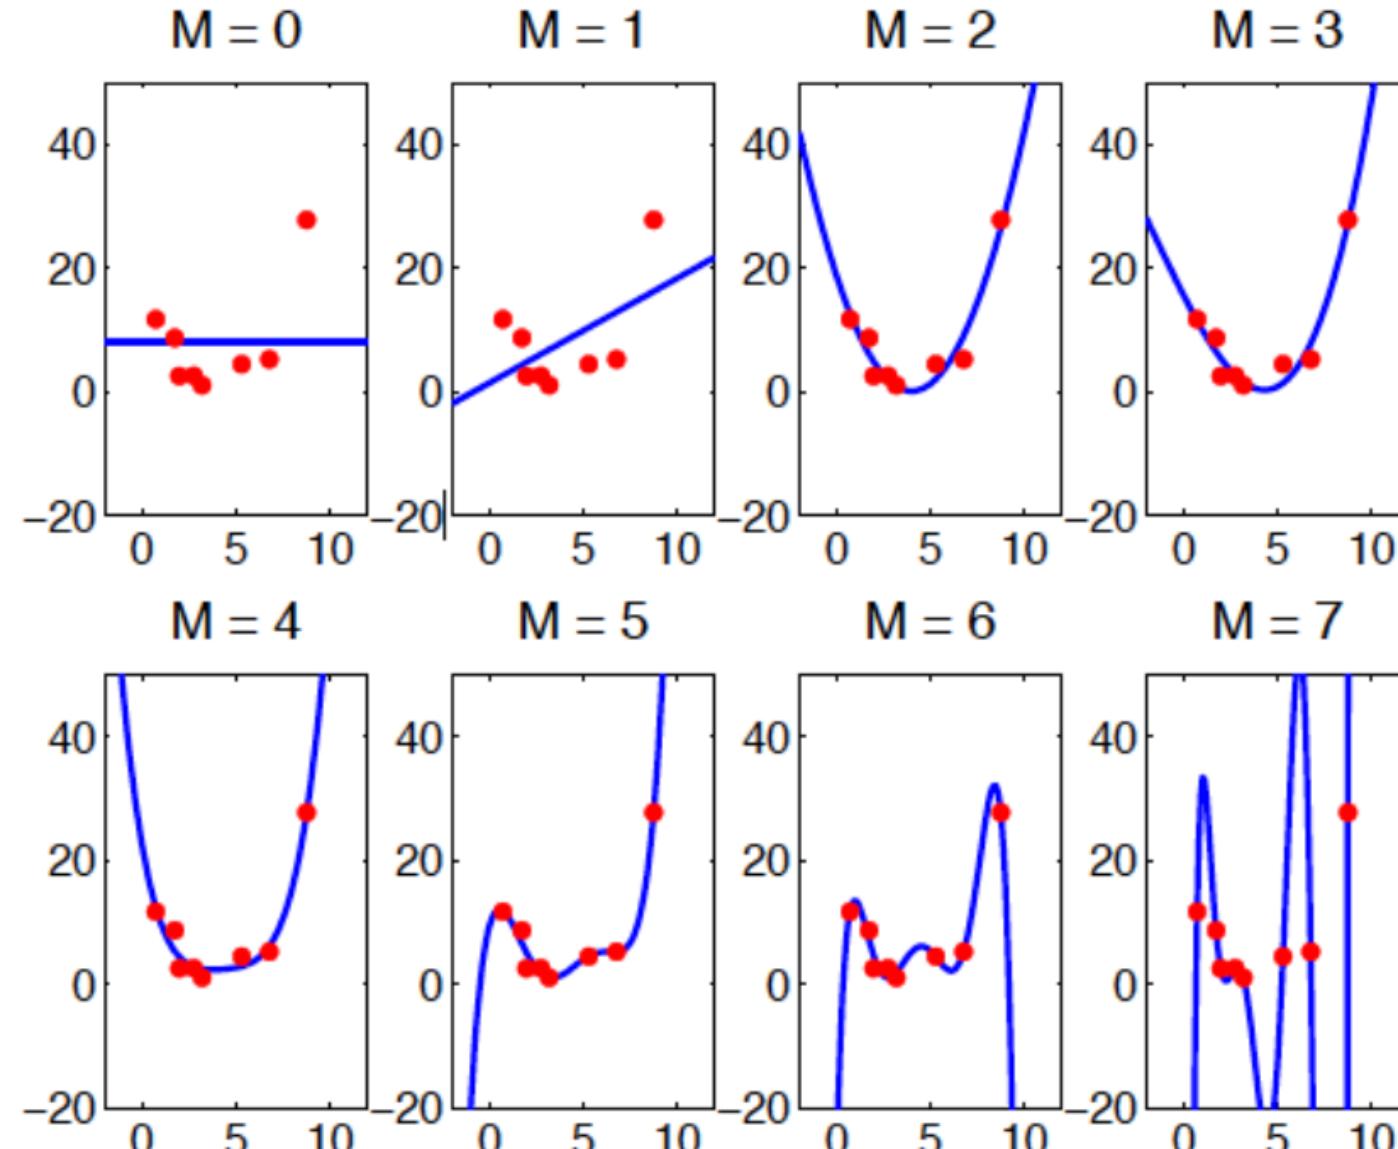
Models with fewer assumptions are to be preferred



- overfitting: learn **too much** from the data
- underfitting: learn **too little** from the data

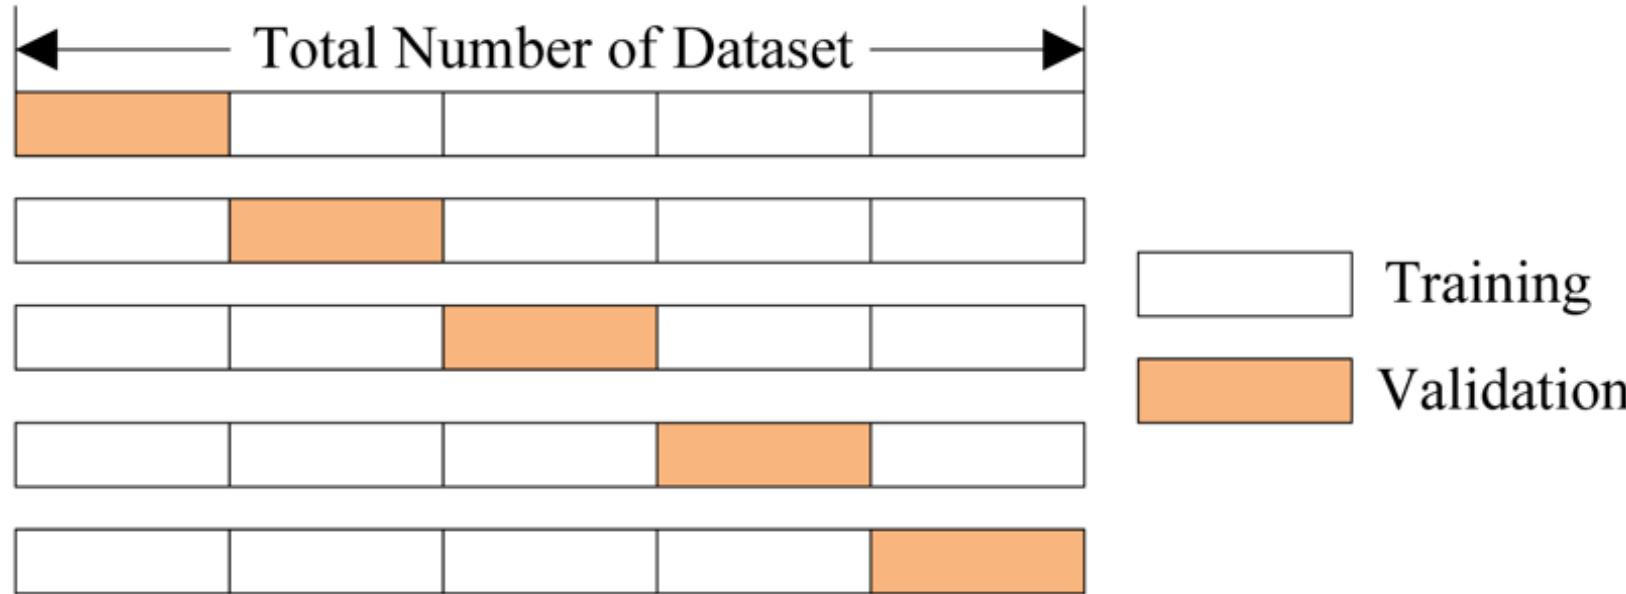
# Which model has the highest predictive power?

cognitive model  
statistics  
computing



# Focusing on Predictive Accuracy

cognitive model
statistics
computing



- Nothing prevents you from doing that in a Bayesian context but holding out data makes your posterior distribution more diffuse
- Bayesians usually condition on *all* the data and evaluate how well a model is expected to **predict out of sample** using "information criteria": model with the **highest expected log predictive density (ELPD)** for new data

# Information Criteria

cognitive model
statistics
computing

AIC – Akaike information criterion

DIC – Deviance Information Criterion

WAIC – Widely Applicable Information Criterion

finding the model that has the highest out-of-sample predictive accuracy

BIC – Bayesian Information Criterion

approximation to LOO

finding the “true” model

# Compute WAIC from Likelihood

cognitive model
statistics
computing

$$\text{WAIC} = -2 \widehat{\text{elpd}}_{\text{waic}}$$

expected log pointwise predictive density

$$\widehat{\text{elpd}}_{\text{waic}} = \widehat{\text{lpd}} - \widehat{p}_{\text{waic}}$$

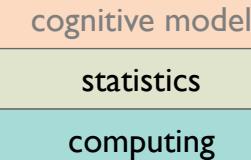
$$\begin{aligned}\widehat{\text{lpd}} &= \text{computed log pointwise predictive density} \\ &= \sum_{i=1}^n \log \left( \frac{1}{S} \sum_{s=1}^S p(y_i | \theta^s) \right).\end{aligned}$$

$$\begin{aligned}\widehat{p}_{\text{waic}} &= \text{estimated effective number of parameters} \\ &= \sum_{i=1}^n V_{s=1}^S (\log p(y_i | \theta^s))\end{aligned}$$

```
lpd <- log(colMeans(exp(log_lik)))
```

```
p_waic <- colVars(log_lik)
```

# \*IC comparisons



		No pooling $(\tau = \infty)$	Complete pooling $(\tau = 0)$	Hierarchical model $(\tau \text{ estimated})$
AIC	$-2 \text{lpd} = -2 \log p(y   \hat{\theta}_{\text{mle}})$	54.6	59.4	
	$k$	8.0	1.0	
	$\text{AIC} = -2 \widehat{\text{elpd}}_{\text{AIC}}$	70.6	61.4	
DIC	$-2 \text{lpd} = -2 \log p(y   \hat{\theta}_{\text{Bayes}})$	54.6	59.4	57.4
	$p_{\text{DIC}}$	8.0	1.0	2.8
	$\text{DIC} = -2 \widehat{\text{elpd}}_{\text{DIC}}$	70.6	61.4	63.0
WAIC	$-2 \text{lppd} = -2 \sum_i \log p_{\text{post}}(y_i)$	60.2	59.8	59.2
	$p_{\text{WAIC } 1}$	2.5	0.6	1.0
	$p_{\text{WAIC } 2}$	4.0	0.7	1.3
	$\text{WAIC} = -2 \widehat{\text{elppd}}_{\text{WAIC } 2}$	68.2	61.2	61.8
LOO-CV	$-2 \text{lppd}$		59.8	59.2
	$p_{\text{loo-cv}}$		0.5	1.8
	$-2 \text{lppd}_{\text{loo-cv}}$		60.8	62.8

# Recording the Log-Likelihood in Stan

cognitive model  
statistics  
computing

```
generated quantities {
  ...
  real log_lik[nSubjects];
  ...

  { # Local section, this saves time and space
    for (s in 1:nSubjects) {
      vector[2] v;
      real pe;

      log_lik[s] = 0;
      v = initV;

      for (t in 1:nTrials) {
        log_lik[s] = log_lik[s] + categorical_logit_lpmf(choice[s,t] | tau[s] * v);

        pe = reward[s,t] - v[choice[s,t]];
        v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;
      }
    }
  }
}
```

# The {loo} Package

cognitive model  
statistics  
computing

```
> library(loo)
> LL1    <- extract_log_lik(stanfit)
> loo1   <- loo(LL1)    # PSIS leave-one-out
> waic1 <- waic(LL1)   # WAIC
```

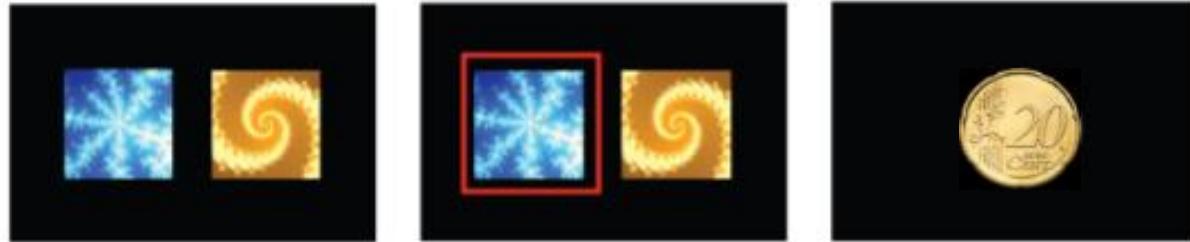
Computed from 4000 by 20 log-likelihood matrix

	Estimate	SE
elpd_loo	-29.5	3.3
p_loo	2.7	1.0
looic	58.9	6.7

Pareto Smoothed Importance Sampling

# Reversal Learning Task

cognitive model  
statistics  
computing



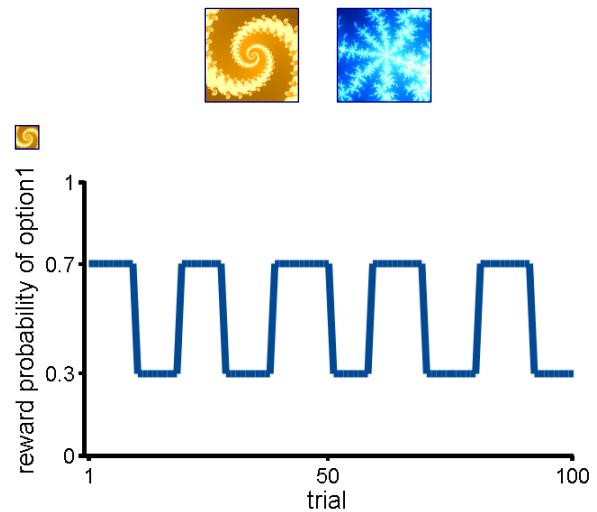
Fictitious RL  
(Counterfactual RL)

Value update:

$$V_{t+1}^c = V_t^c + \alpha * PE$$
$$V_{t+1}^{nc} = V_t^{nc} + \alpha * PEnc$$

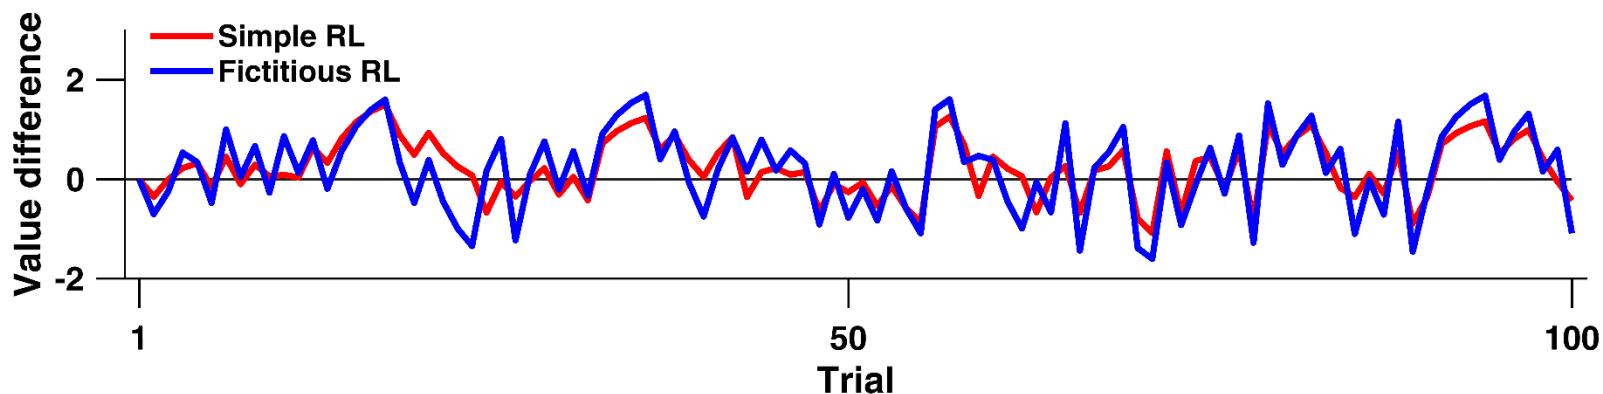
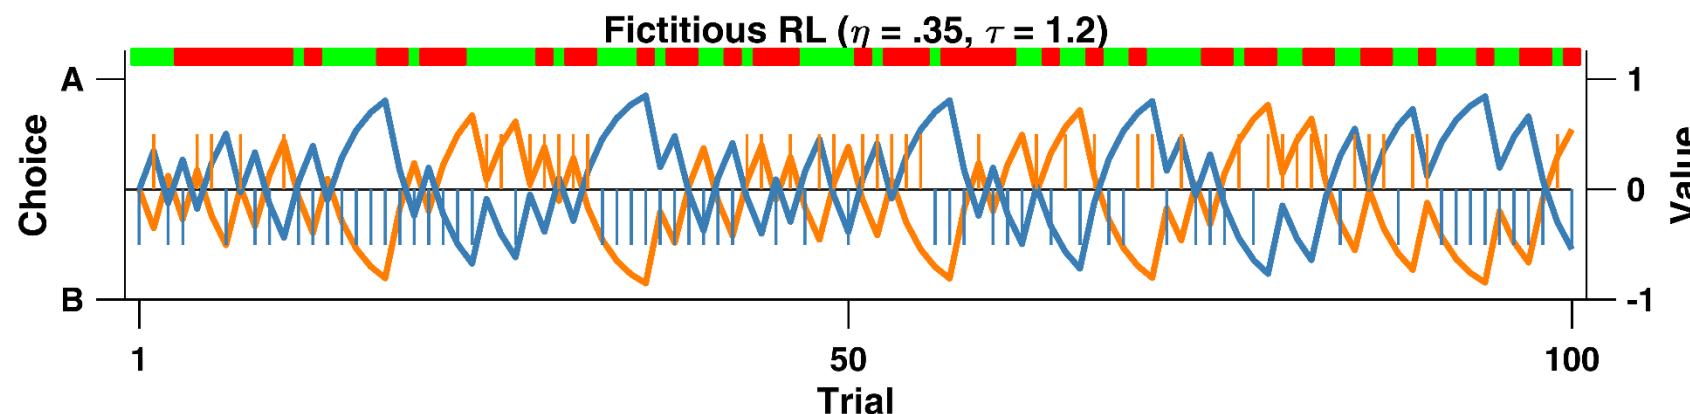
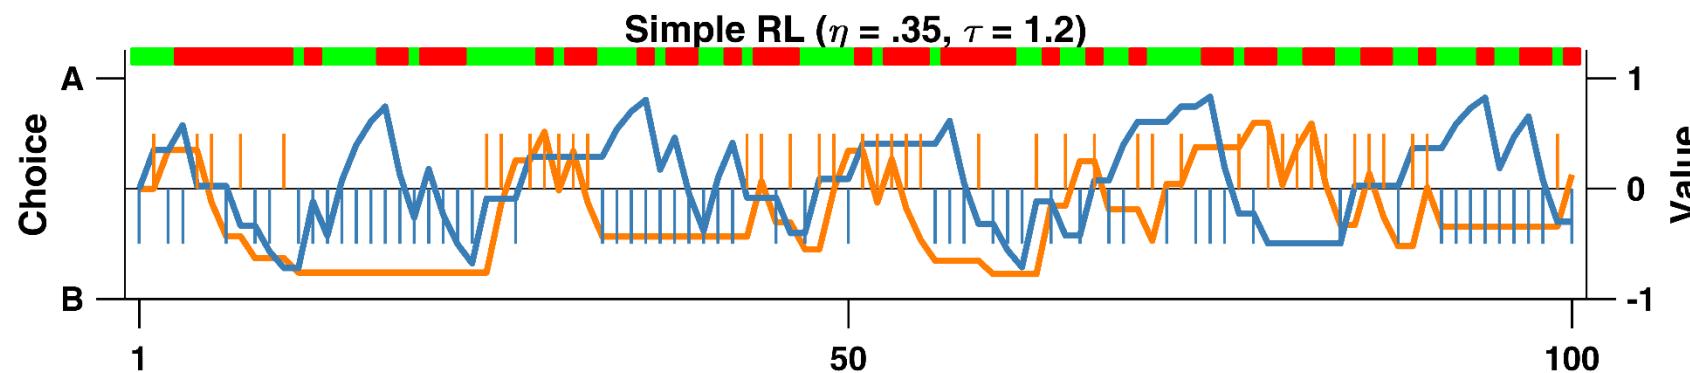
Prediction error:

$$PE = R_t - V_t^c$$
$$PEnc = -R_t - V_t^{nc}$$



# More on Fictitious RL

cognitive model  
statistics  
computing



# Exercise XIII

cognitive model  
statistics  
computing

.../08.compare\_models/\_scripts/compare\_models\_main.R

TASK: (1) complete the fictitious RL model (model2, loglik)  
(2) fit and compare the 2 models

## Exercise XIII – output

```
> LL1 <- extract_log_lik(fit_r1)
> ( loo1 <- loo(LL1) )
Computed from 4000 by 10 log-likelihood matrix
```

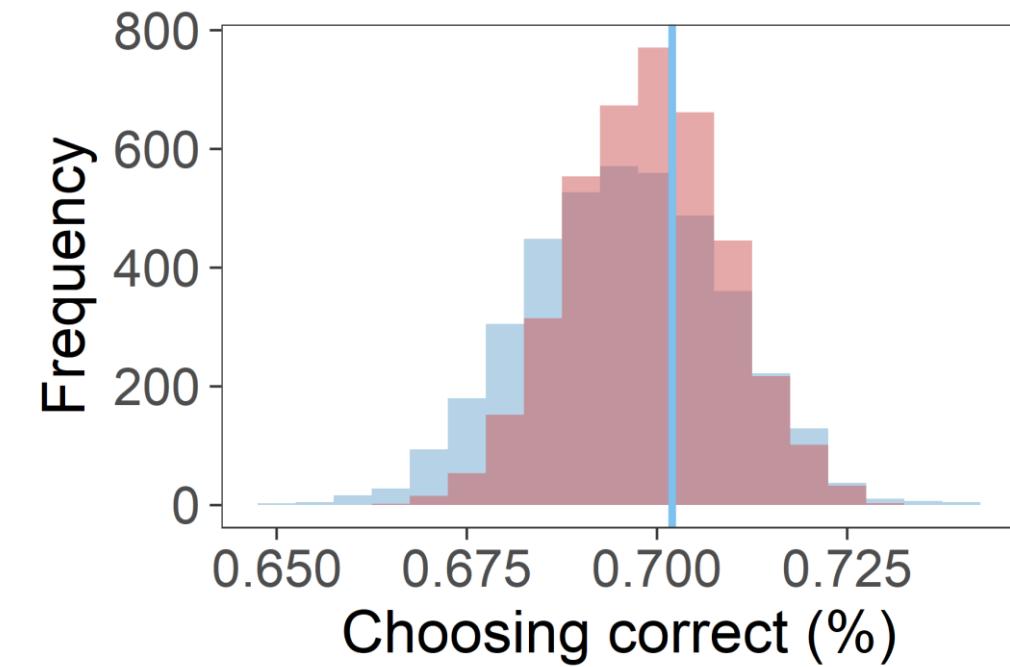
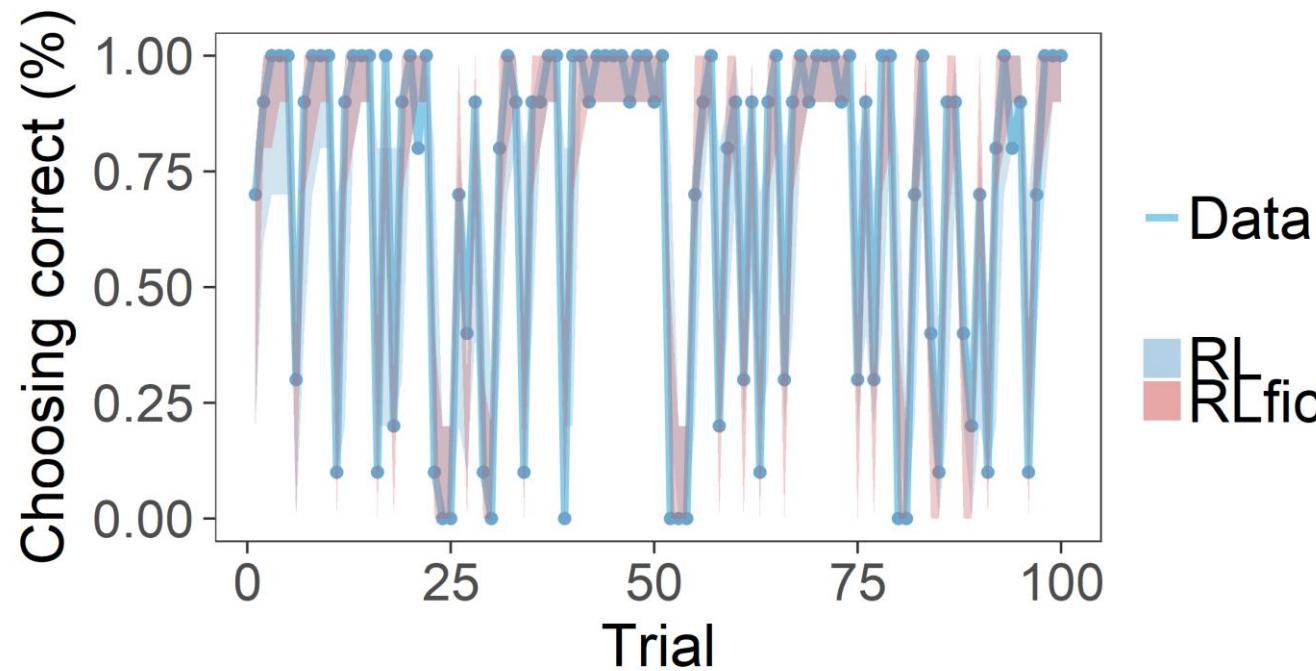
	Estimate	SE
elpd_loo	-389.8	15.4
p_loo	3.8	0.8
looic	779.5	30.9

```
> ( loo2 <- loo(LL2) )
Computed from 4000 by 10 log-likelihood matrix
```

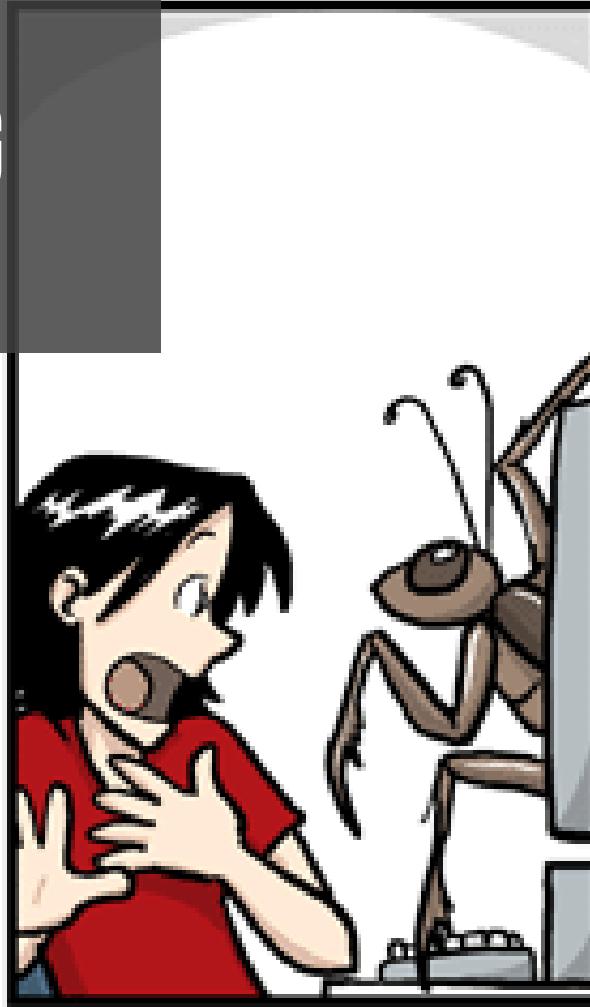
	Estimate	SE
elpd_loo	-281.3	17.5
p_loo	3.4	0.5
looic	562.6	35.0

# Posterior Predictive Check

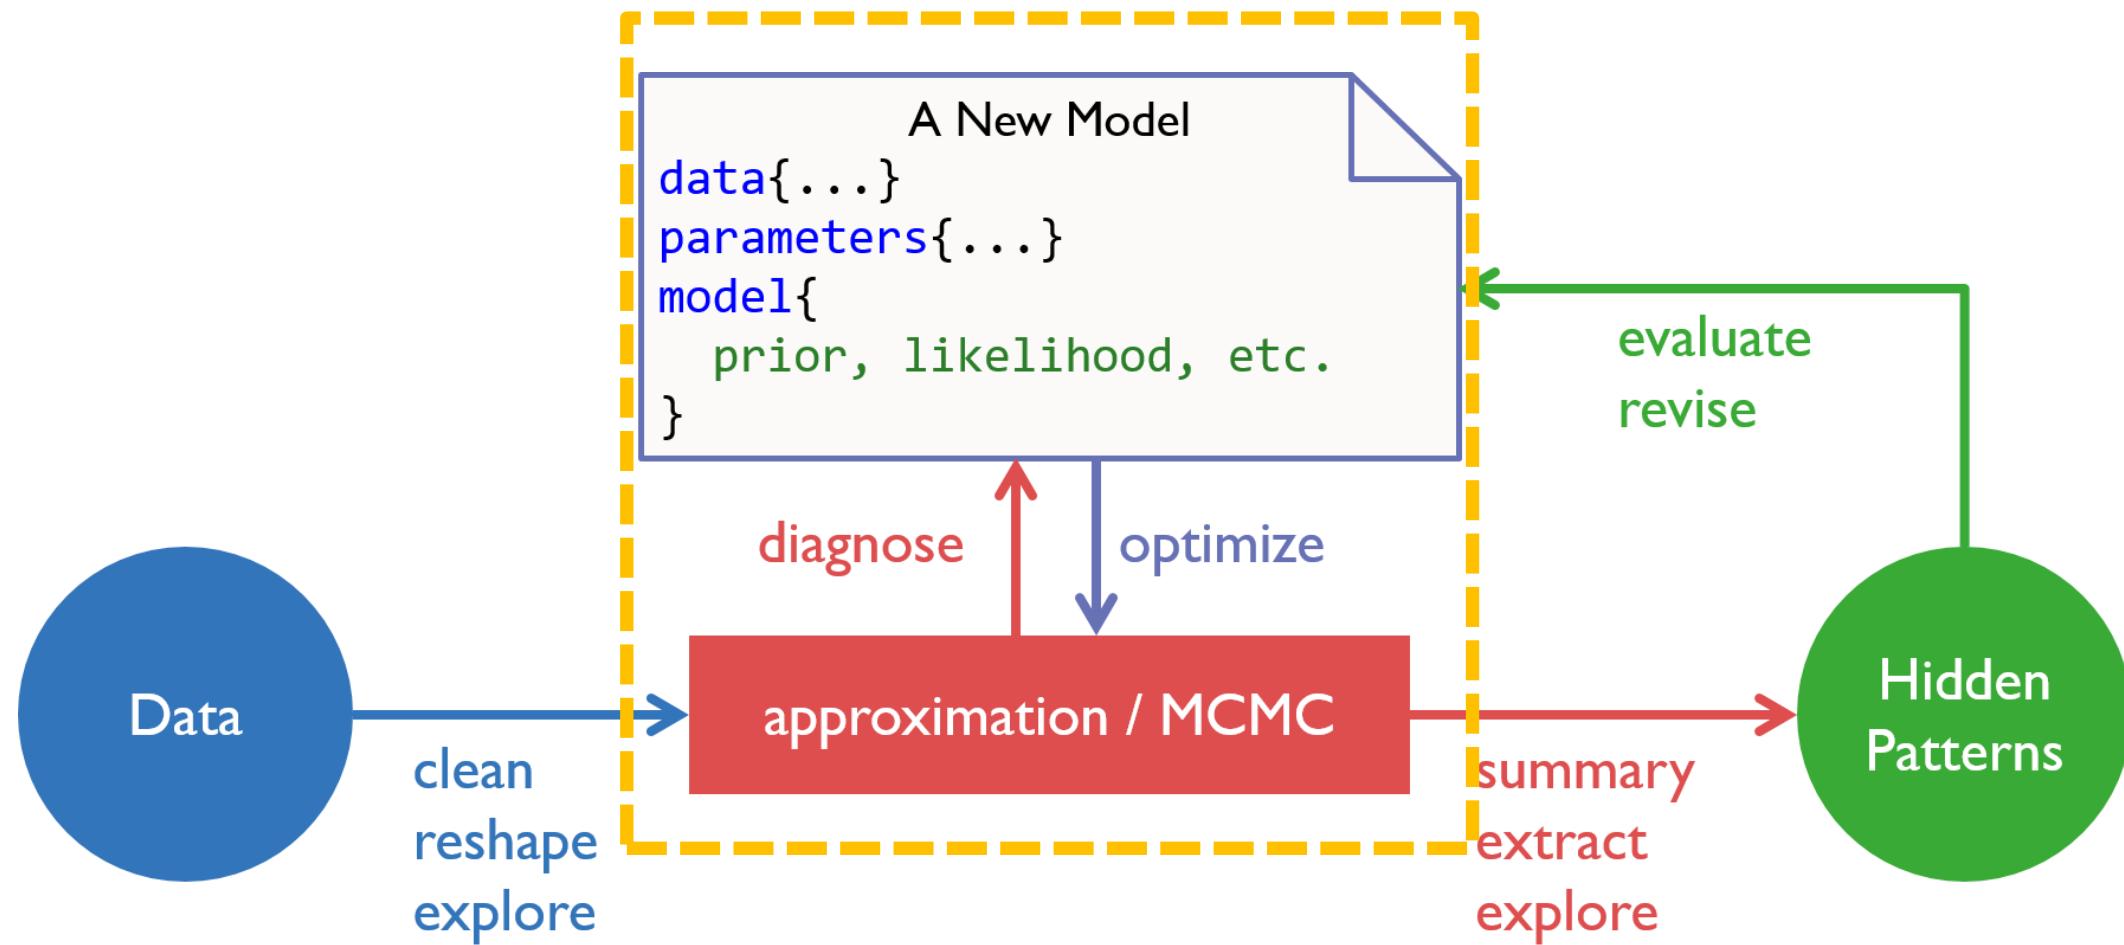
cognitive model  
statistics  
computing



# STAN DEBUGGING



JORGE CHAM © 2005



# Stan Style Tips

cognitive model  
statistics  
computing

## Make it Reproducible

- Scripts are good documentations!
- Save your seed (not cross platform\*)

## Make it Readable

- Choose a consistent style
- Give meaningful variable names

## Start with Simulated Data

## Design Top-Down, Code Bottom-Up

## Write Comments

- Code never lies!



\* Stan seed depends on hardware etc.

# The Editor of your Choice

cognitive model  
statistics  
computing



```
data {  
  int<lower=0> w;  
  int<lower=0> N;  
}  
  
parameters {  
  real<lower=0,upper=1> p;  
}  
  
model {  
  p ~ uniform(0,1);  
  w ~ binomial(N, p);  
}
```



```
data {  
  int<lower=0> w;  
  int<lower=0> N;  
}  
  
parameters {  
  real<lower=0,upper=1> p;  
}  
  
model {  
  p ~ uniform(0,1);  
  w ~ binomial(N, p);  
}
```



```
data {  
  int<lower=0> w;  
  int<lower=0> N;  
}  
  
parameters {  
  real<lower=0,upper=1> p;  
}  
  
model {  
  p ~ uniform(0,1);  
  w ~ binomial(N, p);  
}
```



```
data {  
  int<lower=0> w;  
  int<lower=0> N;  
}  
  
parameters {  
  real<lower=0,upper=1> p;  
}  
  
model {  
  p ~ uniform(0,1);  
  w ~ binomial(N, p);  
}
```

\* Click on each logo to visit their homepage.

\*\* [Comparison](#)

# Common Error / Warning Types

cognitive model  
statistics  
computing

## ERRORS

forget “ ; ”  
mis-indexing: mismatch or constant  
support mismatch  
improper constrain  
improper dimension declaration  
vectorizing when not supported  
wrong data type  
wrong distribution names  
forget priors  
miss spelling

## WARNINGS

forget last blank line  
use earlier version of Stan  
numerical problems  
divergent transitions  
hit max\_treedepth  
BFI too low  
improper prior

see also: [Brief Guide to Stan's Warnings](#)

# Debugging in Stan

- always use a \*.stan file
- press  in RStudio
- use `lookup()`
- start with simulated data
- be careful with copy/paste
- run 1 chain, 1 sample
- debugging by printing

```
for (s in 1:1) {
    vector[2] v;
    real pe;
    v <- initV;

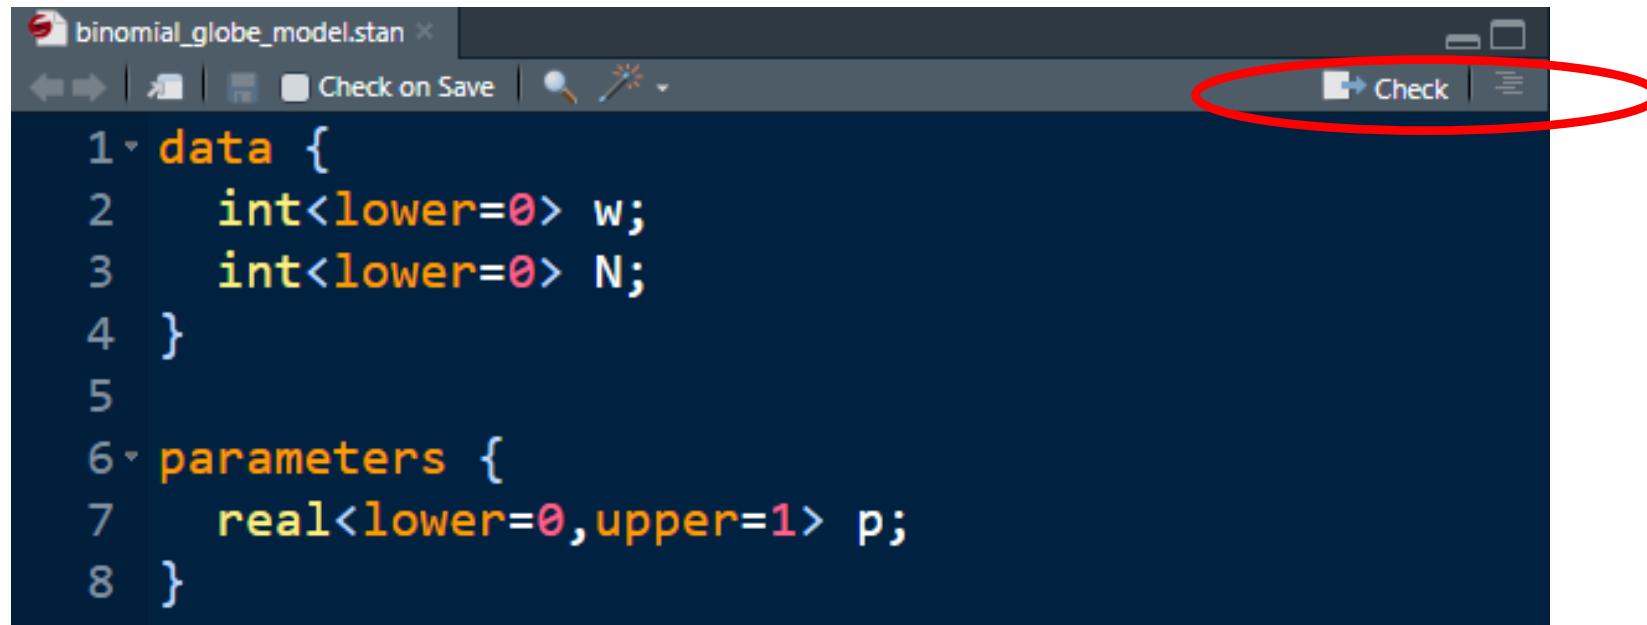
    for (t in 1:nTrials) {
        choice[s,t] ~ categorical_logit( tau[s] * v );

        print("s = ", s, ", t = ", t, ", v = ", v);

        pe <- reward[s,t] - v[choice[s,t]];
        v[choice[s,t]] <- v[choice[s,t]] + lr[s] * pe;
    }
}
```

```
> lookup(dnorm)
      StanFunction          Arguments ReturnType Page SamplingStatement
344     normal           (reals mu, reals sigma)    real   369        TRUE
348     normal_log (reals y, reals mu, reals sigma)    real   369       FALSE
```

# Debugging Stan in RStudio



The screenshot shows the RStudio interface with a Stan script named "binomial\_globe\_model.stan" open. The code is as follows:

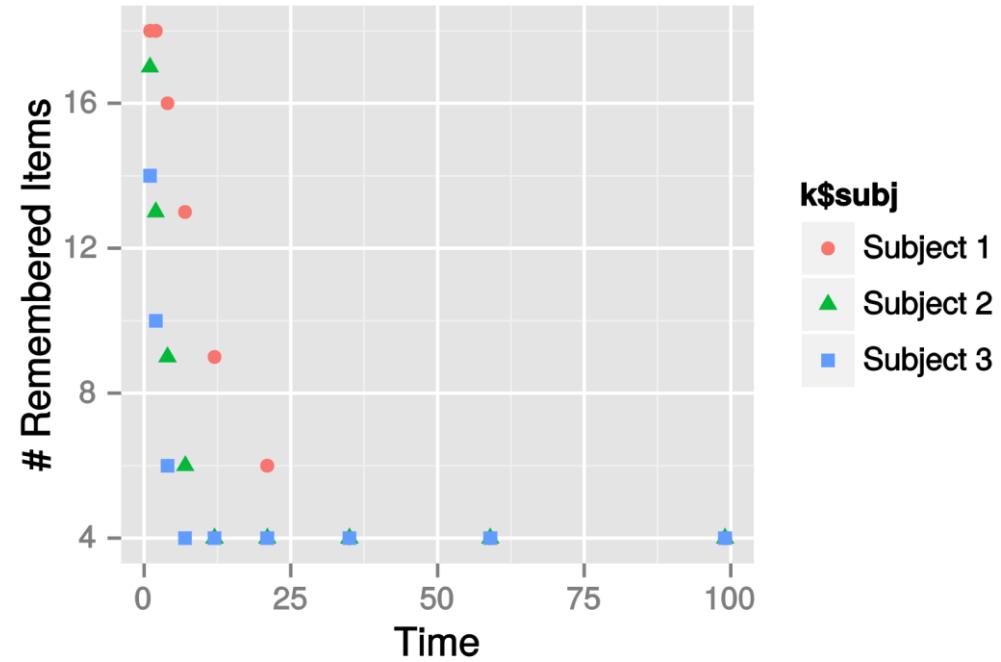
```
1 data {  
2     int<lower=0> w;  
3     int<lower=0> N;  
4 }  
5  
6 parameters {  
7     real<lower=0,upper=1> p;  
8 }
```

The toolbar at the top right includes a "Check" button, which is circled in red to indicate it's the focus of the discussion.

```
rstan::rstudio_stanc("_scripts/binomial_globe_model.stan")
```

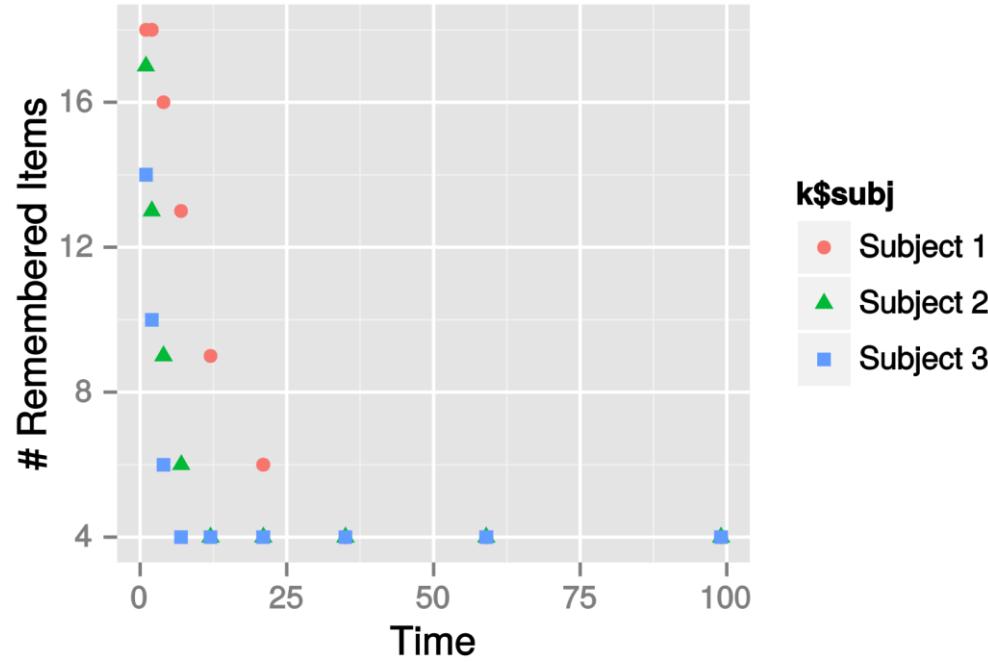


# Example: Memory Retention



Subject	Time Interval								
	1	2	4	7	12	21	35	59	99
1	18	18	16	13	9	6	4	4	4
2	17	13	9	6	4	4	4	4	4
3	14	10	6	4	4	4	4	4	4

# Simple Exponential Decay Model



$$\theta_t = \min(1.0, \exp(-\alpha t) + \beta)$$

↓  
 $p(\text{remember})$        $\text{decay rate}$        $\text{baseline}$

# Exercise XIV

.../09.debugging/\_scripts/exp\_decay\_main.R

**TASK: Debugging the Memory retention model**

**>= 9 errors!**

**Viel Spaß!**

```
> dataList
$`k`
  [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]
[1,]  18   18   16   13    9    6    4    4    4
[2,]  17   13    9    6    4    4    4    4    4
[3,]  14   10    6    4    4    4    4    4    4

$nItem
[1] 18

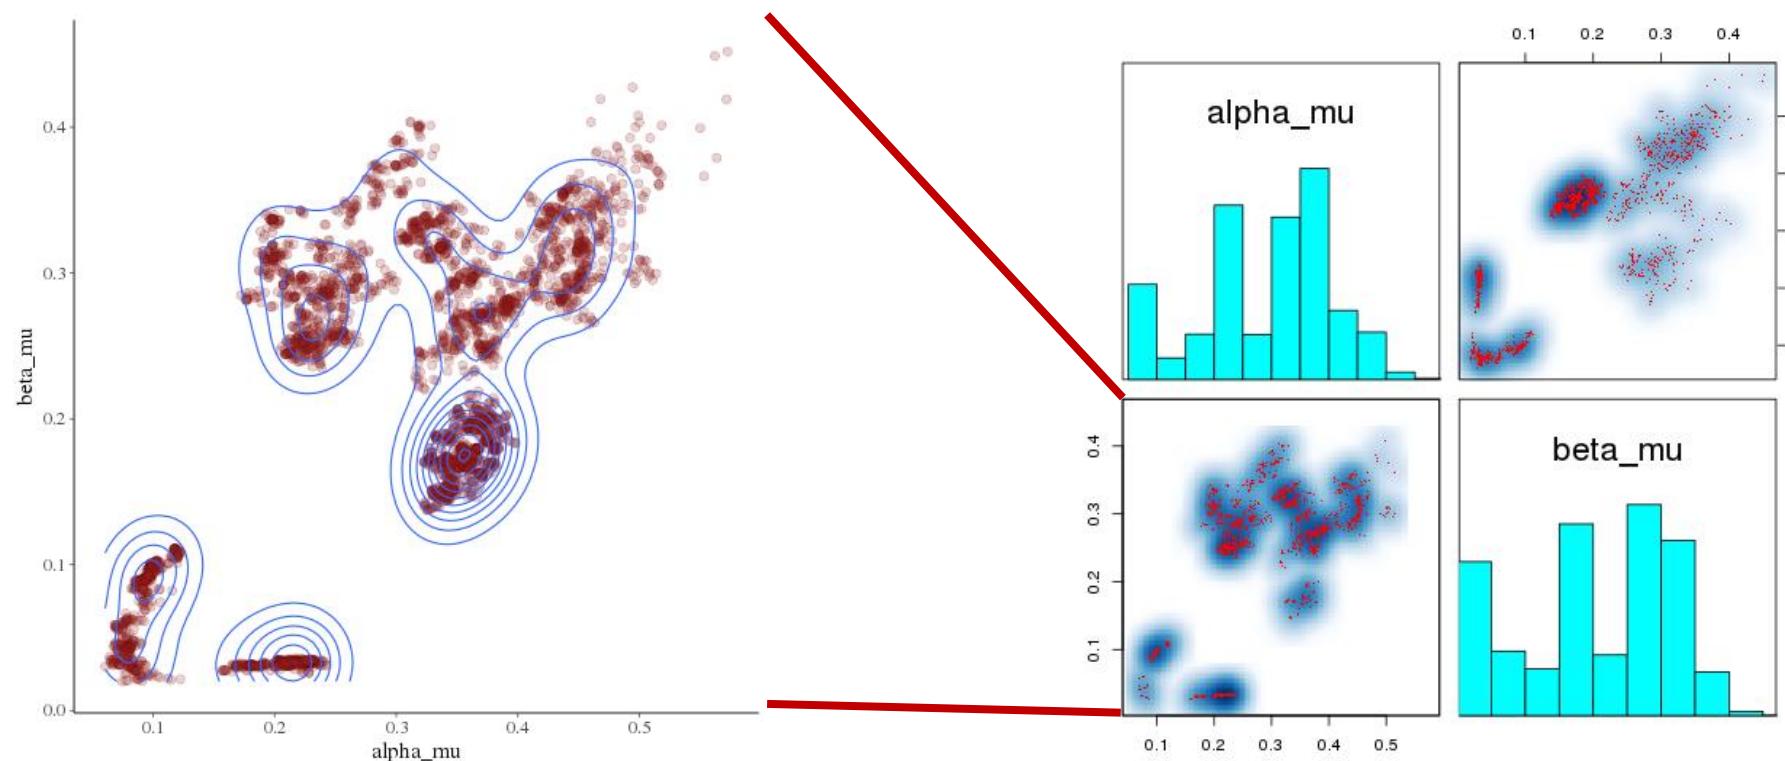
$intervals
[1] 1 2 4 7 12 21 35 59 99

$ns
[1] 3

$nt
[1] 9
```

# Satisfied with the results?

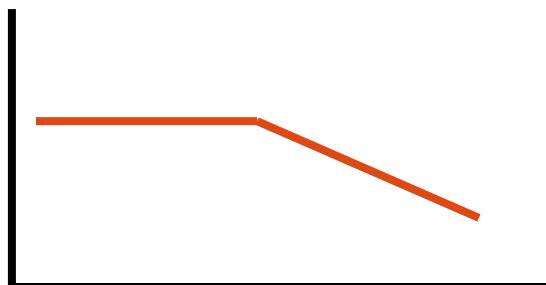
```
Warning messages:  
1: There were 3998 divergent transitions after warmup. Increasing adapt_delta above 0.8 may help. See  
http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup  
2: Examine the pairs() plot to diagnose sampling problems
```



# Why Stan Fails?

```
for (s in 1:ns) {  
  for (t in 1:nt) {  
    theta[s,t] = fmin(1.0, exp(-alpha[s] * intervals[t]) + beta[s]);  
    k[s,t] ~ binomial(nItem, theta[s,t]);  
  }  
}
```

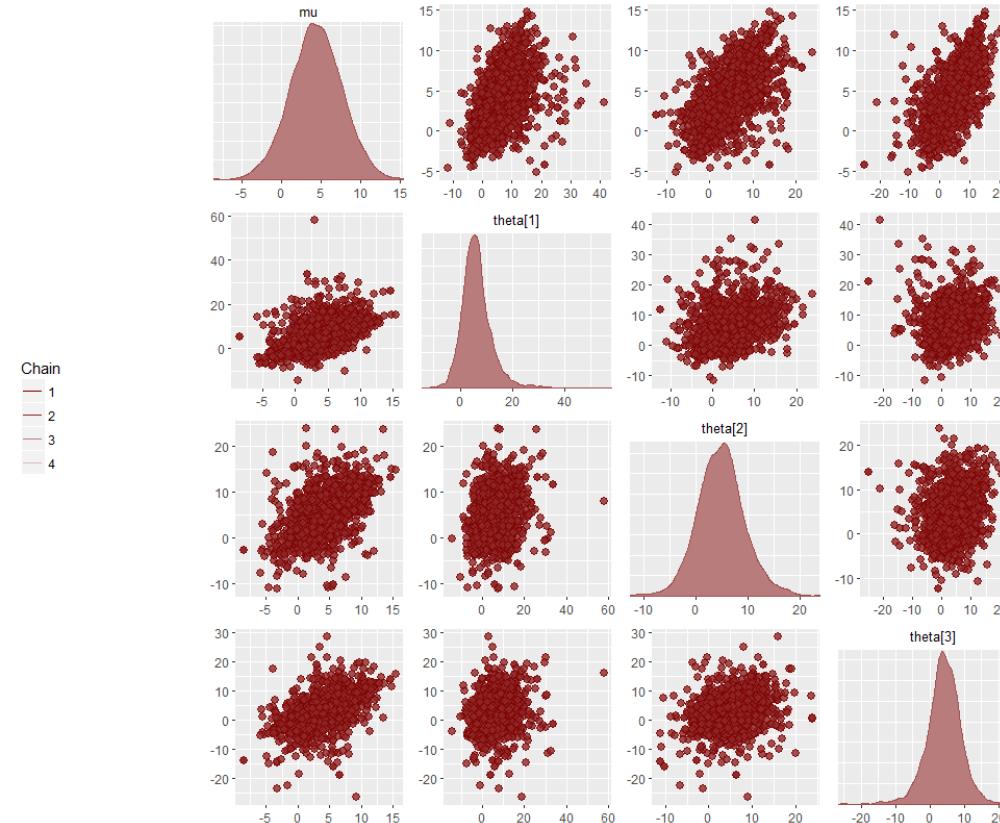
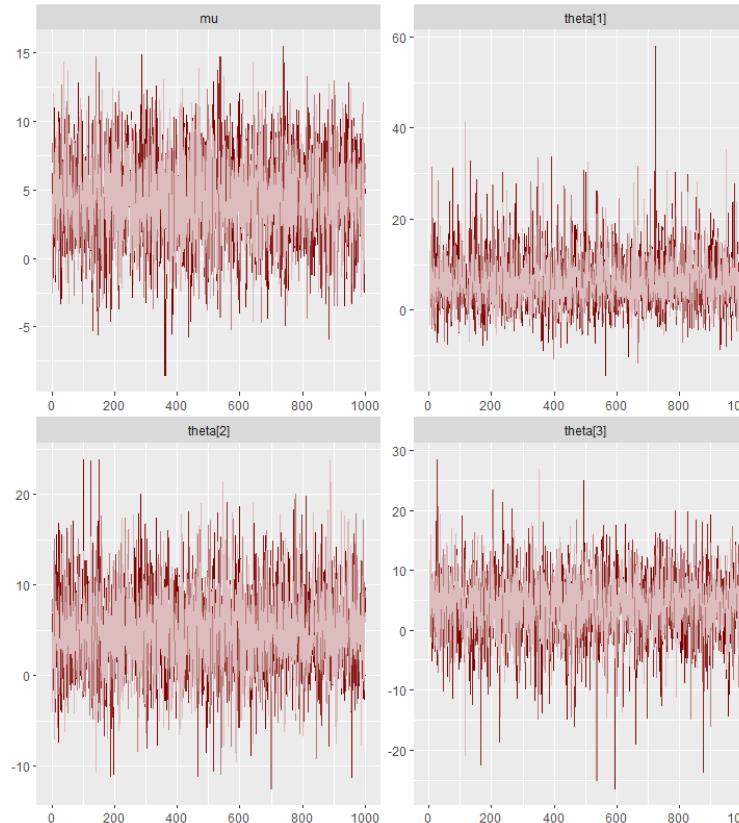
**Non-differentiable** link (likelihood) functions are bad news,  
particularly in Stan, which relies on derivatives.



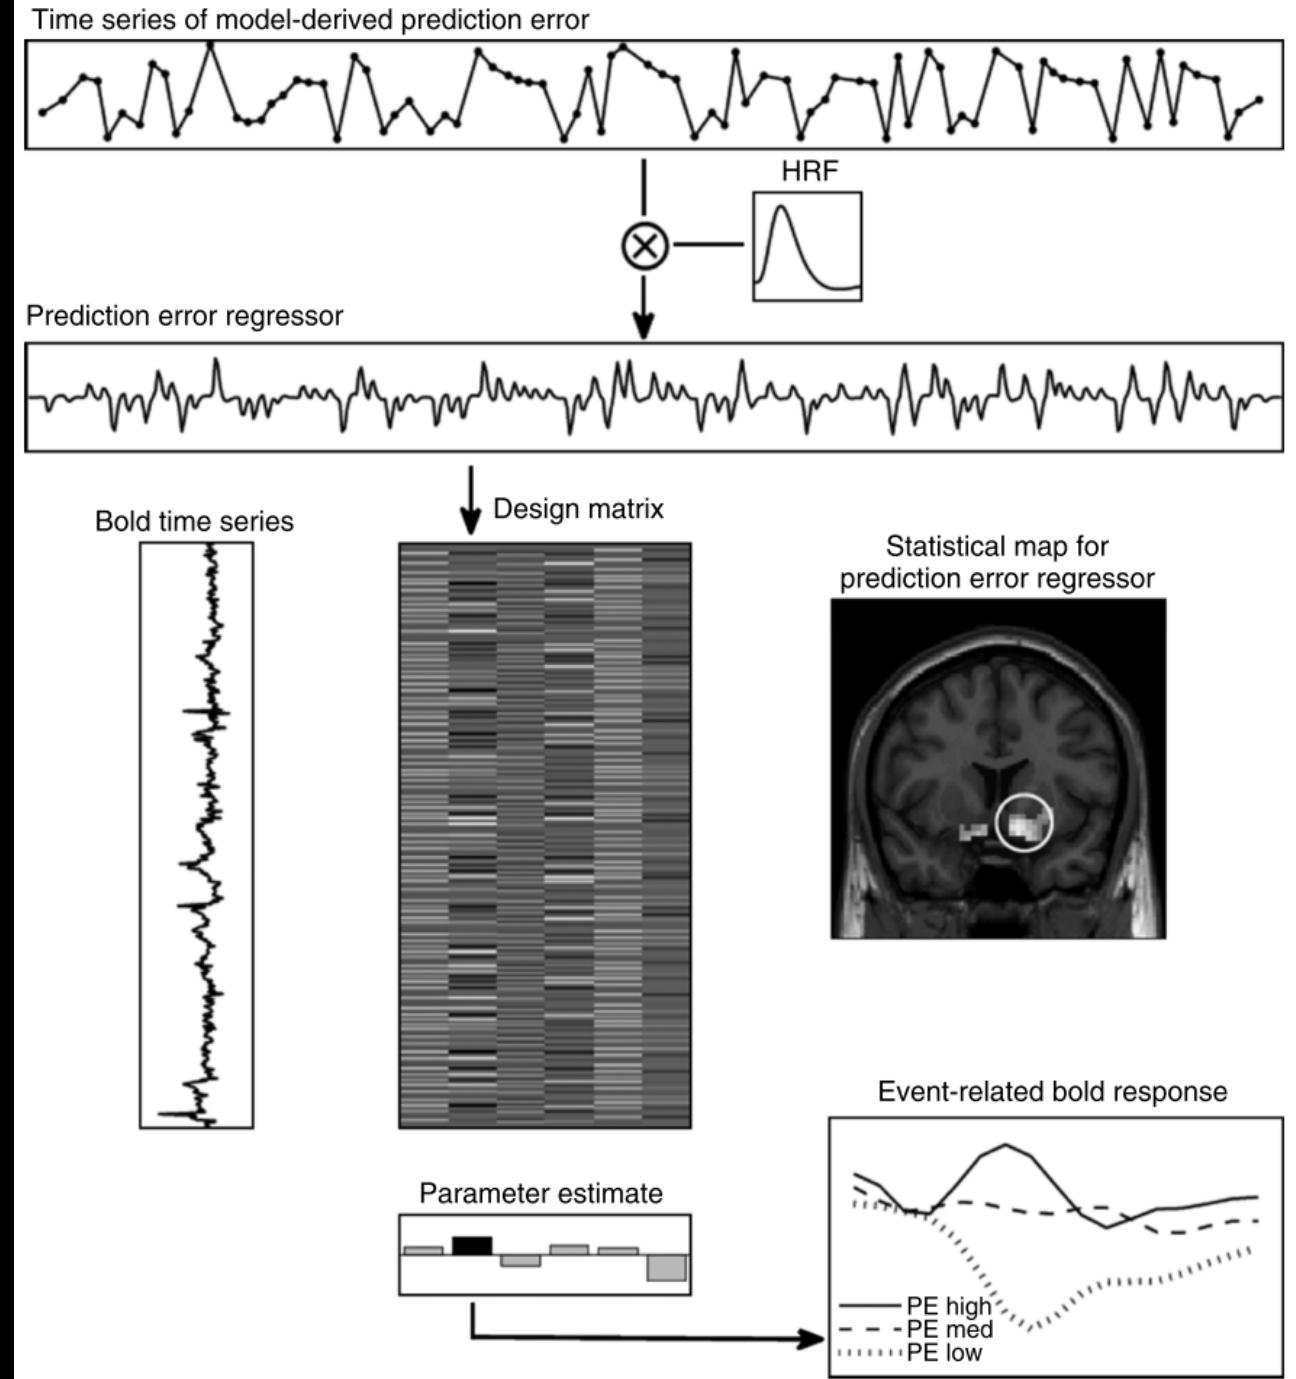
# What to look for?

cognitive model  
statistics  
computing

```
> source('stan_utility.R')
> check_all_diagnostics(fit)
[1] "n_eff / iter looks reasonable for all parameters"
[1] "Rhat looks reasonable for all parameters"
[1] "0 of 4000 iterations ended with a divergence (0%)"
[1] "0 of 4000 iterations saturated the maximum tree depth of 10 (0%)"
[1] "E-BFMI indicated no pathological behavior"
```

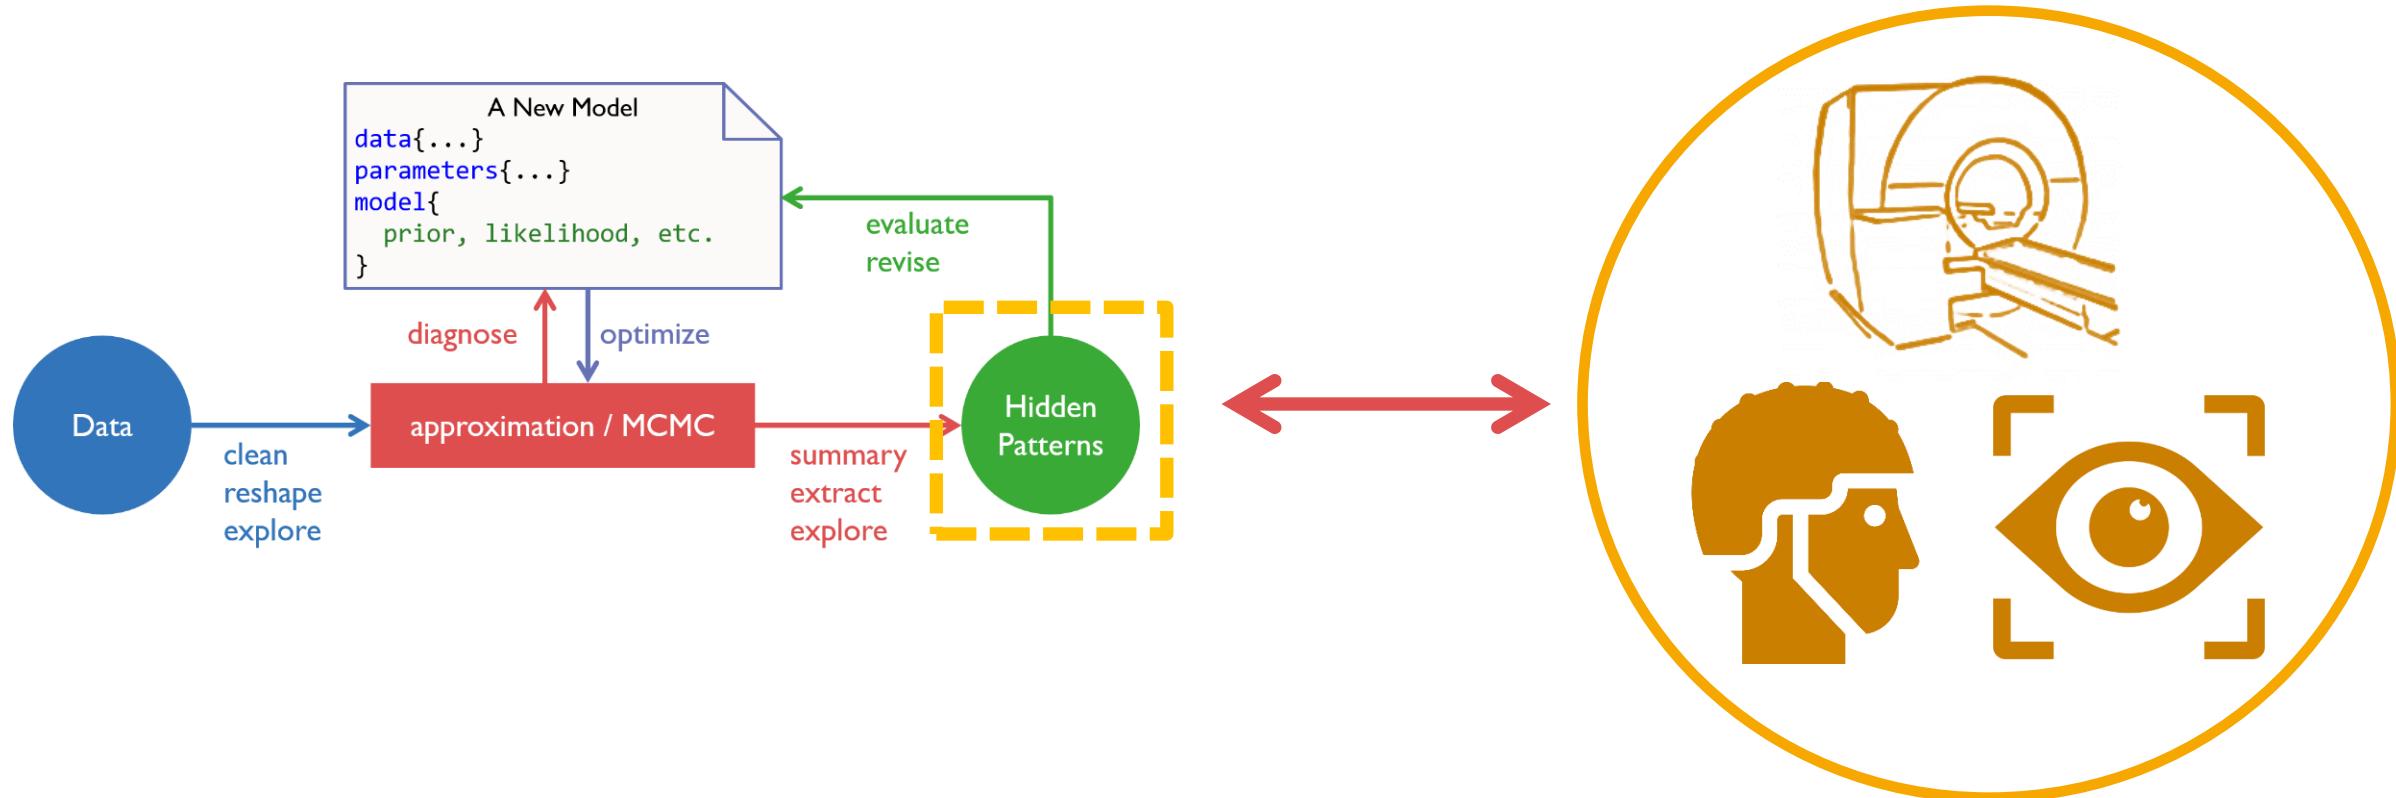


# INTRODUCTION TO MODEL-BASED FMRI



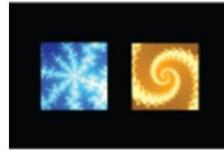
# Model-based Analysis

cognitive model  
statistics  
computing

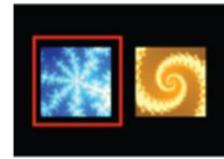


# Perform Model-based fMRI

cognitive model  
statistics  
computing



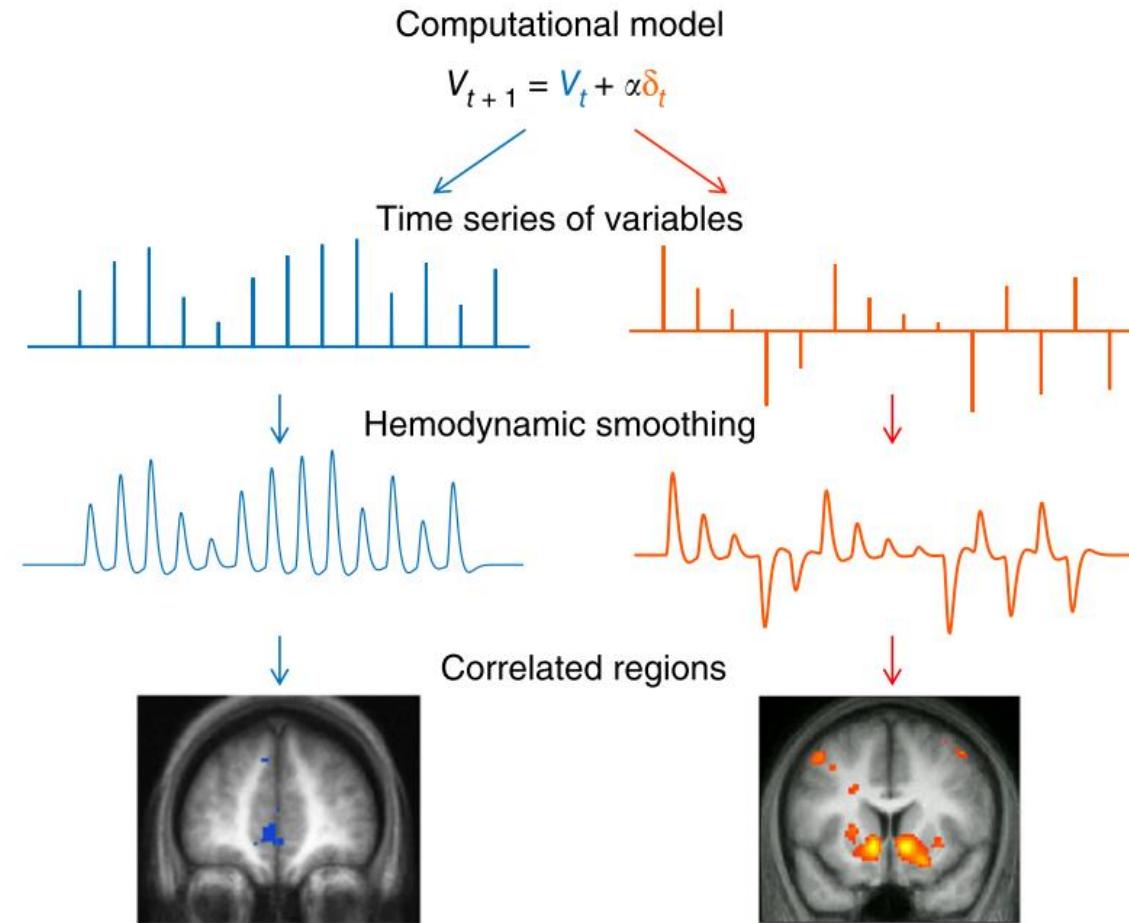
choice presentation



action selection

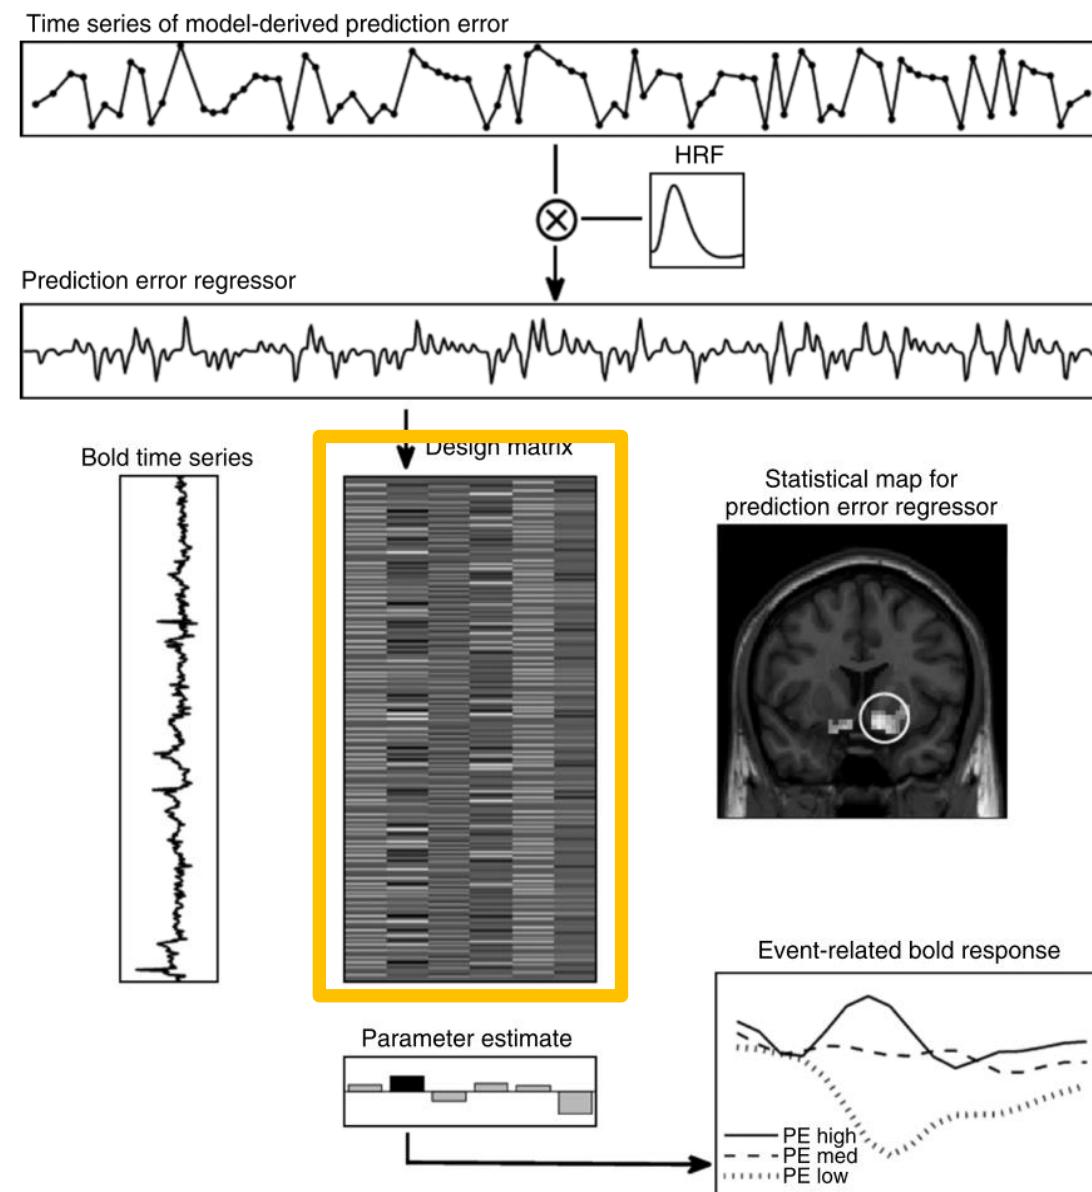


outcome



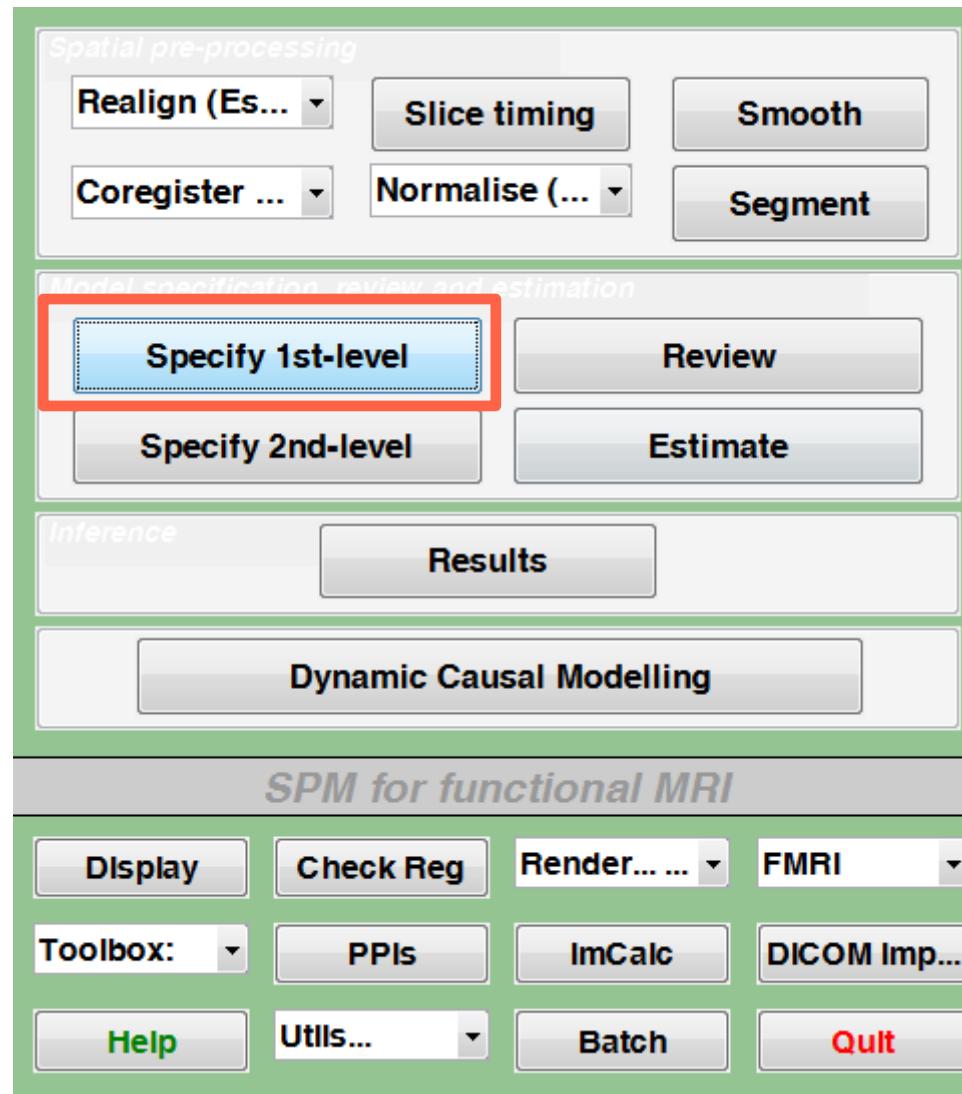
# Perform Model-based fMRI (cont.)

cognitive model  
statistics  
computing



# Implementing in SPM12

cognitive model  
statistics  
computing



The fMRI model specification dialog box shows the following settings:

- fMRI model specific:** Microtime resolution (16), Microtime onset (8).
- Data & Design:** Subject/Session, Scans, Conditions, Condition Name, Onsets, Durations, Time Modulation (No Time Modulation).
- Parametric Modulations:** (highlighted with a red box)
  - Parameter, Name, Values, Polynomial Expansion, Orthogonalise modulations (Yes).
  - multiple conditions, Regressors, Multiple regressors, High-pass filter (128).
- Factorial design:** Basis Functions, Canonical HRF, Model derivatives (No derivatives).
- Current Item:** Parametric Modulations.
- New:** Parameter, Replicate: Parameter (1), Delete: Parameter (1).

# SPM12 – batch scripting

cognitive model  
statistics  
computing

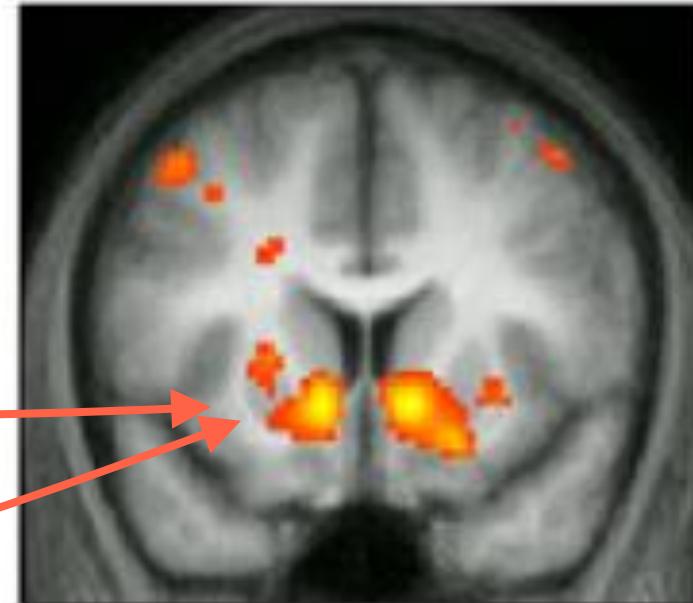
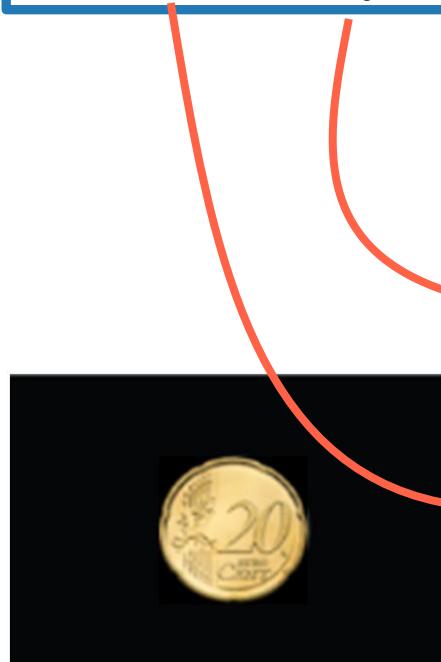
```
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).name = 'onsetPE';
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).onset = onscat.sub(i_sub).cueoutcome;
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).duration = 0;
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).tmod = 0;
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).pmod.name = 'PE';
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).pmod.param = pe(i_sub);
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).pmod.poly = 1;
matlabbatch{1}.spm.stats.fmri_spec.sess.cond(cnt).orth = 0;
```

**make sure:** length(onset) == length(PE)

# A closer look at PE

Prediction error:

$$PE = R_t - V_t$$

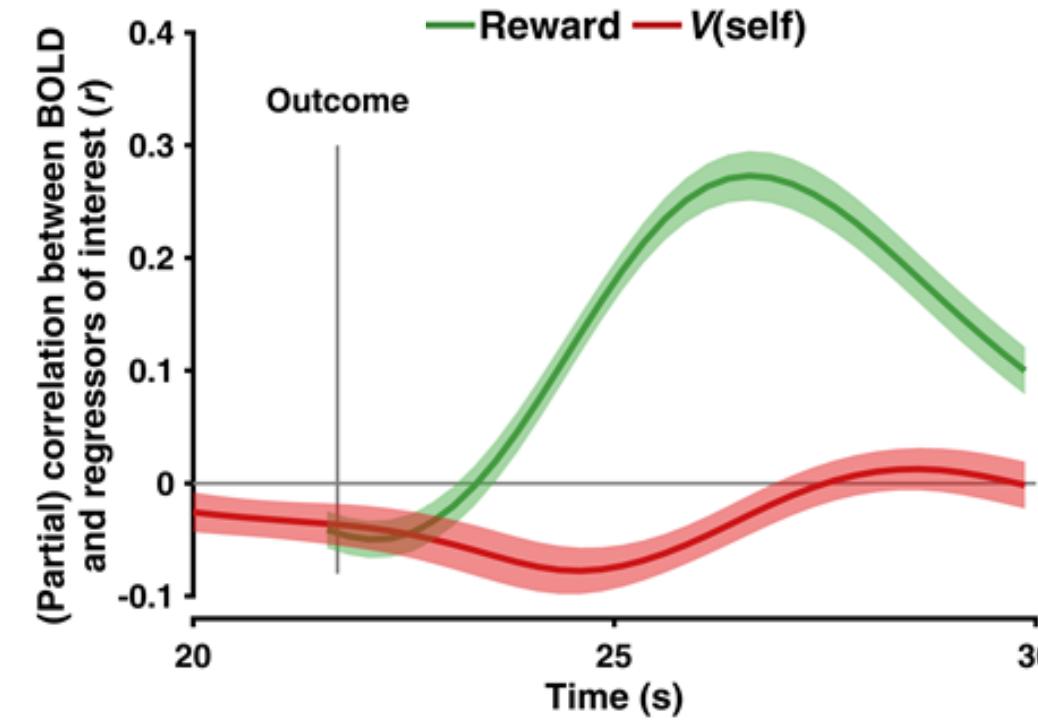
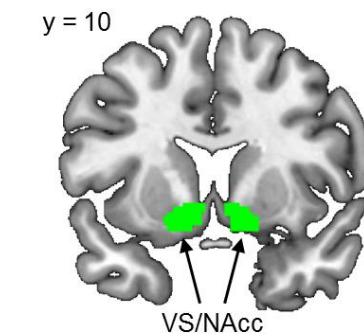
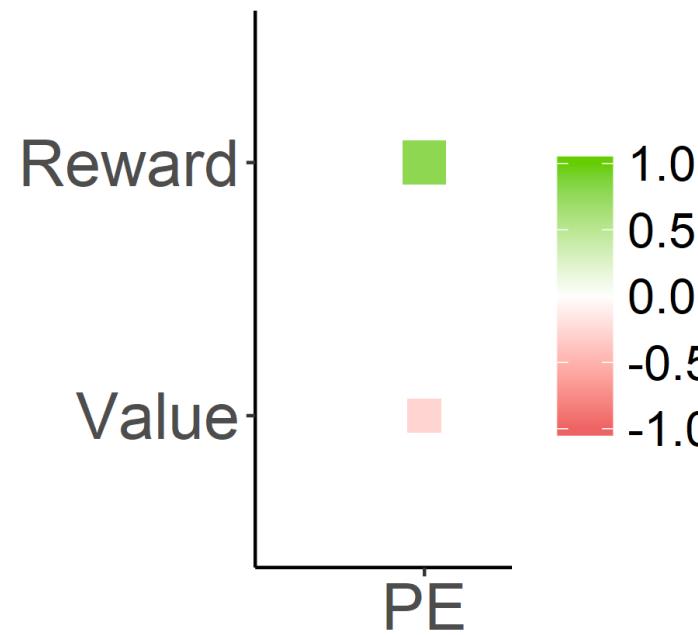


**Q:** how to justify the striatal activity is indeed associated with PE, rather than reward?

# A closer look at PE

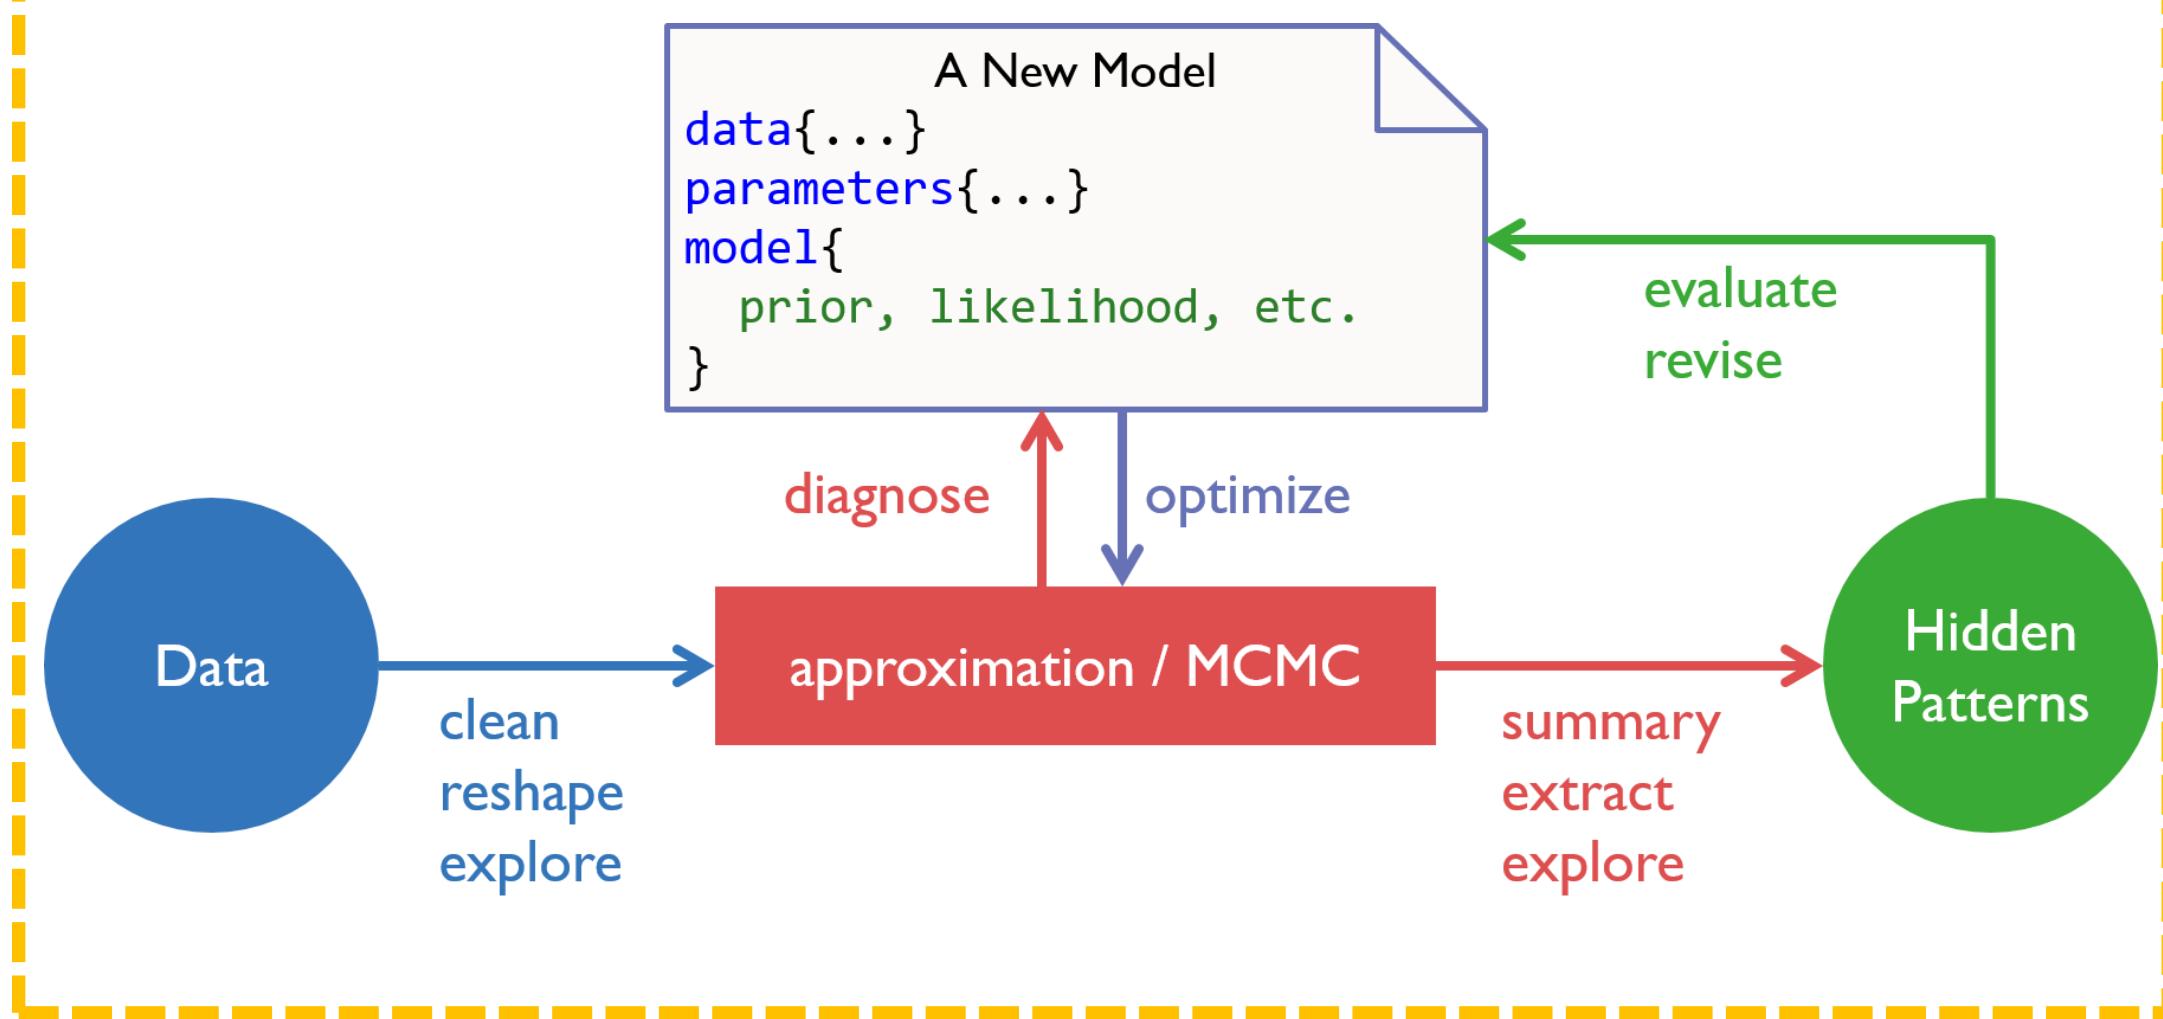
Prediction error:

$$PE = R_t - V_t$$



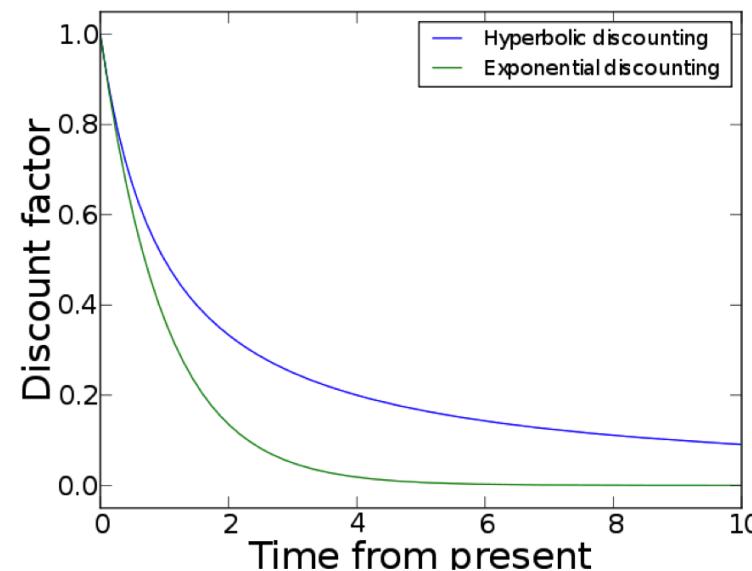
# **DELAY DISCOUNTING**





# Delay Discounting Task and Models

cognitive model  
statistics  
computing



## Hyperbolic Discounting Model

$$SV = \frac{A}{1 + k * delay}$$

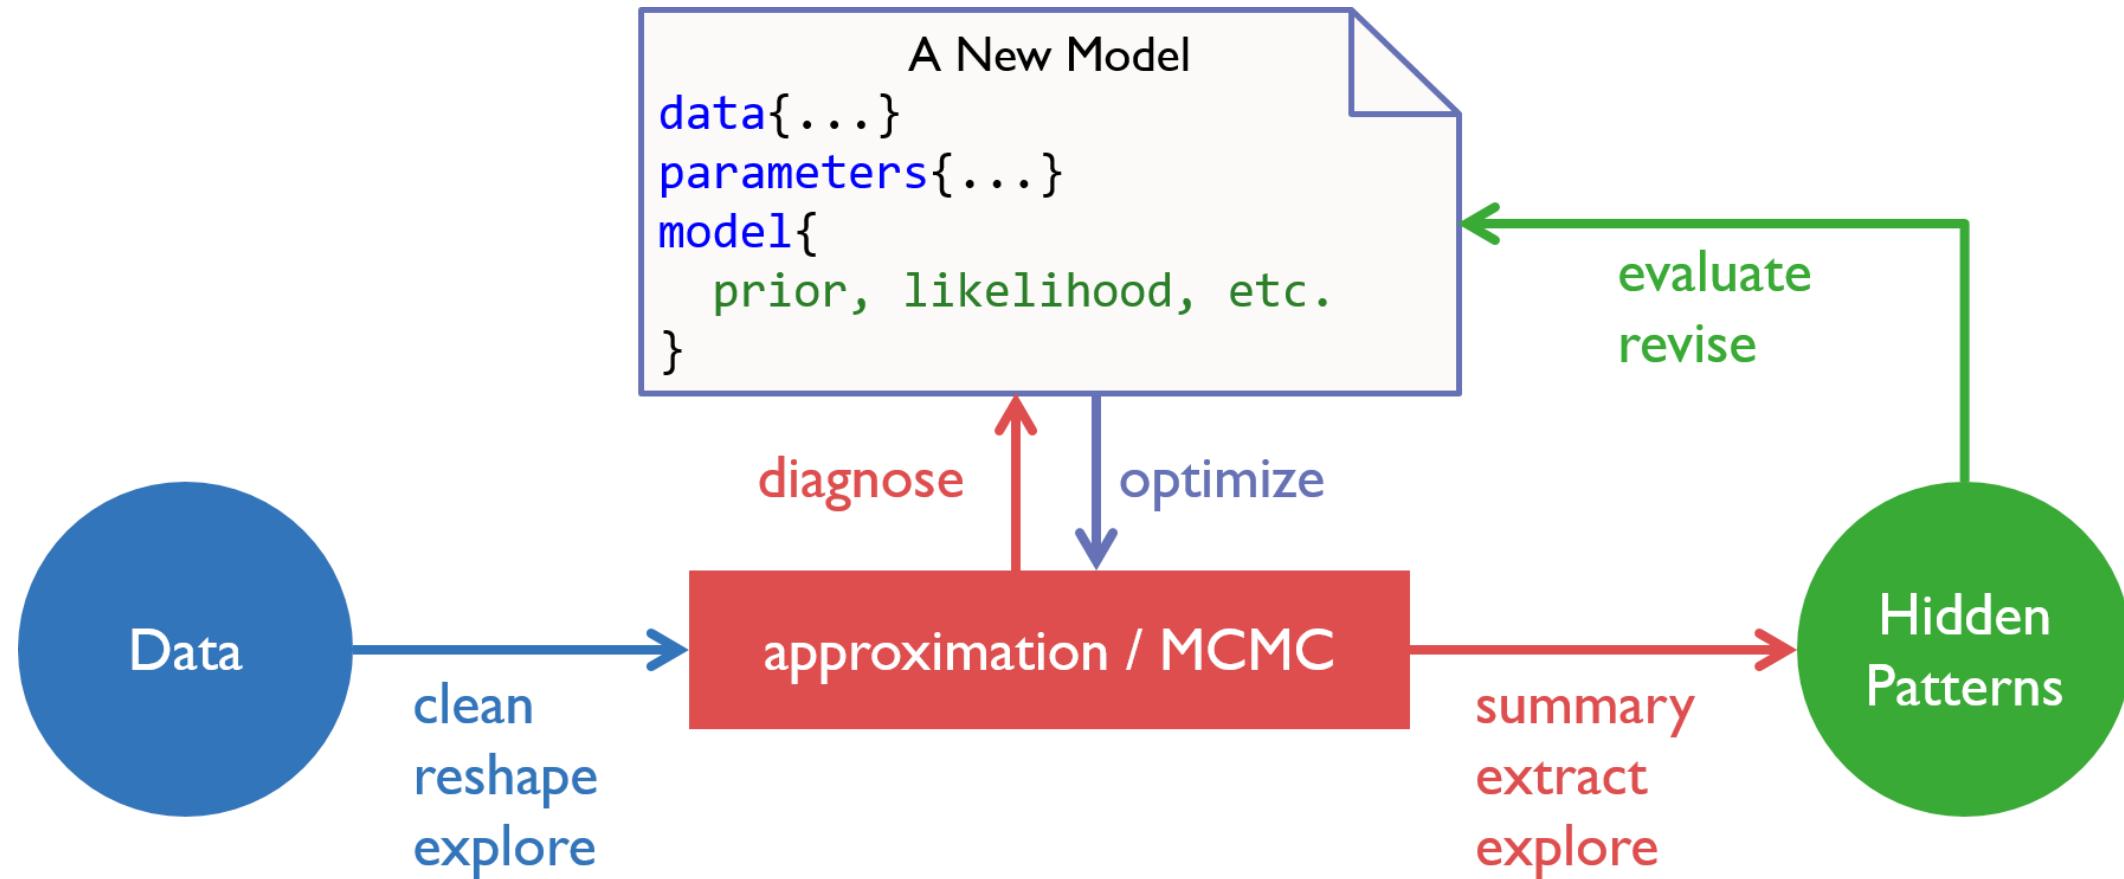
## Exponential Discounting Model

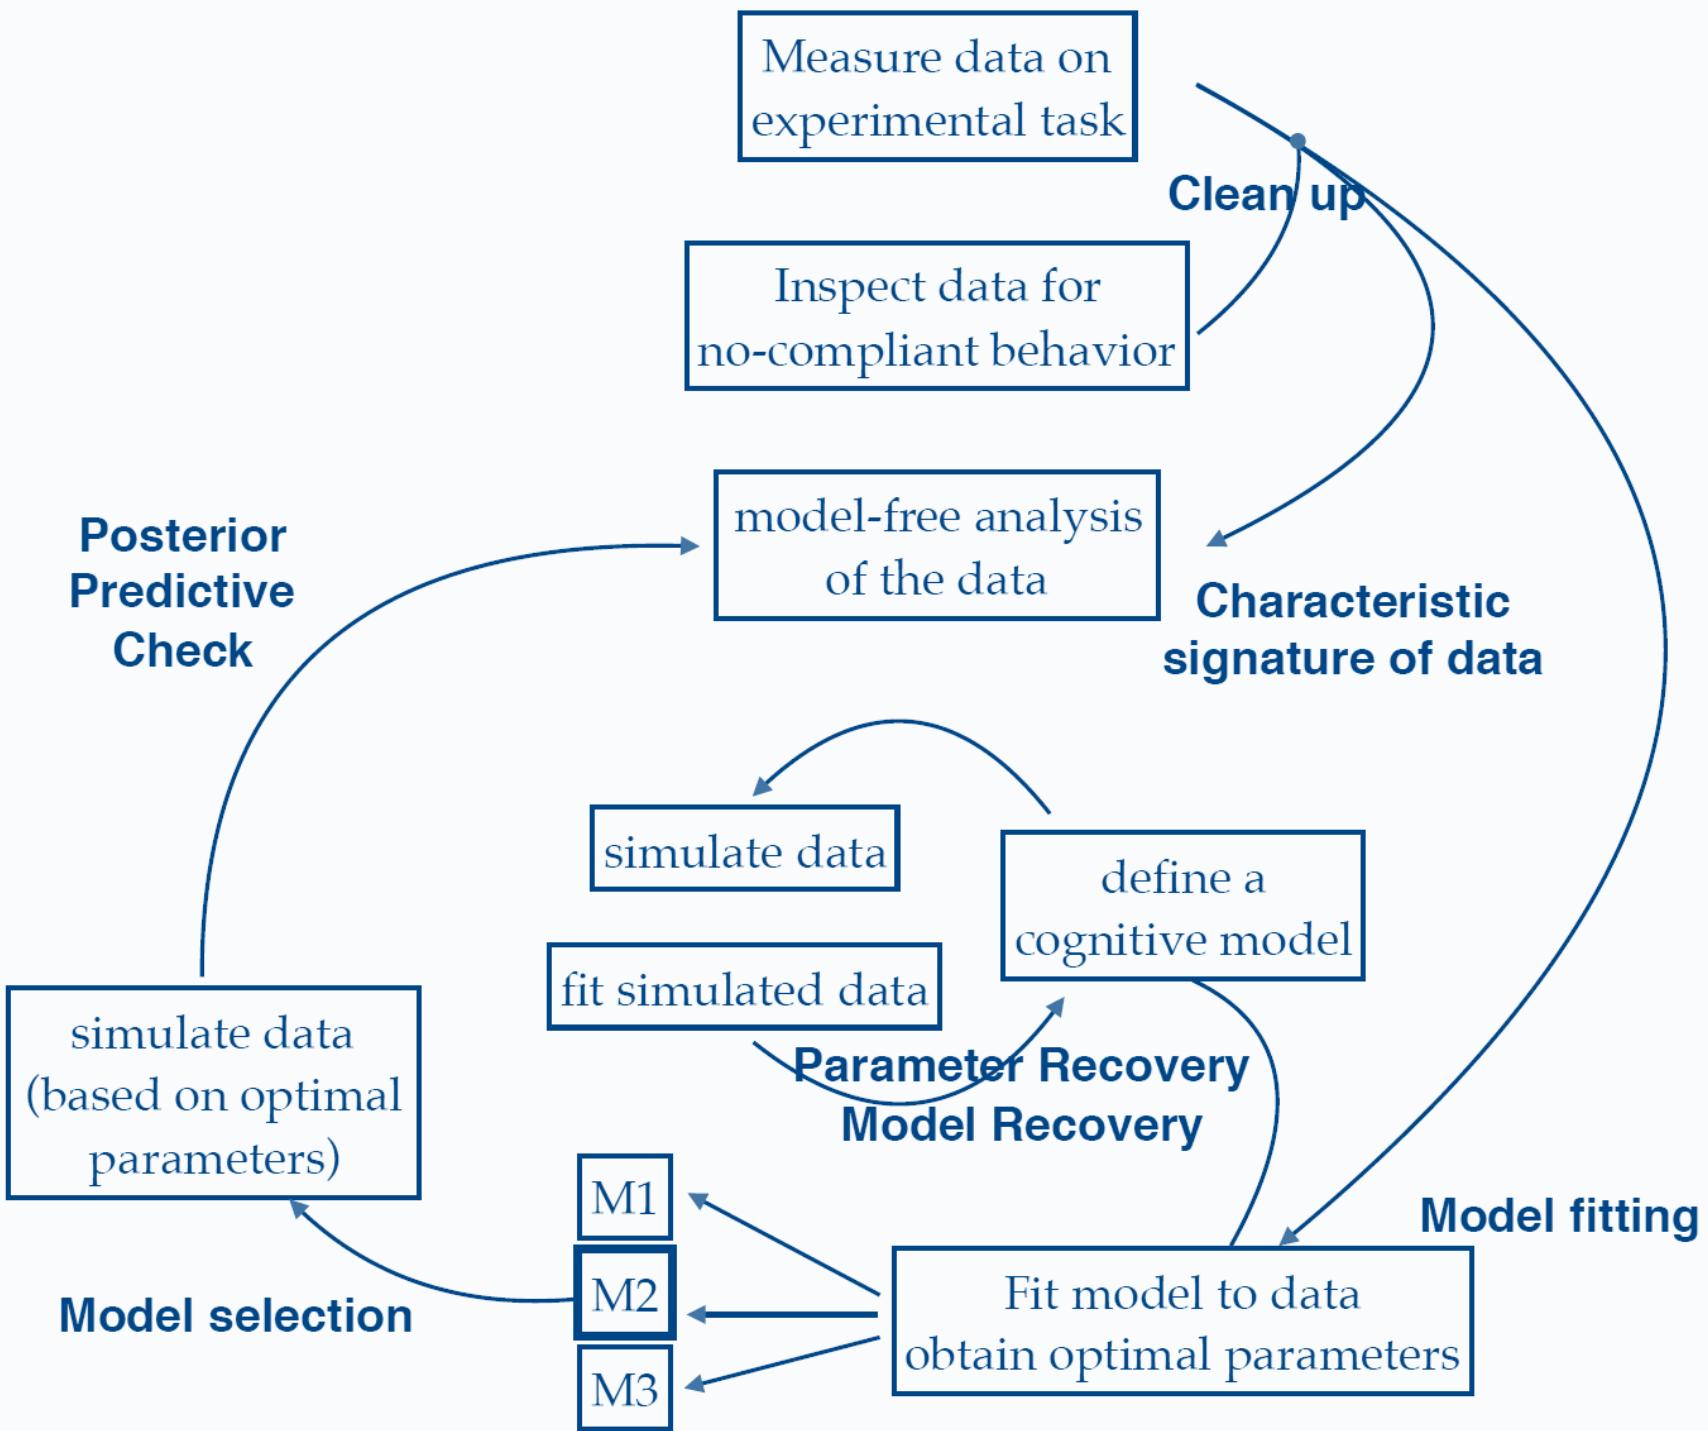
$$SV = A * \exp(-r * delay)$$
$$p(LL) = \frac{1}{1 + \exp^{\text{temp}(v(SS) - v(LL))}}$$

LL - late large option  
SS - soon small option

# Summary

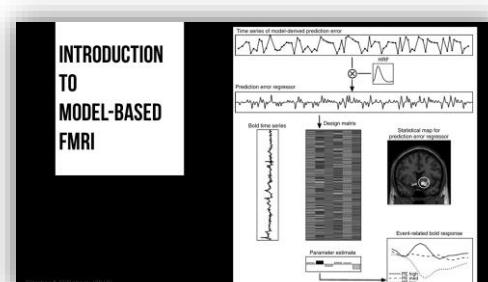
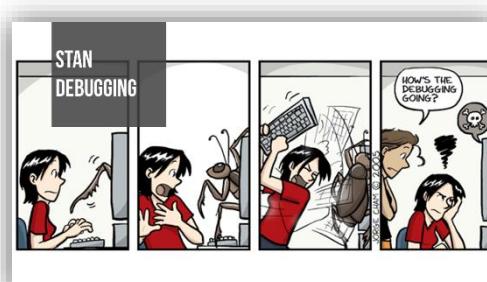
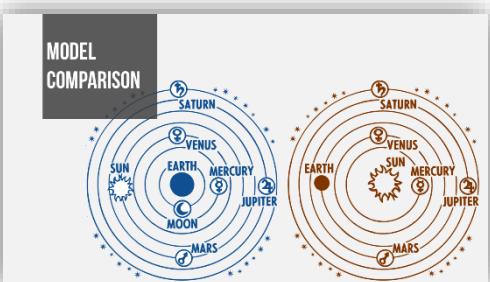
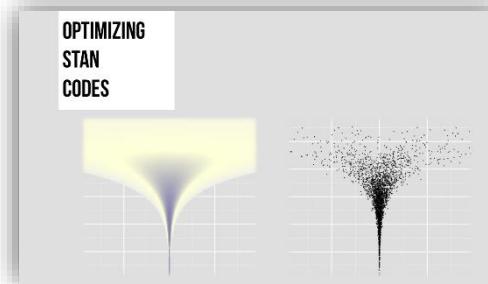
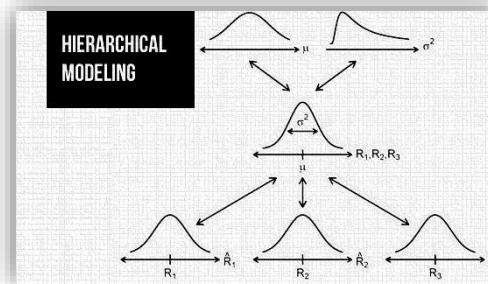
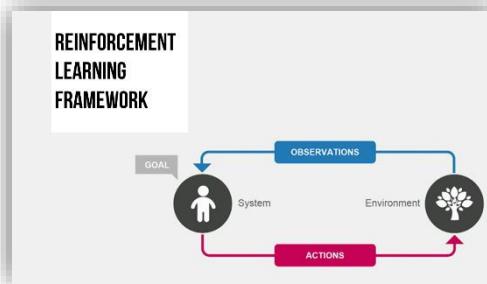
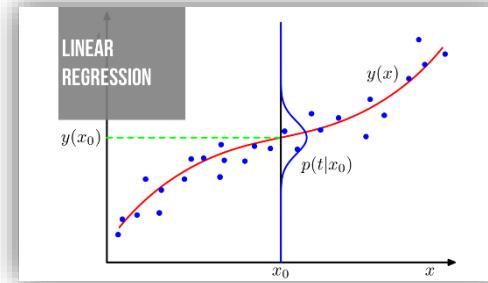
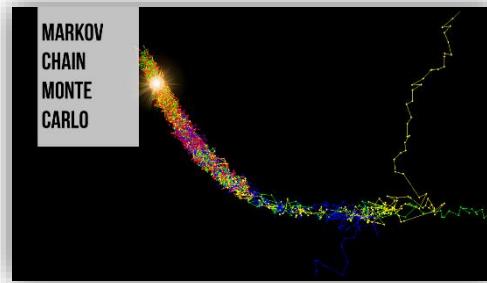
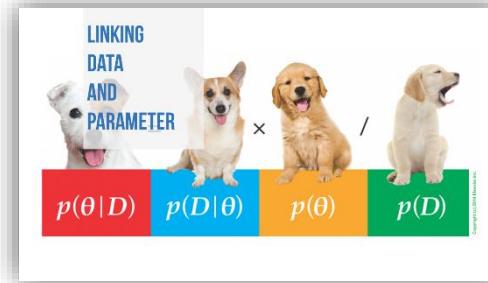
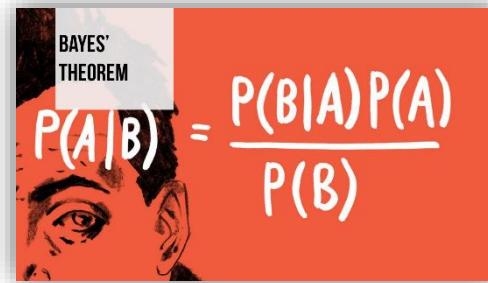
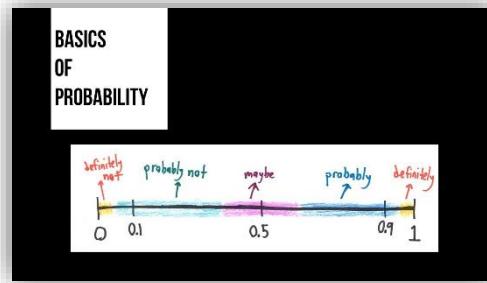
---





Adapted from Jan Gläscher's workshop

# Summary of Topics



# Summary of Examples/Exercises

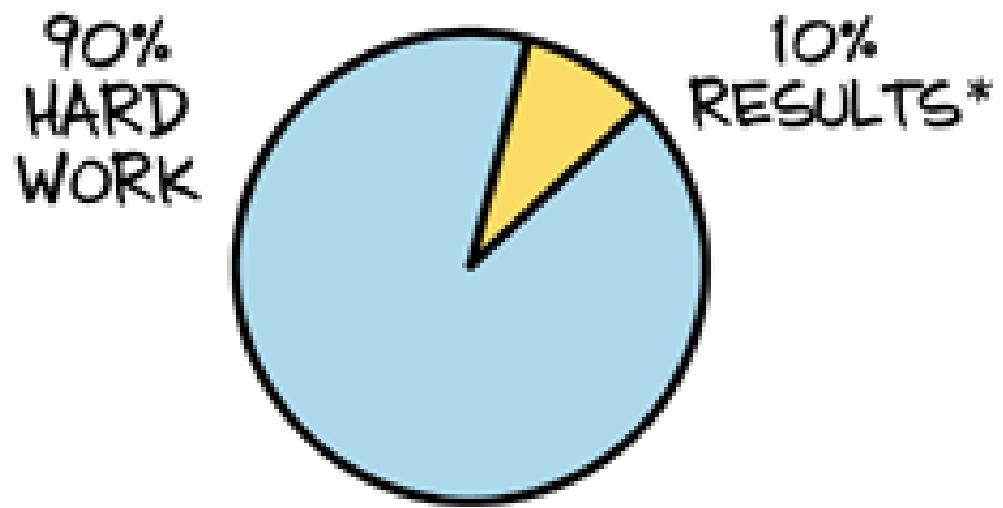
FOLDER	TASK	MODEL
01.R_basics	NA	NA
02.binomial_globe	Globe toss	Binomial Model
03.bernoulli_coin	Coin flip	Bernoulli Model
04.regression_height	Observed weight and height	Linear regression model
05.regression_height_poly		
06.reinforcement_learning	2-armed bandit task	Simple reinforcement learning (RL) model
07.optm_rl		
08.compare_models	Probabilistic reversal learning task	Simple and fictitious RL models
09.debugging	Memory Retention	Exponential decay model
10.model_based	2-armed bandit task	Simple RL model
11.delay_discounting	Delay discounting task	Hyperbolic and exponential discounting model

# After the Workshop, you...

- ...are able to implement your own model
- ...feel comfortable with reading mathematical equations
- ...consider the implementation of the “computational modeling” section
- ...gain insightful understanding of Bayesian stats and modeling
- ...take it as a good start and work on it later

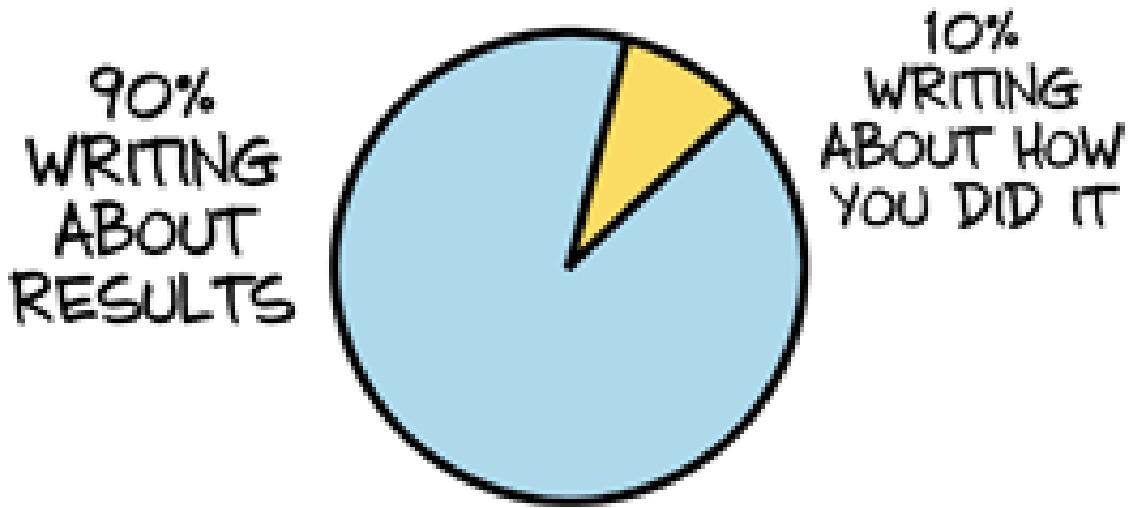
**Remember: practice makes perfect!**

## DOING RESEARCH:



\* BEST CASE SCENARIO

## WRITING ABOUT RESEARCH:



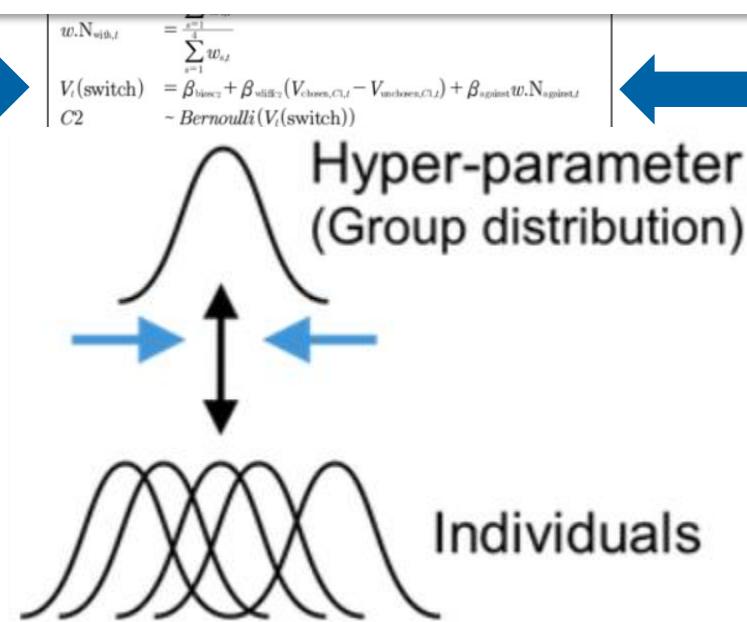
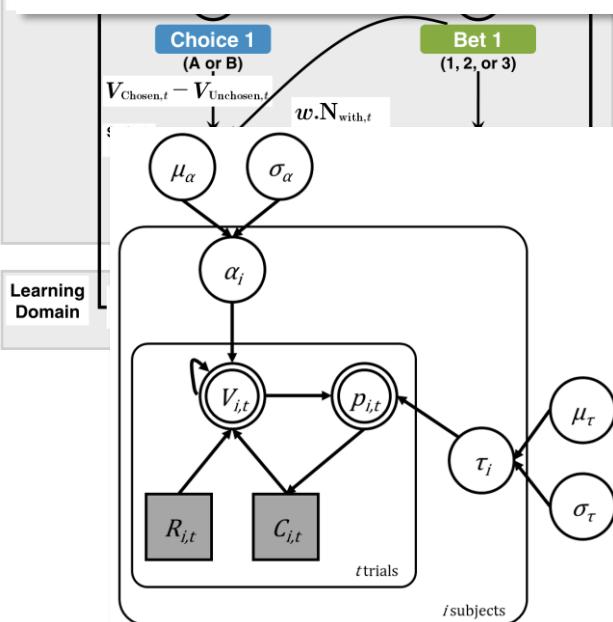
# hBayesDM package

## hBayesDM

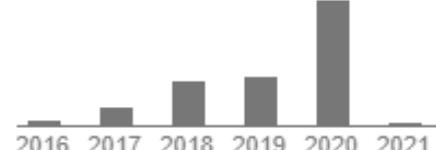
repo status Active build passing CRAN 1.0.2 – 2019-11-13 downloads 33K

DOI 10.1162/CPSY\_a\_00002

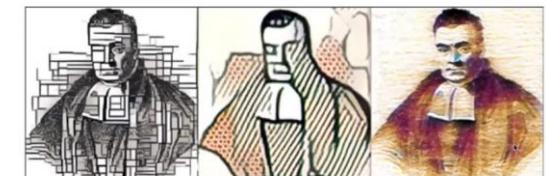
hBayesDM (hierarchical Bayesian modeling of Decision-Making tasks) is a user-friendly package that offers hierarchical Bayesian analysis of various computational models on an array of decision-making tasks. hBayesDM uses Stan for Bayesian inference.



Cited by 115



Stan



Bayesian Statistics and Hierarchical Bayesian Modeling for Psychological Science

Lecture 01

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)  
Department of Cognition, Emotion, and Methods in Psychology

[https://github.com/lei-zhang/BayesCog\\_Wien](https://github.com/lei-zhang/BayesCog_Wien)



## Revealing Neurocomputational Mechanisms of Reinforcement Learning and Decision-Making With the hBayesDM Package

Woo-Young Ahn<sup>1</sup>, Nathaniel Haines<sup>1</sup>, and Lei Zhang<sup>2</sup>

<sup>1</sup>Department of Psychology, The Ohio State University, Columbus, OH

<sup>2</sup>Institute for Systems Neuroscience, University Medical Center Hamburg-Eppendorf, Hamburg, Germany

**Keywords:** Reinforcement learning; Decision-making; Hierarchical Bayesian modeling; Model-based fMRI

cognitive model  
statistics  
computing

Task (alphabetical order)	Model name	hBayesDM function	References (see below for full citations)
Balloon Analogue Risk Task	4 parameter model	bart_4par	Wallsten et al. (2005)
Choice reaction time Task	Drift diffusion model Linear Ballistic Accumulator model	choiceRT_ddm choiceRT_lba	Ratcliff (1978) S. Brown & Heathcote (2008) Annis et al. (2017)
Choice under Risk and Ambiguity (CRA) Task	Linear model Exponential model	cra_linear cra_exp	Levy et al. (2009)
Delay Discounting Task	Constant-Sensitivity (CS) model Exponential model Hyperbolic model	dd_cs dd_exp dd_hyp	Ebert & Prelec (2007) Samuelson (1937) Mazur (1987)
Iowa Gambling Task (IGT)	Prospect Valence Learning-DecayRI Prospect Valence Learning-Delta Value-Plus-Perseverance (VPP) Outcome-Represent. Learning (ORL)	igt_pvl_decay igt_pvl_delta igt_vpp igt_orl	Ahn et al. (2011; 2014) Ahn et al. (2008) Worthy et al. (2013) Haines et al. (in press)
Orthogonalized Go/NoGo Task	RW+noise RW+noise+go bias RW+noise+go bias+Pav. bias M5 (see Table 1 of the reference)	gng_m1 gng_m2 gng_m3 gng_m4	Guitart-Masip et al. (2012) Guitart-Masip et al. (2012) Guitart-Masip et al. (2012) Cavanagh et al. (2013)
Peer influence task	Other-conferred utility (OCU)	peer_ocu	Chung et al. (2015)
Probabilistic Reversal Learning (PRL) Task	Experience-Weighted Attraction Fictitious update Reward-Punishment (Rew.-Pun.) Fictitious + Rew.-Pun. Fictitious + Rew.-Pun. w/o alpha Fictitious w/o alpha	prl_ewa prl_fictitious prl_rp prl_fictitious_rp prl_fictitious_rp_woa prl_fictitious_woa	Ouden et al. (2013) Glässcher et al. (2009) Ouden et al. (2013)
Probabilistic Selection Task	Q-learning with two learning rates	pst_gainloss_Q	M. J. Frank et al. (2007)
Risk-Aversion Task	Prospect Theory (PT) PT without loss aversion (LA) PT without risk aversion (RA)	ra_prospect ra_noLA ra_noRA	Sokol-Hessner et al. (2009) Tom et al. (2007)
Risky Decision Task	Happiness model	rdt_happiness	Rutledge et al. (2014)
Two-Armed Bandit (Experience-based) Task	Rescorla-Wagner (delta) model	bandit2arm_delta	Erev et al. (2010) Hertwig et al. (2004)
Two Step (TS) Task	7 parameter model 6 parameter model 4 parameter model	ts_7par ts_6par ts_4par	Daw et al. (2011) Wunderlich et al. (2012)
Four-Armed Bandit (Experience-based) Task	Fictive upd.+rew/pun sens. Fictive upd.+rew/pun sens.+lapse	bandit4arm_4par bandit4arm_lapse	Seymour et al. (2012) Seymour et al. (2012)
Ultimatum Game	Ideal Bayesian observer model Rescorla-Wagner (delta) model	ug_bayes ug_delta	Xiang et al. (2013) Gu et al. (2015)
Wisconsin Card Sorting Task	Sequential learning model	wcs_sql	A. J. Bishara et al. (2010)



# Workshops / Summer schools

- Cognitive Modeling Academy Hamburg (inaugural, 2021!)
- JAGS and WinBUGS Workshop @ Amsterdam, NL (annual)
- Model-based Neuroscience Summer School @ Amsterdam, NL (annual)
- European Summer School on Computational and Mathematical Modeling of Cognition @ multiple EU sites (biannual)
- Computational Psychiatry Course @ Zürich, CH (annual)
- London Computational Psychiatry Course @ London, UK (annual?)
- Methods in Neuroscience at Dartmouth Computational Summer School @ Dartmouth, USA (annual)
- Brains, Minds & Machines Summer Course @ MIT, USA (annual)
- Kavli Summer Institute in Cognitive Neuroscience @UCSB, USA (annual)
- Neuromatch Academy on Computational Neuroscience @online! (annual)

# References

- Ahn, W.-Y., Haines, N., & Zhang, L. (2017). Revealing Neurocomputational Mechanisms of Reinforcement Learning and Decision-Making With the hBayesDM Package. *Computational Psychiatry*, 1, 24–57.
- Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., ... & Riddell, A. (2017). Stan: A probabilistic programming language. *Journal of statistical software*, 76(1).
- Cohen, J. D., Daw, N., Engelhardt, B., Hasson, U., Li, K., Niv, Y., ... & Willke, T. L. (2017). Computational approaches to fMRI analysis. *Nature Neuroscience*, 20(3), 304.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., & Rubin, D. B. (2014). *Bayesian data analysis* (3rd ed.). New York, NY: CRC Press.
- Gelman, A., & Rubin, D. B. (1992). Inference from iterative simulation using multiple sequences. *Statistical Science*, 7, 457–472.
- Gläscher, J., Hampton, A. N., & O'Doherty, J. P. (2009). Determining a role for ventromedial prefrontal cortex in encoding action-based value signals during reward-related decision making. *Cerebral Cortex*, 19(2), 483-495.
- Gläscher, J. & O'Doherty, J. P. (2010). Model-based approaches to neuroimaging: combining reinforcement learning theory with fMRI data. *Wiley Interdisciplinary Reviews: Cognitive Science*, 1(4), 501-510.
- Hampton, A. N., Adolphs, R., Tyszka, J. M., & O'Doherty, J. P. (2007). Contributions of the amygdala to reward expectancy and choice signals in human prefrontal cortex. *Neuron*, 55(4), 545-555.
- Kruschke, J. (2014). *Doing Bayesian data analysis: A tutorial with R, JAGS, and Stan*. Academic Press.
- Lee, M. D., & Wagenmakers, E. J. (2014). *Bayesian cognitive modeling: A practical course*. Cambridge university press.
- Lewandowsky, S., & Farrell, S. (2010). *Computational modeling in cognition: Principles and practice*. Sage Publications.
- McElreath, R. (2018). *Statistical Rethinking: A Bayesian Course with Examples in R and Stan*. CRC Press.
- O'Doherty, J. P., Hampton, A., & Kim, H. (2007). Model-based fMRI and its application to reward learning and decision making. *Annals of the New York Academy of Sciences*, 1104(1), 35-53.
- Rescorla, R. A., & Wagner, A. R. (1972). A theory of Pavlovian conditioning: Variations in the effectiveness of reinforcement and nonreinforcement. *Classical conditioning II: Current research and theory*, 2, 64-99.
- Sutton, R. S., Barto, A. G., & Bach, F. (1998). *Reinforcement learning: An introduction*. MIT press.
- Vehtari, A., Gelman, A., & Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, 27(5), 1413-1432.

# Acknowledgement

Antonius Wiegler (MMB, Paris )

Jan Gläscher (UKE, Hamburg)

Woo Young Ahn (SNU, Seoul)

Nate Haines (OSU, OH)

ANY  
QUESTIONS  
?

Happy Computing!